

Fast algorithms using orthogonal polynomials

Sheehan Olver

*Department of Mathematics,
Imperial College,
London SW7 2AZ, UK*
E-mail: s.olver@imperial.ac.uk

Richard Mikaël Slevinsky

*Department of Mathematics,
University of Manitoba,
Winnipeg, Manitoba, Canada R3T 2M8*
E-mail: Richard.Slevinsky@umanitoba.ca

Alex Townsend

*Department of Mathematics,
Cornell University,
Ithaca, NY 14853, USA*
E-mail: townsend@cornell.edu

We review recent advances in algorithms for quadrature, transforms, differential equations and singular integral equations using orthogonal polynomials. Quadrature based on asymptotics has facilitated optimal complexity quadrature rules, allowing for efficient computation of quadrature rules with millions of nodes. Transforms based on rank structures in change-of-basis operators allow for quasi-optimal complexity, including in multivariate settings such as on triangles and for spherical harmonics. Ordinary and partial differential equations can be solved via sparse linear algebra when set up using orthogonal polynomials as a basis, provided that care is taken with the weights of orthogonality. A similar idea, together with low-rank approximation, gives an efficient method for solving singular integral equations. These techniques can be combined to produce high-performance codes for a wide range of problems that appear in applications.

CONTENTS

1	Introduction	574
2	Orthogonal polynomials	576
3	Quadrature	591
4	Orthogonal polynomial transforms	598
5	Ordinary differential equations	625
6	Partial differential equations	642
7	Singular integral equations	653
8	Conclusions	668
A	Classical orthogonal polynomial recurrence relationships	669
B	Multivariate orthogonal polynomial and their recurrence relationships	677
C	Semiclassical orthogonal polynomials	686
D	Notation	687
	References	691

1. Introduction

Orthogonal polynomials (OPs) play a fundamental role in numerical analysis and computational mathematics with applications in quadrature, function approximation and differential equations (Szegő 1975). In the past ten years there has been a resurgence of research attention on the classical properties of orthogonal polynomials, which has led to more efficient algorithms in classical applications as well as new directions. This article aims to give an extensive survey of these recent results. There are three main reasons for the revival of orthogonal polynomials: (1) newfound uses for rigorous asymptotic expansions, (2) the exploitation of sparse and low-rank structure in recurrence relationships, and (3) the introduction of Chebfun, and packages such as ApproxFun and Dedalus, that leverage orthogonal polynomials to automate computations involving functions and differential equations. A general underlying theme is to combine recurrence relationships and asymptotic expansions with structured linear algebra to achieve fast algorithms.

Rigorous asymptotics of special functions is a traditional area of numerical analysis, used to develop trustworthy numerical evaluation schemes for special functions. Two essential types of asymptotic expansions are trigonometric (where a special function is written by a linear combination of cosines and sines) and uniform. The resurgence of interest in these expansions is due to their use in exploiting regions where the special function is trigonometric-like, allowing for a range of fast algorithms based on the fast Fourier transform (FFT) and state-of-the-art numerical linear algebra.

Sparsity in recurrence relationships is also classical, with their use in fast numerical algorithms for differential equations going back to Clenshaw (1957). The success of finite element methods and the development of fast black-box linear algebra reduced their usage. However, OPs have remained popular for computing with functions and solving differential equations on simple geometries (*i.e.* rectangles, disks, spheres and balls) and in spectral element methods (see Karniadakis and Sherwin 2013 for an extensive overview). In particular, they are important in the astrophysics and computational fluid dynamics communities, where simple geometries are prevalent (Burns *et al.* 2020).

The success of software packages such as ApproxFun (Olver, Goretkin, Slevinsky and Townsend 2016), Chebfun (Driscoll, Hale and Trefethen 2014) and Dedalus (Burns *et al.* 2020) has renewed interest in polynomial approximation and OPs, as the software systems allow one to manipulate functions and accurately solve differential equations with a few lines of Julia, MATLAB and Python code, respectively. These software systems have also led to novel ways to exploit sparsity in the recurrence relationships for OPs. While this survey is primarily on the underlying mathematics, there is a strong link between mathematical research and software design in this subject.

1.1. Outline of paper

Section 2. We give an overview of univariate and multivariate OPs and their properties, including the classical OPs: Chebyshev, Legendre, ultraspherical, Jacobi, Laguerre and Hermite. Unlike standard introductions, we phrase their construction and properties in terms of *quasimatrices*. We find this convenient when connecting structured recurrence relationships to structured numerical linear algebra.

Section 3. We discuss the efficient calculation of quadrature rules using OPs, beginning with the well-known FFT-based algorithms for Clenshaw–Curtis quadrature. We then discuss more recent algorithms for calculating Gauss quadrature rules, building on careful use of asymptotic expansions for calculating the nodes and weights.

Section 4. We discuss fast transforms for expansions in OPs and spherical harmonics. There are two categories of transforms: transform of coefficients between different families of OPs (*e.g.* between Chebyshev and Legendre expansions, or between spherical harmonics and tensored Fourier expansions) and transforms between values on a grid and coefficients. The latter typically consists of first converting to a simpler basis using a transform based on the FFT, whereas the former utilizes hidden structure in change-of-basis (connection) matrices.

Section 5. Univariate OPs can be used to transform ordinary differential equations to sparse linear systems using recurrence relationships, facilitating optimal complexity solvers: $O(N)$ cost for N degrees of freedom. In this vein, we present three methods in a unified framework: integral reformulation as introduced by Clenshaw (1957), the ultraspherical spectral method of Olver and Townsend (2013) and a one-dimensional version of the p -finite element method of Shen (1994). Each method has its benefits and trade-offs: integral reformulation is the most straightforward, the ultraspherical spectral method the most versatile. In contrast, the p -finite element method preserves self-adjointness and positive definiteness of the operator. We also discuss domain decomposition, applications to eigenvalue problems and symmetrizing non-symmetric discretizations, which improve the complexity of algorithms for computing eigenvalues.

Section 6. Similarly, multivariate OPs can be used to transform partial differential equations to sparse linear systems. In the case of rectangular geometries this proceeds using a tensor product basis, in which case the sparse linear systems arise from Kronecker products of ordinary differential equations. Naïve direct linear algebra achieves $O(N^2)$ complexity for N degrees of freedom. This can be reduced to $O(N^{3/2})$ for operators that can be expressed as a sum of two Kronecker products using the Schur factorization, and improved further to quasi-optimal complexity for the Poisson equation. Going beyond rectangles, the ultraspherical spectral method and p -finite element method have natural extensions to triangles using recurrence relationships, leading to sparse operators. We demonstrate these techniques on simple examples such as the inhomogeneous and variable coefficient Helmholtz equation and the biharmonic equation.

Section 7. OPs have a natural link to singular integral operators such as the Hilbert, Cauchy/Stieltjes and log-kernel transforms, which can be exploited for the calculation of singular integrals and solution of singular integral equations such as those arising from boundary integral reformulations of partial differential equations. In particular, an ultraspherical-like spectral method is achievable, reducing singular integral equations to sparse linear systems that can be solved efficiently. Low-rank approximation techniques play a role here for handling general kernels.

2. Orthogonal polynomials

A sequence of OPs is a family of polynomials that are orthogonal to each other under some inner product (Szegő 1975). The field of OPs developed from the study of continued fractions, and are now used throughout approximation theory, numerical analysis and computational physics.

2.1. Univariate orthogonal polynomials

Let $p_0(x), p_1(x), p_2(x), \dots$ be a sequence of real-valued polynomials such that $p_n(x)$ is exactly of degree n , that is,

$$p_n(x) = k_n x^n + O(x^{n-1}), \quad k_n \neq 0.$$

We say that $p_0(x), p_1(x), p_2(x), \dots$ are *orthogonal polynomials* with respect to a non-negative, continuous, weight function w over a (possibly infinite) interval (a, b) if

$$\langle p_m, p_n \rangle_w = \int_a^b w(x) p_m(x) p_n(x) dx = \begin{cases} 0 & m \neq n, \\ h_n & m = n. \end{cases} \quad (2.1)$$

If $h_n = 1$, then the polynomials are not just orthogonal but orthonormal. Given a weight function, the orthonormal polynomials associated with w are unique up-to-sign.

Any set of OPs can be rescaled to be orthonormal; however, the standard normalizations (Szegő 1975) are determined by prescribing their values at special points. With the standard definitions, classical OPs have simple rational recurrence relationships that are extremely efficient to generate, whereas their normalized counterparts involve square roots. In practical computations, this is typically helpful, though we deviate from this convention for legitimate reasons from time to time.

2.2. Quasimatrix notation and orthogonal polynomial expansions

Many applications of OPs use them as a convenient basis in which to represent functions. Let $p_0(x), p_1(x), p_2(x), \dots$ be a sequence of univariate OPs on $[a, b]$ with respect to the weight function w ; then

$$f(x) = \sum_{n=0}^{\infty} p_n(x) f_n = \mathbf{P}(x) \mathbf{f}, \quad \mathbf{P}(x) = (p_0(x) \mid p_1(x) \mid \dots),$$

where $\mathbf{f} = (f_0, f_1, \dots)^{\top}$ is an infinite vector and $\mathbf{P}(x)$ is a $1 \times \infty$ row-vector.

We view \mathbf{P} as an $(a, b) \times \infty$ quasimatrix,¹ a formalism that allows us to view it as a linear map from coefficients to functions. That is, we can view \mathbf{P} as the *synthesis operator*: a one-to-one linear map from the space of infinite vectors with a finite number of non-zero entries ($\ell_F \subset \mathbb{C}^\infty$) to the space of polynomials ($\mathbb{C}[x]$). Of course, we can go further and relate other coefficient spaces to function spaces (*e.g.* with Legendre polynomials it is a one-to-one map from ℓ^2 to $L^2[-1, 1]$) and quasimatrices allow us to

¹ An $[a, b] \times n$ quasimatrix is a matrix with n columns, whose columns are functions defined on $[a, b]$ as opposed to vectors.

discuss these maps collectively, but for our discussion it suffices to interpret the map as $\mathbf{P}: \ell_F \rightarrow \mathbb{C}[x]$. We take the convention that

$$\mathbf{P}^\top f := \begin{pmatrix} \langle p_0, f \rangle \\ \langle p_1, f \rangle \\ \vdots \end{pmatrix},$$

where $\langle \cdot, \cdot \rangle$ is the $L^2(a, b)$ inner product. We therefore have the *analysis operator* $\mathbf{P}^\top \underline{w}$, where the underbar indicates that it is the multiplication by w operator. That is,

$$\mathbf{P}^\top \underline{w} f := \begin{pmatrix} \langle p_0, w f \rangle \\ \langle p_1, w f \rangle \\ \vdots \end{pmatrix}.$$

Given a polynomial $f(x)$ we can solve the following linear system:

$$\mathbf{P} \mathbf{f} = f$$

for the expansion coefficients, \mathbf{f} . Due to the orthogonality, we find that

$$\mathbf{P}^\top \underline{w} \mathbf{P} = D_h = \text{diag}(h_0, h_1, \dots),$$

where h_0, h_1, \dots are the orthogonality constants in (2.1). This means that

$$\mathbf{f} = \mathbf{P}^{-1} f = (\mathbf{P}^\top \underline{w} \mathbf{P})^{-1} \mathbf{P}^\top \underline{w} f = D_h^{-1} \mathbf{P}^\top \underline{w} f.$$

Equivalently, for $n \geq 0$, we have

$$f_n = \frac{1}{h_n} \int_a^b p_n(x) w(x) f(x) dx, \quad h_n = \int_a^b w(x) (p_n(x))^2 dx. \quad (2.2)$$

Due to the orthogonality of p_0, p_1, \dots the coefficient f_n can be computed independently of the other entries of \mathbf{f} .

For non-polynomial functions f , one needs to use truncated orthogonal polynomial expansions to approximate f . In particular, if f is a sufficiently smooth function on $[a, b]$, then one expects that the entries in \mathbf{f} decay to zero. The expansion can be truncated after N terms to obtain a polynomial approximation to f ,

$$f(x) \approx \sum_{n=0}^{N-1} f_n p_n(x) = \mathbf{P} P_N \mathbf{f},$$

where $P_N \mathbf{f} = (f_0, \dots, f_{N-1}, 0, 0, \dots)^\top$ is the truncation operator $P_N: \mathbb{C}^\infty \rightarrow \ell_F$. A polynomial approximation to f , formed in this way, can serve as a proxy to f and make subsequent computations more efficient (Boyd 2013). In principle, to compute a polynomial approximation to f , the coefficients f_n for $0 \leq n \leq N - 1$ need to be computed via the integral expression

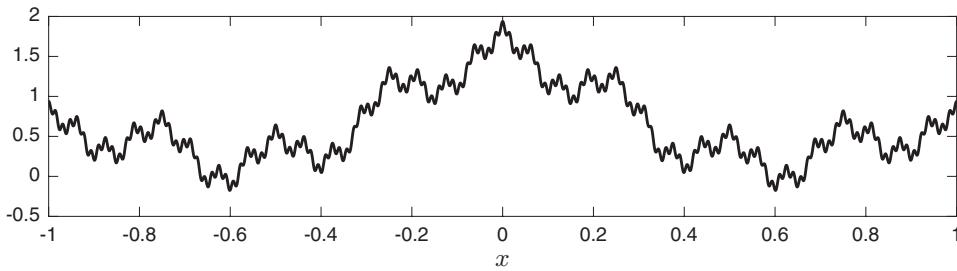


Figure 2.1. The function $f(x) = \sum_{k=0}^5 2^{-k} \cos((\pi/2)4^k x)$ approximated by a polynomial of degree 476 to within machine precision, where the polynomial is represented using a Chebyshev expansion.

in (2.2). However, this is computationally prohibitive for large N . Instead, fast transformation from samples of f to expansion coefficients are desirable (see Section 4). In Figure 2.1 we show a function approximated by a finite Chebyshev expansion. The degree of the polynomial used to represent this function was adaptively selected using the effective heuristic of Aurentz and Trefethen (2017).

Quasimatrices provide convenient notation for representing the action of continuous operators by discrete operators. For example, consider the quasimatrix formed by the monomials $\mathbf{Z}(z) = (1 | z | z^2 | \dots)$ defined on the unit disk. A simple relationship for differentiation is

$$\mathcal{D}\mathbf{Z}(z) = (0 \mid 1 \mid 2z \mid 3z^2 \mid \dots) = \mathbf{Z}(z) \begin{pmatrix} 0 & 1 & & \\ & 2 & & \\ & & \ddots & \end{pmatrix}, \quad (2.3)$$

where \mathcal{D} is the differentiation operator. This is clearly true when viewing this as an operator $\mathcal{D}\mathbf{Z}: \ell_F \rightarrow \mathbb{C}[z]$, with suitable extensions to sufficiently fast decaying coefficient spaces.

Remark 2.1. In what follows, we use quasimatrices without being specific about the spaces they act on. We note that it is usually straightforward to deduce these from the actual entries of the resulting operators. For example, for monomials, we have $\mathbf{Z}: \ell^2 \rightarrow H_+$, where H_+ is the Hardy space on the unit circle. One can also see from the growth in the operator in (2.3) that $\mathcal{D}\mathbf{Z}: \ell_1^2 \rightarrow H_+$ is bounded, where

$$\ell_1^2 = \left\{ \mathbf{f} \in \mathbb{C}^\infty : \sum_{n=0}^{\infty} ((n+1)\mathbf{f}_n)^2 < \infty \right\}.$$

That is, one can combine known relationships between spaces coming from approximation theory with trivial relationships deduced from the growth properties of banded operators to deduce precise statements.

2.3. The Jacobi operator associated to orthogonal polynomials

All sets of orthogonal univariate polynomials satisfy a three-term recurrence relationship. That is, all univariate OPs satisfy a recurrence of the form

$$xp_n(x) = c_n p_{n-1}(x) + a_n p_n(x) + b_n p_{n+1}(x), \quad n \geq 1, \quad (2.4)$$

$$xp_0(x) = a_0 p_0(x) + b_0 p_1(x). \quad (2.5)$$

One of the simplest ways to evaluate $p_n(x_0)$ at a point x_0 is to use the three-term recurrence (perhaps with some care) to successively evaluate $p_0(x_0), \dots, p_{n-1}(x_0)$, before finally computing $p_n(x_0)$ (Gautschi 1967). The recurrence in (2.5) can be encoded as a tridiagonal operator, known as a (non-symmetric) Jacobi operator.

Definition 2.2. A (non-symmetric) *Jacobi operator*, J , associated with a family of OPs is the tridiagonal operator given by

$$\mathbf{P}(x)^\top \underline{x} = \underbrace{\begin{pmatrix} a_0 & b_0 & & \\ c_1 & a_1 & b_1 & \\ & c_2 & a_2 & \ddots \\ & \ddots & \ddots & \ddots \end{pmatrix}}_J \mathbf{P}(x)^\top. \quad (2.6)$$

If the family of OPs are orthonormal, then J is a symmetric tridiagonal operator (Gautschi 2004).

The transpose of a Jacobi operator, denoted by $X := J^\top$, is the discrete operator that represents multiplication by x in the OP basis. This can be seen from (2.5) by transposing (2.6) to obtain the relationship $x\mathbf{P} = \mathbf{P}X$. That means that if $f(x) = \mathbf{P}(x)\mathbf{f}$, then the expansion coefficients of $xf(x)$ are $X\mathbf{f}$, that is,

$$xf(x) = x\mathbf{P}(x)\mathbf{f} = \mathbf{P}(x)X\mathbf{f}.$$

To multiply a function f expressed in an OP expansion by x , we can apply the operator X to the expansion coefficients of f .

2.4. Classical univariate orthogonal polynomials

Classical OPs are the most widely used univariate OPs and consist of Hermite polynomials, generalized Laguerre polynomials and Jacobi polynomials (including as special cases the ultraspherical, Chebyshev and Legendre polynomials). These polynomials have several characterizing properties that make them computationally useful and single them out from other OPs: (1) they are the only univariate OPs (up to a linear change of variables) whose derivatives are also OPs (Sonine 1887), and (2) they satisfy

Table 2.1. The classical OPs on bounded and unbounded real intervals. Up to a linear change of variables, the classical OPs are the only univariate OPs whose derivatives are also OPs.

Family	Notation	Interval	$w(x)$
Chebyshev (1st kind)	$T_n(x)$	$[-1,1]$	$(1-x^2)^{-1/2}$
Chebyshev (2nd kind)	$U_n(x)$	$[-1,1]$	$(1-x^2)^{1/2}$
Chebyshev (3rd kind)	$V_n(x)$	$[-1,1]$	$((1+x)/(1-x))^{1/2}$
Chebyshev (4th kind)	$W_n(x)$	$[-1,1]$	$((1-x)/(1+x))^{1/2}$
Legendre	$P_n(x)$	$[-1,1]$	1
Ultraspherical	$C_n^{(\lambda)}(x), \lambda > -\frac{1}{2}, \lambda \neq 0$	$[-1,1]$	$(1-x^2)^{\lambda-1/2}$
Jacobi	$P_n^{(\alpha,\beta)}(x), \alpha, \beta > -1$	$[-1,1]$	$(1-x)^\alpha(1+x)^\beta$
Laguerre	$L_n^{(\alpha)}(x), \alpha > -1$	$[0,\infty)$	$x^\alpha e^{-x}$
Hermite	$H_n(x)$	$(-\infty,\infty)$	e^{-x^2}

a fixed second-order Sturm–Liouville differential equation. The first property makes them ideal in spectral methods for solving differential equations, while (2) ensures that there are asymptotic expansions that hold for the entire family of OPs. Table 2.1 details the classical OPs, along with their orthogonality.

The Jacobi polynomials $P_n^{(\alpha,\beta)}(x)$ are orthogonal with respect to the weight function $w(x) = (1-x)^\alpha(1+x)^\beta$ with $\alpha, \beta > -1$ and are scaled so that

$$P_n^{(\alpha,\beta)}(x) = \frac{(n+\alpha+\beta+1)_n}{2^n n!} x^n + \dots,$$

where $(\alpha)_\beta$ is the Pochhammer symbol. They are useful special functions for singular integral equations, quadrature rules, building multivariable OPs, and random matrix theory. In quasimatrix notation, we may write

$$\mathbf{P}^{(\alpha,\beta)}(x) := (P_0^{(\alpha,\beta)}(x) | P_1^{(\alpha,\beta)}(x) | \dots).$$

Special cases of Jacobi polynomials include ultraspherical

$$C_n^{(\lambda)}(x) = (2\lambda)_n / (\lambda + 1/2)_n P_n^{(\lambda-1/2, \lambda-1/2)}(x),$$

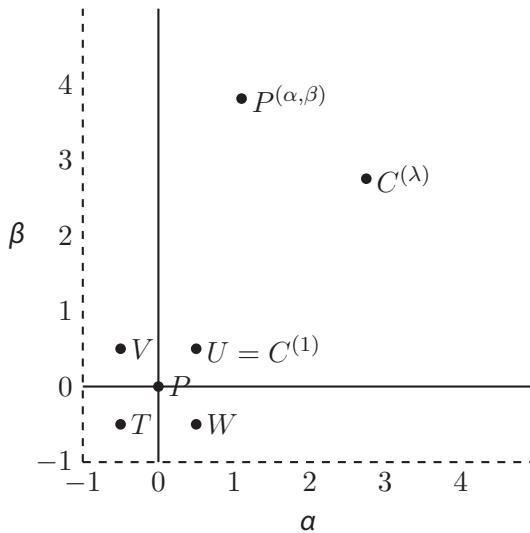


Figure 2.2. Jacobi polynomials $P^{(\alpha,\beta)}$ for $\alpha, \beta > -1$ are orthogonal with respect to $w(x) = (1-x)^\alpha(1+x)^\beta$. Along the diagonal $\alpha = \beta = \lambda - 1/2$ the OPs are referred to as *ultraspherical polynomials* and are denoted by $C^{(\lambda)}$ for $\lambda > -1/2$, $\lambda \neq 0$.

Legendre $P_n(x) = P_n^{(0,0)}(x)$, and Chebyshev polynomials (four kinds):²

$$\begin{aligned} T_n(x) &= P_n^{(-1/2, -1/2)}(x)/P_n^{(-1/2, -1/2)}(1), \\ U_n(x) &= (n+1)P_n^{(1/2, 1/2)}(x)/P_n^{(1/2, 1/2)}(1), \\ V_n(x) &= P_n^{(-1/2, 1/2)}(x)/P_n^{(-1/2, 1/2)}(1), \\ W_n(x) &= (2n+1)P_n^{(1/2, -1/2)}(x)/P_n^{(1/2, -1/2)}(1). \end{aligned}$$

Associated to these OPs are the quasimatrices $\mathbf{C}^{(\lambda)}(x)$, $\mathbf{P}(x)$, $\mathbf{T}(x)$, $\mathbf{U}(x)$, $\mathbf{V}(x)$ and $\mathbf{W}(x)$, respectively. Figure 2.2 summarizes the special cases of Jacobi polynomials. The figure is useful for understanding and exploring sparse recurrence relationships for derivatives, weighted derivatives and change-of-basis (conversion) operators.

2.5. Multivariate orthogonal polynomials

OPs are also possible in multiple variables, and Dunkl and Xu (2014) give a comprehensive overview. For a multi-index $\alpha = (\alpha_1, \dots, \alpha_d) \in \mathbb{N}_0^d$, where $\mathbb{N}_0 = \{0, 1, 2, \dots\}$, and $\mathbf{x} = (x_1, \dots, x_d)$, a monomial term is given by $\mathbf{x}^\alpha =$

² Unfortunately, Olver, Lozier, Boisvert and Clark (2010) have interchanged the notation for V_n and W_n . Therefore, the convention chosen in the Digital Library of Mathematical Functions does not match the original definitions (Gautschi 1992). This is one of the only times we deviate from the notation found in Olver *et al.* (2010).

$x_1^{\alpha_1} \cdots x_d^{\alpha_d}$. The number $|\alpha| = \alpha_1 + \cdots + \alpha_d$ is called the (total) degree of \mathbf{x}^α . The space of polynomials of (total) degree at most n is denoted by

$$\Pi_n^d = \text{Span}\{\mathbf{x}^\alpha : |\alpha| \leq n, \alpha \in \mathbb{N}_0^d\}.$$

A polynomial $p \in \Pi_n^d$ is said to be an *orthogonal polynomial* of degree n with respect to a weight function $w(\mathbf{x})$ if

$$\int \cdots \int p(\mathbf{x}) w(\mathbf{x}) q(\mathbf{x}) d\mathbf{x} = 0$$

for any $q \in \Pi_{n-1}^d$. This means that p is orthogonal to all polynomials of lower degree, though it does not have to be orthogonal to polynomials of the same degree. We will only use *mutually orthogonal polynomials* which are also orthogonal with polynomials of the same degree.

We find it convenient to express a family of multivariable OPs as a block-quasimatrix, where each block contains the polynomials of the same total degree. In two variables we write

$$\mathbf{P}(\mathbf{x}) = (p_{00}(\mathbf{x}) | p_{10}(\mathbf{x}) | p_{11}(\mathbf{x}) | p_{20}(\mathbf{x}) | p_{21}(\mathbf{x}) | p_{22}(\mathbf{x}) | \cdots),$$

while in three variables we have

$$\mathbf{P}(\mathbf{x}) = (p_{00}(\mathbf{x}) | p_{10}(\mathbf{x}) | p_{11}(\mathbf{x}) | p_{12}(\mathbf{x}) | \cdots).$$

A sequence of multivariable OPs is said to be orthonormal with respect to w if all the polynomials of a given degree are orthogonal and of unit length, i.e. $\mathbf{P}^\top \underline{w} \mathbf{P} = I$, where I is the identity operator. OPs in multiple variable are only unique up to an orthogonal recombination. In two variables, given any sequence of $n \times n$ orthogonal matrices Q_n , we find that if \mathbf{P} are orthonormal then

$$\tilde{\mathbf{P}} := \mathbf{P} \begin{pmatrix} Q_0 & & \\ & Q_1 & \\ & & \ddots \end{pmatrix}$$

is also orthonormal with respect to the same weight, since

$$\tilde{\mathbf{P}}^\top \underline{w} \tilde{\mathbf{P}} = \begin{pmatrix} Q_0^\top & & \\ & Q_1^\top & \\ & & \ddots \end{pmatrix} \mathbf{P}^\top \underline{w} \mathbf{P} \begin{pmatrix} Q_0 & & \\ & Q_1 & \\ & & \ddots \end{pmatrix} = I.$$

Therefore, the choice of basis in multivariable OPs is always somewhat arbitrary, though different choices lead to different sparsity properties in recurrence relationships.

Multivariable OPs satisfy a three-term recurrence relationship in each variable and hence have an associated (non-symmetric) Jacobi operator for each variable.

Definition 2.3. The *Jacobi operators* J_1, \dots, J_d associated with a family of multivariate OPs are block-tridiagonal operators such that

$$\mathbf{P}^\top(\mathbf{x})\underline{x}_k = \underbrace{\begin{pmatrix} A_0^k & B_0^k & & \\ C_1^k & A_1^k & B_1^k & \\ & C_2^k & A_2^k & \ddots \\ & & \ddots & \ddots \end{pmatrix}}_{J_k} \mathbf{P}^\top(\mathbf{x}), \quad 1 \leq k \leq d.$$

The sizes of A_n^k , B_n^k and C_n^k are determined by the number of ways d non-negative integers can be chosen to sum to k . When $d = 2$, we have $A_n^k \in \mathbb{R}^{(n+1) \times (n+1)}$, $B_n^k \in \mathbb{R}^{(n+1) \times (n+2)}$ and $C_n^k \in \mathbb{R}^{(n+1) \times n}$.

Again, the transpose of J_k for $1 \leq k \leq d$, denoted by $X_k := J_k^\top$, is the discrete operator that represents multiplication by x_k in the OP basis. That is, if $f(\mathbf{x}) = \mathbf{P}(\mathbf{x})\mathbf{f}$, then the expansion coefficients for $x_k f(\mathbf{x})$ are $X_k \mathbf{f}$, since

$$x_k f(\mathbf{x}) = x_k \mathbf{P}(\mathbf{x})\mathbf{f} = \mathbf{P}(\mathbf{x})X_k \mathbf{f}.$$

If the family of OPs are orthonormal, then $J_k \equiv X_k$ is symmetric, that is, $(C_{n+1}^k)^\top = B_n^k$ and $(A_n^k)^\top = A_n^k$ for $1 \leq k \leq d$ and $n \geq 0$.

2.6. Constructing multivariate orthogonal polynomials

Classical multivariate OPs are defined on simple geometries such as rectangles, disks, triangles, hypercubes, balls and tetrahedra. There are, at least, two advantages to constructing multivariate OPs from univariate ones: (1) they typically have simple expressions in terms of univariate OPs, and (2) recurrence relationships can be derived from known univariate relationships. Two useful ways to construct multivariate OPs are tensor products and Koornwinder's construction (Koornwinder 1975).

2.6.1. Tensor product construction

OPs on tensor product domains (*e.g.* squares and hypercubes) can be constructed using the tensor product of univariate OPs. If $\phi_0(x), \phi_1(x), \dots$ and $\psi_0(x), \psi_1(x), \dots$ are two sequences of OPs with respect to the weight w_1 on $[a, b]$ and w_2 on $[c, d]$, then $\phi_j(x)\psi_k(y)$ for $j, k \geq 0$ are orthogonal with weight $w_1(x)w_2(y)$ on $[a, b] \times [c, d]$. Let $\Phi(x) = (\phi_0(x) | \phi_1(x) | \dots)$ and $\Psi(x) = (\psi_0(x) | \psi_1(x) | \dots)$. Then in terms of quasimatrices we can define the tensor product $\Phi \otimes \Psi$ as a generalization of a quasimatrix defined on \mathbb{R}^2 :

$$(\Phi \otimes \Psi)_{j,k}(x, y) = \phi_j(x)\psi_k(y), \quad j, k \geq 0.$$

One could then express coefficients as a matrix with multiplication defined tensor-wise. That is, if $F \in \mathbb{C}^\infty \otimes \mathbb{C}^\infty$ we would (formally) have

$$f(x, y) = (\Phi \otimes \Psi)(x, y)F,$$

thus we can view $\Phi \otimes \Psi: \ell_F^2 \rightarrow \mathbb{C}[x, y]$, where ℓ_F^2 denotes the space of matrices with finite number of non-zero entries.

It is often convenient to identify $\mathbb{C}^\infty \otimes \mathbb{C}^\infty$ with infinite matrices $\mathbb{C}^{\infty \times \infty}$, in which case we write

$$f(x, y) = \sum_{j, k=0}^{\infty} F_{j, k} \phi_j(x) \psi_k(y), = \Phi(x)F\Psi(y)^\top.$$

One still needs to determine how the entries of F are stored sequentially in memory. A natural approach is to replace the infinite matrix F with an $n \times m$ truncation and order the entries column by column (*column major*) or row by row (*row major*), with column major being the modern standard. We can interpret this as a block-quasimatrix, with blocks of size n , and denote it by $\text{vec}(\cdot)$ as follows:

$$\begin{aligned} & \text{vec}(\Phi_n \otimes \Psi_m)(x, y) \\ &:= (\phi_0(x)\psi_0(y), \dots, \phi_n(x)\psi_0(y) \mid \dots \mid \phi_0(x)\psi_m(y), \dots, \phi_n(x)\psi_m(y)). \end{aligned}$$

This point of view is computationally convenient for OPs on tensor product domains when working with truncated expansions, and naturally leads to the discretization of certain partial differential equations as generalized Sylvester equations (see Section 6.2).

An alternative viewpoint is to traverse the basis in an anti-diagonal way:

$$\begin{aligned} \mathbf{P}(x, y) &= \text{diagtrav}(\Phi(x) \otimes \Psi(y)) \\ &= (\phi_0(x)\psi_0(y) \mid \phi_1(x)\psi_0(y) \mid \phi_0(x)\psi_1(y) \mid \dots). \end{aligned}$$

That is, $\text{diagtrav}(\cdot)$ traverses a square infinite matrix and writes it out in anti-diagonal order. This has the benefit of grouping coefficients in blocks of polynomial degree, which fits the definition of multivariate OPs and extends naturally to non-tensor domains.

Remark 2.4. Another basis ordering is the hyperbolic cross, which comes up in high-dimensional quadrature to circumvent the curse of dimensionality. We focus on low-dimensional settings, where these are less useful.

2.6.2. Koornwinder's construction

Another construction of multivariate OPs is given by Koornwinder (1975). Let w_1 and w_2 be weight functions on $[a, b]$ and $[c, d]$, respectively, and let

$\rho(x)$ be a function on (a,b) such that either

- (1) ρ is a polynomial of degree 1, or
- (2) ρ^2 is non-negative polynomial of degree ≤ 2 with $c = -d > 0$ and w_2 is an even function.

Then one can construct multivariate OPs with respect to the weight function $W(x,y) = w_1(x)w_2(y/\rho(x))$ on $\Omega = \{(x,y) : c\rho(x) \leq y \leq d\rho(x), a \leq x \leq b\}$. To form the multivariate OPs, first write down all the univariate OPs $p_0^{(k)}(x), p_1^{(k)}(x), \dots$, associated to $\rho^{2k+1}(x)w_1(x)$ on $[a,b]$ and the univariate OPs $q_0(x), q_1(x), \dots$, associated to $w_2(x)$ on $[c,d]$. It can be shown that the bivariate polynomials given by

$$P_{n,k}(x,y) = p_{n-k}^{(k)}(x)\rho(x)^k q_k(y/\rho(x)), \quad 0 \leq k \leq n,$$

for each $n \geq 0$ are OPs on Ω with respect to $W(x,y)$ (Dunkl and Xu 2014). Note that there are natural extensions to higher dimensions, as outlined in Dunkl and Xu (2014).

In the degenerate case, $\rho(x) = 1$, the construction derives tensor product OPs for tensor product domains. However, the non-degenerate cases include OPs on spheres, disks, triangles and more exotic domains.

'Legendre' polynomials on a right-angled triangle. Let

$$\Omega = \{(x,y) : 0 \leq y \leq x, 0 \leq x \leq 1\}$$

be a right-angled triangle. Using Koornwinder's construction with $\rho(x) = x$ and $w_1(x) = w_2(x) = 1$, we find that

$$P_{n,k}(x,y) = P_{n-k}^{(2k+1,0)}(2x-1)x^k P_k\left(2\frac{y}{x}-1\right), \quad 0 \leq k \leq n, n \geq 0,$$

are OPs on the right-angled triangle with $W(x,y) = 1$.

'Jacobi' polynomials on a right-angled triangle. Let

$$\Omega = \{(x,y) : 0 \leq y \leq x, 0 \leq x \leq 1\}$$

be a right-angled triangle. Letting $\rho(x) = x$, $w_1(x) = x^a(1-x)^b$ and $w_2(x) = x^c(1-x)^d$, with $a,b,c,d > -1$, we find that

$$P_{n,k}^{(a,b,c,d)}(x,y) = P_{n-k}^{(2k+a+1,b)}(2x-1)x^k P_k^{(c,d)}\left(2\frac{y}{x}-1\right), \quad 0 \leq k \leq n, n \geq 0,$$

are OPs on the right-angled triangle with respect to $W(x,y) = x^a(1-x)^b(y/x)^c(1-y/x)^d$. An affine translation of this construction and setting $d = 0$ gives the standard Jacobi polynomials on the unit simplex $0 \leq x \leq 1, 0 \leq y \leq 1-x$, which are orthogonal with respect to $x^a y^b (1-x-y)^c$:

$$P_{n,k}^{(a,b,c)}(x,y) = P_{n-k}^{(2k+b+c+1,a)}(2x-1)(1-x)^k P_k^{(b,c)}\left(2\frac{y}{1-x}-1\right).$$

‘Legendre’ polynomials on the unit disk. The unit disk can be described as

$$\Omega = \{(x, y) : -\rho(x) \leq y \leq \rho(x), -1 \leq x \leq 1\},$$

where $\rho(x) = \sqrt{1-x^2}$. Following Koornwinder’s construction with $w_1(x) = w_2(x) = 1$, we find that

$$Z_{n,k}(x, y) = C_{n-k}^{(k+1)}(x)(1-x^2)^{k/2}P_k\left(\frac{y}{\sqrt{1-x^2}}\right), \quad 0 \leq k \leq n, n \geq 0,$$

are OPs on the unit disk with $W(x, y) = 1$.

Alternatively, it is common to use polar coordinates on the disk so that

$$Z_{n,k}(r, \theta) = P_{(n-|k|)/2}^{(0, |k|)}(2r^2 - 1)r^{|k|}e^{ik\theta}, \quad -n \leq k \leq n, n \geq 0,$$

where $n - k$ is even. These polynomials are also referred to as the Zernike polynomials (Zernike 1934).

‘Ultraspherical’ polynomials on the unit disk. The unit disk can be described as

$$\Omega = \{(x, y) : -\rho(x) \leq y \leq \rho(x), -1 \leq x \leq 1\},$$

where $\rho(x) = \sqrt{1-x^2}$. Following Koornwinder’s construction with $w_1(x) = w_2(x) = (1-x^2)^{\lambda-1/2}$, we find that

$$Z_{n,k}^{(\lambda)}(x, y) = C_{n-k}^{(k+1)}(x)(1-x^2)^{k/2}C_k^{(\lambda)}\left(\frac{y}{\sqrt{1-x^2}}\right), \quad 0 \leq k \leq n, n \geq 0,$$

are OPs on the unit disk with $W(x, y) = (1-\|x\|^2)^\lambda$.

And on the unit sphere

$$\Omega = \{(x, y, z) \in \mathbb{R}^3 : x^2 + y^2 + z^2 = 1\},$$

parametrized by co-latitude $\theta \in [0, \pi]$ and longitude $\varphi \in [0, 2\pi)$, spherical harmonics (Atkinson and Han 2012, Dai and Xu 2013) share a similar structural property with Koornwinder’s constructions and the periodic representation of Zernike polynomials

$$Y_n^k(\theta, \varphi) = P_{n-|k|}^{(|k|, |k|)}(\cos \theta)(\sin \theta)^{|k|}e^{ik\varphi}, \quad -n \leq k \leq n, n \geq 0.$$

The same construction procedure can be used to derive multivariate OPs on domains with parabolic boundaries (Koornwinder 1975), quadratic curves (Olver and Xu 2020b), quadratic surfaces of revolution (Olver and Xu 2020a) and disk slices and trapeziums (Snowball and Olver 2020).

‘Lambda’ functions in \mathbb{R}^3 . The so-called Lambda functions (Hylleraas 1929, Shull and Löwdin 1955, Löwdin and Shull 1956) provide a complete orthogonal basis for $L^2(\mathbb{R}^3)$ that are defined in terms of generalized Laguerre polynomials and spherical harmonics

$$\Lambda_{k,\ell}^m(r, \theta, \varphi) = e^{-r/2}r^\ell L_{k-\ell}^{(2\ell+2)}(r)Y_\ell^m(\theta, \varphi).$$

Here, $k \geq 0$, $0 \leq \ell \leq k$, and $-m \leq \ell \leq m$. For a recent survey of bases for Sobolev spaces in \mathbb{R}^3 , we refer the interested reader to Weniger (2019).

2.7. Evaluation of orthogonal polynomials and Clenshaw's algorithm

The three-term recurrence relationship for OPs provides an immediate algorithm to evaluate OPs. In the univariate setting, assuming that $p_0(x) = 1$, the OPs satisfy the following lower-triangular system (see (2.5)):

$$\underbrace{\begin{pmatrix} 1 & & & \\ a_0 - x & b_0 & & \\ c_1 & a_1 - x & b_1 & \\ c_2 & a_2 - x & b_2 & \\ \ddots & \ddots & \ddots & \ddots \end{pmatrix}}_{L_x} \begin{pmatrix} p_0(x) \\ p_1(x) \\ p_2(x) \\ p_3(x) \\ \vdots \end{pmatrix} = \begin{pmatrix} 1 \\ 0 \\ 0 \\ 0 \\ \vdots \end{pmatrix}, \quad L_x := \begin{pmatrix} e_0^\top \\ J - xI \end{pmatrix}. \quad (2.7)$$

This is equivalent to $\mathbf{P}(x)^\top = L_x^{-1}e_0$, where e_0 is the first unit vector. The linear system in (2.7) can be solved using forward recurrence to evaluate $p_0(x), \dots, p_{N-1}(x)$ in $O(N)$ operations, which we observe to be computationally more efficient in practice than using explicit formulas such as $T_n(x) = \cos(n \cos^{-1} x)$ for $-1 \leq x \leq 1$. One can also use (2.7) to efficiently evaluate functions, that is,

$$f(x) = \mathbf{P}(x)\mathbf{f} = e_0^\top L_x^{-\top} \mathbf{f}.$$

This procedure is known as *Clenshaw's algorithm* (Clenshaw 1955). Therefore, to evaluate $f(x)$ one can perform back-substitution using the linear system $L_x^\top \mathbf{v} = \mathbf{f}$ and then taking the leading entry. This requires $O(N)$ operations and no extra memory allocations if \mathbf{f} is such that $\mathbf{f}_k = 0$ for $k \geq N$.

Clenshaw's algorithm works in multiple dimensions as well, though it is significantly more involved as one needs to work simultaneously with all recurrence relationships, *i.e.* $J_k \mathbf{P}(x)^\top = \mathbf{P}(x)^\top \underline{x}_k$ for $1 \leq k \leq d$. In two variables, the analogue of (2.7) is the following block system:

$$\underbrace{\begin{pmatrix} 1 & & & \\ A_0^x - xI_1 & B_0^x & & \\ A_0^y - yI_1 & B_0^y & & \\ C_0^x & A_1^x - xI_2 & B_1^x & \\ C_0^y & A_1^y - yI_2 & B_1^y & \\ \ddots & \ddots & \ddots & \ddots \end{pmatrix}}_{L_{x,y}} \mathbf{P}(x,y)^\top = \begin{pmatrix} 1 \\ \mathbf{0}_{1 \times 1} \\ \mathbf{0}_{1 \times 1} \\ \mathbf{0}_{2 \times 1} \\ \mathbf{0}_{2 \times 1} \\ \vdots \end{pmatrix}, \quad (2.8)$$

which is built by interlacing the blocks of

$$J_x \mathbf{P}(x, y)^\top = \mathbf{P}(x, y)^\top \underline{x} \quad \text{and} \quad J_y \mathbf{P}(y)^\top = \mathbf{P}(x)^\top \underline{y}.$$

The issue here is that $L_{x,y}$ is not a lower-triangular matrix. Fortunately, the blocks given by

$$B_n = \begin{pmatrix} B_n^x \\ B_n^y \end{pmatrix} \in \mathbb{R}^{(2n+2) \times (n+2)}$$

have full column rank (Dunkl and Xu 2014, Theorem 3.3.4) and hence B_n has a left-inverse B_n^+ so that $B_n^+ B_n = I_{n+2}$. Thus we can modify (2.8) to the following lower-triangular linear system:

$$\tilde{L}_{x,y} \mathbf{P}(x, y)^\top = \begin{pmatrix} 1 \\ \mathbf{0}_{1 \times 1} \\ \mathbf{0}_{2 \times 1} \\ \vdots \end{pmatrix}, \quad \tilde{L}_{x,y} = \begin{pmatrix} 1 & & & \\ & B_0^+ & & \\ & & B_1^+ & \\ & & & \ddots \end{pmatrix} L_{x,y}.$$

This can now be solved by forward substitution. A multivariate analogue of Clenshaw's algorithm follows immediately:

$$f(x, y) = \mathbf{P}(x, y) \mathbf{f} = \mathbf{e}_0^\top \tilde{L}_{x,y}^{-\top} \mathbf{f}.$$

The choice of left-inverse B_n^+ is a delicate issue as it can impact the complexity and stability of multivariable Clenshaw's algorithm. For the multivariate OPs we often have that B_n in (2.8) is a sparse matrix. However, if one naïvely uses the Moore–Penrose pseudo-inverse, then B_n^+ is typically dense, which leads to a disappointing complexity for evaluating OP expansions. In specific cases, it can be possible to build a sparse B_n^+ ; see Olver, Townsend and Vasil (2019) for an example on the triangle.

2.8. Recurrences and differential operators

Classical orthogonal polynomials satisfy numerous recurrences – going beyond the three-term recurrence relationship – for differentiation, multiplication and other operations; see Olver *et al.* (2010, Chapter 18) for a comprehensive summary. The relevant ones for this paper are reviewed in Appendix A and are expressed in the quasimatrix notation.

As a brief example, families of classical orthogonal polynomials have the beautiful property that their derivatives are also classical orthogonal polynomials. For example, we have the following formula (Olver *et al.* 2010, 18.9.19):

$$\frac{dC_n^{(\lambda)}}{dx} = 2\lambda C_{n-1}^{(\lambda+1)},$$

which can be expressed as (see Proposition A.11)

$$\mathcal{D}\mathbf{C}^{(\lambda)} = \mathbf{C}^{(\lambda+1)} \underbrace{\begin{pmatrix} 0 & 2\lambda & & \\ & 2\lambda & & \\ & & 2\lambda & \\ & & & \ddots \end{pmatrix}}_{D_{(\lambda)}^{(\lambda+1)}}.$$

In other words, if $f(x) = \mathbf{C}^{(\lambda)}\mathbf{f}$, then

$$f'(x) = \mathcal{D}\mathbf{C}^{(\lambda)}\mathbf{f} = \mathbf{C}^{(\lambda+1)}D_{(\lambda)}^{(\lambda+1)}\mathbf{f}.$$

That is, we can efficiently compute the coefficients of the derivative of f in another classical OP expansion via the product $D_{(\lambda)}^{(\lambda+1)}\mathbf{f}$.

Another example of an operator is the change from one basis to another (see Proposition A.9):

$$\mathbf{C}^{(\lambda)} = \mathbf{C}^{(\lambda+1)} \underbrace{\begin{pmatrix} 1 & 0 & -\frac{\lambda}{2+\lambda} & & \\ & \frac{\lambda}{1+\lambda} & 0 & -\frac{\lambda}{3+\lambda} & \\ & & \frac{\lambda}{2+\lambda} & 0 & -\frac{\lambda}{5+\lambda} \\ & & & \ddots & \ddots \\ & & & & \ddots \end{pmatrix}}_{R_{(\lambda)}^{(\lambda+1)}}.$$

That is, we have

$$f(x) = \mathbf{C}^{(\lambda)}(x)\mathbf{f} = \mathbf{C}^{(\lambda+1)}(x)R_{(\lambda)}^{(\lambda+1)}\mathbf{f}.$$

This is a simple example of a transform, which we discuss further in Section 4.4. Note that these relationships can be composed to deduce more general operators, for example,

$$\begin{aligned} f''(x) + xf(x) &= \mathcal{D}^2\mathbf{C}^{(\lambda)}(x)\mathbf{f} + x\mathbf{C}^{(\lambda)}(x)\mathbf{f} \\ &= \mathcal{D}\mathbf{C}^{(\lambda+1)}(x)D_{(\lambda)}^{(\lambda+1)}\mathbf{f} + \mathbf{C}^{(\lambda)}(x)X_{(\lambda)}\mathbf{f} \\ &= \mathbf{C}^{(\lambda+2)}(\underbrace{D_{(\lambda)}^{(\lambda+2)}D_{(\lambda)}^{(\lambda+1)}}_{D_{(\lambda)}^{(\lambda+2)}} + \underbrace{R_{(\lambda+1)}^{(\lambda+2)}R_{(\lambda)}^{(\lambda+1)}}_{R_{(\lambda)}^{(\lambda+2)}}X_{(\lambda)})\mathbf{f}, \end{aligned}$$

where $X_{(\lambda)}$ is the multiplication-by- x operator for $\mathbf{C}^{(\lambda)}$. Such manipulations lead to fast numerical methods for differential equations, as discussed in Section 5.

3. Quadrature

A quadrature rule approximates an integral by a weighted sum of samples of f . The classic interpolatory quadrature rules are Newton–Cotes, Gaussian quadrature and Clenshaw–Curtis, which all take the form

$$\int_a^b w(x)f(x)dx \approx \sum_{k=1}^N w_k f(x_k), \quad (3.1)$$

where $w(x)$ is a non-negative integrable weight function and $\{x_k\}$ and $\{w_k\}$ are called the quadrature nodes and weights, respectively. It is known that an N -point Gaussian quadrature exactly integrates polynomials of degree $\leq 2N - 1$. When $w(x) = 1$, the corresponding Gauss quadrature rule is called Gauss–Legendre.

3.1. Weighted Clenshaw–Curtis quadrature

Clenshaw–Curtis quadrature sets the nodes $\{x_k\}$ in (3.1) to be Chebyshev points that are trivial to generate. There are at least three main variations on Chebyshev points:

Fejér's 1st rule	$x_k = \cos\left(\frac{k\pi}{N+1}\right)$,	$1 \leq k \leq N$,
Fejér's 2nd rule	$x_k = \cos\left(\frac{(k-1/2)\pi}{N}\right)$,	$1 \leq k \leq N$,
Cloesshaw–Curtis	$x_k = \cos\left(\frac{(k-1)\pi}{(N-1)}\right)$,	$1 \leq k \leq N$.

The differences between these choices for the quadrature nodes are minor, though the Cloesshaw–Curtis rule has the advantage that the N -point and $2N$ -point nodal sets are nested, allowing for re-use of function evaluation when N is doubled, which can have a significant impact on speed if calculating f is expensive. Therefore Cloesshaw–Curtis has become the most popular variant.

The weights $\{w_k\}$ are selected so that any polynomial p of degree $\leq N - 1$ is exactly integrated. That is,

$$\int_{-1}^1 w(x)p(x)dx = \sum_{k=1}^N w_k p(x_k).$$

In particular, the quadrature must exactly integrate the Chebyshev polynomials $T_0(x), \dots, T_{N-1}(x)$. We find that $\{w_k\}$ satisfy the following linear

system:

$$\mathbf{T}_N(\mathbf{x})^\top \mathbf{w} = \mathbf{b}, \quad \mathbf{T}_N(\mathbf{x}) = \begin{pmatrix} T_0(x_1) & \cdots & T_{N-1}(x_1) \\ \vdots & \ddots & \vdots \\ T_0(x_N) & \cdots & T_{N-1}(x_N) \end{pmatrix}, \quad (3.2)$$

$$\mathbf{b}_n = \int_{-1}^1 w(x) T_n(x) dx,$$

where $\mathbf{x} = (x_1, \dots, x_N)^\top$ and $\mathbf{w} = (w_1, \dots, w_N)^\top$.

The original papers by Fejér (1933) and Clenshaw and Curtis (1960) were written before the rediscovery of the fast Fourier transform (Cooley and Tukey 1965); however, soon afterwards Gentleman (1972) noted that the linear system $\mathbf{T}_N(\mathbf{x})^\top \mathbf{w} = \mathbf{b}$ could be solved using the FFT (see Section 4.1). Clenshaw–Curtis quadrature is popular because the rule is accurate, and the nodes and weights can be easily computed in $O(N \log N)$ operations. In the 1970s it was one of the only accurate rules for which large quadrature rules could be computed for large N . Today, with fast algorithms to compute many Gauss quadrature rules, we regard Gauss and Clenshaw–Curtis formulas as both valuable for numerical integration (Trefethen 2008).

Waldvogel (2003) showed a factor of 2 computational saving in computing $\{w_k\}$ when $w(x) = 1$ by exploiting the fact that $b_k = 0$ for odd k in (3.2), which allows one to compute the weights using an FFT of half the size. Waldvogel’s speedup can be extended to symmetric weights of the form $w(x) = (1 - x^2)^\alpha$ for $\alpha > -1$ (Sommariva 2013).

3.2. The original derivation of Gauss–Legendre quadrature

In contrast to Clenshaw–Curtis quadrature, an N -point Gauss–Legendre quadrature rule exactly integrates polynomials of degree $\leq 2N - 1$ by carefully selecting *both* the nodes and weights. Most textbooks in numerical analysis present the topic in its modern formulation using orthogonal polynomials, as derived by Jacobi (1826). However, the original derivation of the Gauss–Legendre quadrature rule by Gauss (1815) is a beautiful calculation involving continued fractions (Gautschi 1981a).

Let p be any polynomial of degree $\leq 2N - 1$ and let Γ be the circle of radius $1 + \delta$ (for arbitrary small $\delta > 0$) centred at 0 and parametrized in an anticlockwise orientation. Then, from Cauchy’s integral formula we can write

$$\int_{-1}^1 p(x) dx = \int_{-1}^1 \frac{1}{2\pi i} \int_{\Gamma} \frac{p(z)}{z-x} dz dx = \frac{1}{2\pi i} \int_{\Gamma} p(z) \phi(z) dz, \quad (3.3)$$

where $\phi(z) = \log((z+1)/(z-1))$ is the analytic function in $\mathbb{C} \cup \{\infty\} \setminus [-1, 1]$ given by the standard branch of the logarithm, so that ϕ is positive for

$z \in (1, \infty)$. Since we want the quadrature rule to be exact, we need

$$\int_{-1}^1 p(x)dx = \sum_{k=1}^N w_k p(x_k) = \frac{1}{2\pi i} \int_{\Gamma} p(z) r_N(z) dz, \quad r_N(z) = \sum_{k=1}^N \frac{w_k}{z - x_k}, \quad (3.4)$$

where $r_N(z)$ is a $(N, N-1)$ rational function with poles at the quadrature nodes $\{x_k\}$. By combining (3.3) and (3.4) and noting that p could be any polynomial, we require that

$$\frac{1}{2\pi i} \int_{\Gamma} z^j (\phi(z) - r_N(z)) dz = 0, \quad 0 \leq j \leq 2N-1.$$

Due to the orthogonality of the monomials on the unit circle, we need $\phi(z) - r_N(z) = O(z^{2N})$. Gauss would have immediately known how to solve this using continued fractions, which were a popular topic in the early 1800s. In particular, $\phi(z)$ can be expressed as the following continued fraction:

$$\phi(z) = \log\left(\frac{z+1}{z-1}\right) = \cfrac{a_1}{z + \cfrac{a_2}{z + \cfrac{a_3}{z + \cfrac{a_4}{z + \dots}}}}, \quad a_k = -\frac{k^2}{4k^2 - 1}.$$

From this expression one can find the N -convergent, which is an $(N, N-1)$ rational function $r_N(z) = A_N(z)/B_N(z)$ such that $\phi(z) - r_N(z) = O(z^{2N})$. Moreover, it is known that

$$B_{k+1}(z) = zB_k(z) - \frac{k^2}{4k^2 - 1} B_{k-1}(z), \quad k \geq 0, \quad B_0(z) = 1, \quad B_{-1}(z) = 0. \quad (3.5)$$

Gauss used this three-term recurrence to derive $B_N(z)$ and then found its zeros to calculate the quadrature nodes. For example, $B_3(z) = z^3 - (3/5)z$, so the three-point Gauss quadrature rule has nodes $\{-\sqrt{3/5}, 0, \sqrt{3/5}\}$. One can then find the corresponding Gauss–Legendre weights by solving the linear system

$$(x^0, x^1, \dots, x^{N-1})^\top \mathbf{w} = \mathbf{b}, \quad b_k = \int_{-1}^1 x^k dx = \frac{1 - (-1)^{k+1}}{k+1}.$$

For the three-point quadrature, we find that the weights are $\{5/9, 8/9, 5/9\}$. Today, tables of Gauss–Legendre quadrature nodes and weights can be found throughout the internet and textbooks, though it is more practical to calculate them numerically.

3.3. The Golub–Welsch algorithm

Only a few years after Gauss's derivation, Jacobi (1826) realized the connection between Gauss quadrature and OPs. Jacobi first observed that the recurrence relationship in (3.5) is satisfied by the Legendre polynomials

when scaled to be monic. That means that the N -point Gauss–Legendre quadrature nodes satisfy $P_N(x_k) = 0$ for $1 \leq k \leq N$. At the time, Jacobi’s observation did not help one compute Gauss quadrature nodes and weights; however, it did allow Jacobi to generalize Gauss quadrature to other integrable weight functions. In particular, the Gauss quadrature nodes associated to (3.1) are given by the zeros of $p_N(x)$, where p_0, p_1, \dots , is a family of orthogonal polynomials associated to the weight function $w(x)$. In other words, the Gauss quadrature nodes are the zeros of an OP.

The most well-known algorithm for computing Gauss quadrature rules is the Golub–Welsch algorithm (Golub and Welsch 1969). It is based on the fact that

$$J_{0:N-1,0:N} \begin{pmatrix} p_0(\lambda) \\ \vdots \\ p_N(\lambda) \end{pmatrix} = \lambda \begin{pmatrix} p_0(\lambda) \\ \vdots \\ p_{N-1}(\lambda) \end{pmatrix}$$

and hence, by rearranging, we find

$$(J_N - \lambda I) \begin{pmatrix} p_0(\lambda) \\ \vdots \\ p_{N-2}(\lambda) \\ p_{N-1}(\lambda) \end{pmatrix} = \begin{pmatrix} 0 \\ \vdots \\ 0 \\ -b_{N-1}p_N(\lambda) \end{pmatrix}. \quad (3.6)$$

Here, J_N is the $N \times N$ truncation of the Jacobi operator associated with the OPs. The right-hand side of (3.6) is the zero vector if and only if λ is a root $p_N(\lambda)$. This means that the quadrature nodes can be computed in $O(N^3)$ operations as an eigenvalue problem involving a tridiagonal matrix. Moreover, if the OPs are orthonormal (which does not change the zeros), then J_N is a symmetric tridiagonal matrix reducing the cost to only $O(N^2)$ operations.

From (3.6), the eigenvector of J_N corresponding to eigenvalue x_j is

$$\begin{pmatrix} p_0(x_j) \\ \vdots \\ p_{N-1}(x_j) \end{pmatrix}.$$

This means that the quadrature weights can be obtained from the eigenvectors of J_N . Since an N -point Gauss quadrature is exact for polynomials up to degree $\leq 2N - 1$, we know that the quadrature weights $\mathbf{w} = (w_1, \dots, w_N)^\top$ satisfy

$$\mathbf{S}_N(\mathbf{x}) D_{\mathbf{w}} \mathbf{S}_N^\top(\mathbf{x}) = D_{\mathbf{b}}, \quad b_k = \int_{-1}^1 w(x)p_k(x)dx,$$

where

$$(\mathbf{S}_N(\mathbf{x}))_{jk} = p_{k-1}(x_j), \quad 1 \leq j, k \leq N.$$

Therefore $D_{\mathbf{w}}^{-1} = \mathbf{S}_N^\top(\mathbf{x})\mathbf{S}_N(\mathbf{x})$, that is,

$$\frac{1}{w_j} = \sum_{i=0}^{N-1} p_i(x_j)^2, \quad 1 \leq j \leq N.$$

In other words, the reciprocal of w_j is equal to the sum of the squares of the entries of the eigenvector of J_N corresponding to eigenvalue x_j , where the first entry of the eigenvector is scaled to have the value of 1.

3.4. Computing Gauss quadrature rules associated with classical orthogonal polynomials

The most important examples of Gauss quadrature rules are those associated with classical OPs. By now, there are hundreds of publications on the subject of computing classical Gauss quadrature rules, and the most efficient algorithms are tailor-made for each quadrature rule. In Table 3.1 we have surveyed the literature on computing Gauss–Legendre quadrature to observe the continuing progression to larger and larger quadrature rules (Townsend 2015). Also, see Figure 3.1. The surprise is that the Golub–Welsch algorithm is not, and never was, the fastest or most accurate algorithm for computing Gauss–Legendre quadrature rules. However, its generality and ease of implementation makes it one of the most widely used algorithms for computing Gauss–Legendre quadrature to this day.

There are now spectacularly fast algorithms for computing Gauss quadrature rules associated with classical OPs. The fastest methods are based on high-order asymptotic formulas for the quadrature nodes (the zeros of the classical OPs) and weights. For example, Bogaert derived asymptotic formulas for the n -point Gauss–Legendre nodes, which are accurate to about 16 digits of precision for any $N \geq 20$ (Bogaert 2014). Let $x_{N,k} = \cos(\theta_{N,k})$ be the k th node in the N -point Gauss quadrature rule. Then

$$\theta_{N,k} = \alpha_{N,k} + \sum_{m=1}^5 F_m \left(\alpha_{N,k}, \frac{\cos(\alpha_{N,k})}{\sin(\alpha_{N,k})} \right) v_N^{2m} + O(v_N^{12}), \quad v_N = \frac{1}{N+1/2}, \quad (3.7)$$

where $\alpha_{N,k} = v_N j_{0,k}$ and $j_{0,k}$ is the k th positive zero of the Bessel function $J_0(x)$. Here, the functions F_1, \dots, F_5 are known explicitly. For example,

$$F_1(x, u) = \frac{1}{8} \frac{ux - 1}{x}, \quad F_2(x, u) = \frac{1}{384} \frac{6x^2(1+u^2) + 25 - u(31u^2 + 33)x^3}{x^3},$$

where the remaining formulas for F_3 , F_4 and F_5 along with an asymptotic formula for the quadrature weights can be found in Bogaert (2014). For all practical purposes, (3.7) are essentially explicit formulas for the nodes and weights as they can be used for any $N \geq 20$ (in double precision) and the cost to evaluate them involves basic arithmetic and trigonometric evaluations.

Table 3.1. A brief survey of computing Gauss–Legendre quadrature. Originally, small Gauss–Legendre quadrature rules were calculated by brute-force hand calculations; today, a billion quadrature nodes and weights can be computed in a fraction of second on a standard laptop. The value N is the largest reported Gauss–Legendre quadrature rule in the publication.

N	Algorithm	Reference
7	by hand	Gauss (1815)
8	by hand	Tallqvist (1905)
10	by hand	Moors (1905)
7	by hand	Nyström (1930)
12	by hand	Bayly (1938)
16	by hand	Lowan, Davids and Levenson (1942)
16	Newton–Raphson	Davis and Rabinowitz (1956)
64	Newton–Raphson	Gawlik (1958)
96	Newton–Raphson	Davis and Rabinowitz (1958)
10	QD algorithm	Rutishauser (1962)
200	Newton–Raphson	Love (1966)
512	Newton–Raphson	Stroud and Secrest (1966)
50	Golub–Welsch	Golub and Welsch (1969)
256	Newton–Raphson	Lether (1978)
48	asymptotics for nodes	Gatteschi (1979)
640	Golub–Welsch	Gautschi (1982)
32	asymptotics	Förster and Petras (1990)
32	asymptotics for nodes	Förster and Petras (1993)
8000	higher-order Newton–Raphson	Yakimiw (1996)
8000	optimized Golub–Welsch	Yakimiw (1996)
200	Newton–Raphson	Petras (1999)
2048	Golub–Welsch	Swarztrauber (2003)
16384	variant of Newton–Raphson	Swarztrauber (2003)
12288	Newton–Raphson	Bailey, Jeyabalan and Li (2004)
10^6	GLR algorithm	Glaser, Liu and Rokhlin (2007)
10^8	Newton–Raphson with asymptotics	Bogaert, Michiels and Fostier (2012)
10^6	Newton–Raphson with asymptotics	Hale and Townsend (2013)
10^8	asymptotics formulas for the nodes	Bogaert (2014)

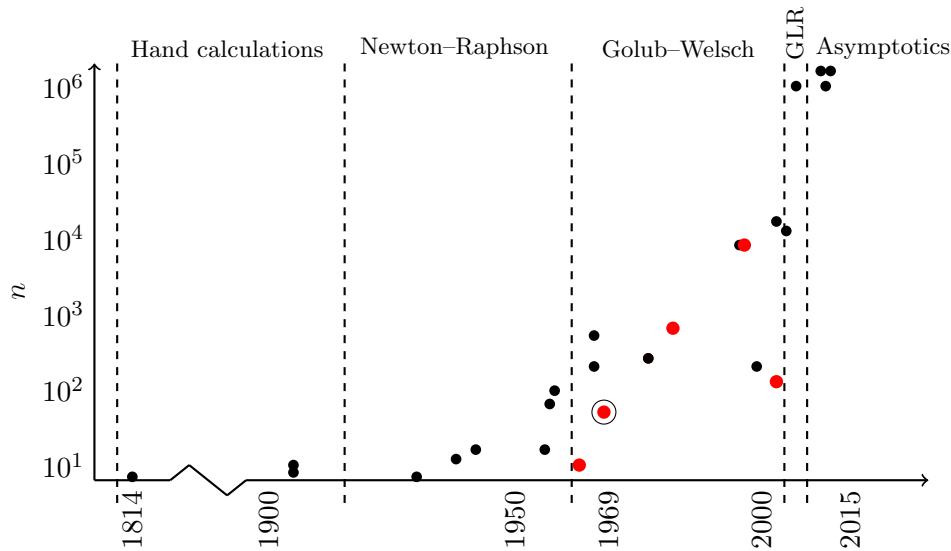


Figure 3.1. The history of computing high-order Gauss–Legendre quadrature. A dot represents published work, located at the publication year and the largest Gauss–Legendre rule reported therein. A large red dot is a paper based on variants of the Golub–Welsch algorithm, and the circled red dot is the original paper by Golub and Welsch. GLR represents the Glaser, Liu and Rokhlin (2007) algorithm. Figure reproduced from Townsend (2015) with permission. Copyright © 2015 Society for Industrial and Applied Mathematics.

If one wants to compute the nodes to more than 16 digits of precision, then one recommends using $\theta_{N,k}$ in (3.7) as an initial guess for Newton’s method on the rootfinding problem $P_N(\cos\theta) = 0$ (Bogaert *et al.* 2012, Hale and Townsend 2013).

Other asymptotic formulas have been derived for Gauss–Jacobi quadrature (Gil, Segura and Temme 2019*b*) as well as Gauss–Laguerre (Gil, Segura and Temme 2019*a*) and Gauss–Hermite (Gil, Segura and Temme 2018). Again, for high-precision calculations, one can take these asymptotic formulas as initial guesses in Newton’s method (Hale and Townsend 2013, Townsend, Trogdon and Olver 2016).

3.5. Computing non-classical Gauss quadrature rules

The previous section shows that having control over the asymptotic behaviour of the zeros of classical OPs can lead to fast iterative-free algorithms for computing Gauss quadrature rules associated with classical weights. Going beyond classical weights, one can build on the recent advances in deriving the asymptotics of general OPs using the Riemann–Hilbert approach (Deift 1999) to determine the asymptotic behaviour of the zeros for

general weights on the real line. One can alternatively use the numerical solution of the underlying Riemann–Hilbert problem (Olver 2012) to calculate numerical asymptotic approximations to OPs and thereby calculate associated Gauss quadrature nodes for relatively general weights (Townsend *et al.* 2016).

4. Orthogonal polynomial transforms

Orthogonal polynomial transforms involve computing expansion coefficients in one basis either from function samples or from coefficients in another basis. There are two main types.

- (1) *Synthesis* and *analysis*. The infinite-dimensional synthesis operator is the map that takes OP expansion coefficients to the corresponding function by summing up the expansion. The natural finite-dimensional analogue of the synthesis operator is the map from expansion coefficients to values of the expansion at a specified grid. The infinite-dimensional analysis operator takes a function to its expansion coefficients, and the finite-dimensional counterpart takes function samples to expansion coefficients.
- (2) *Connection problems*. It is possible to represent a polynomial in different OP bases. The connection problem is the change-of-basis transform from one OP basis to another.

Our general strategy for synthesis, that is, sampling a polynomial represented in a given OP expansion, is to first represent the polynomial in an OP expansion in which a fast evaluation scheme can be derived based on the FFT. Therefore, we tackle synthesis in two steps: (1) convert one OP basis into another, and (2) evaluate the polynomial in its new OP basis using the FFT. Our strategy for analysis is, therefore, to reverse the process: (1) compute the coefficients in a new OP basis using the FFT, and (2) convert from the new basis into the desired basis.

4.1. Synthesis and analysis based on the fast Fourier transform

The synthesis and analysis transforms that are closely connected to FFTs are the ones related to the four types of Chebyshev expansions and the associated four types of Gauss quadrature nodes (Mason 1993). That is, consider the four kinds of Chebyshev nodes,

$$\begin{aligned} x_k^I &= -\cos\left(\frac{k-1/2}{N}\pi\right), & x_k^{II} &= -\cos\left(\frac{k}{N+1}\pi\right), & 1 \leq k \leq N, \\ x_k^{III} &= \cos\left(\frac{N-k+1/2}{N+1/2}\pi\right), & x_k^{IV} &= \cos\left(\frac{k+1/2}{N+1/2}\pi\right), & 1 \leq k \leq N, \end{aligned} \tag{4.1}$$

Table 4.1. Connection between the discrete cosine and sine transforms of various types and the evaluation of the four types of Chebyshev series at the four kinds of Chebyshev nodes. Equivalence of these transforms holds up to a diagonal scaling.

Transform	OP	Nodes	Inverse	Transform	OP	Nodes	Inverse
DCT-I	T_n	2nd	DCT-I	DST-I	U_n	2nd	DST-I
DCT-II	V_n	2nd	DCT-III	DST-II	W_n	2nd	DST-III
DCT-III	T_n	1st	DCT-II	DST-III	U_n	1st	DST-II
DCT-IV	V_n	1st	DCT-IV	DST-IV	W_n	1st	DST-IV
DCT-V	T_n	4th	DCT-V	DST-V	U_n	4th	DST-V
DCT-VI	V_n	4th	DCT-VII	DST-VI	W_n	4th	DST-VII
DCT-VII	T_n	3rd	DCT-VI	DST-VII	U_n	3rd	DST-VI
DCT-VIII	V_n	3rd	DCT-VIII	DST-VIII	W_n	3rd	DST-VIII

which are the zeros of $T_N(x)$, $U_N(x)$, $V_N(x)$ and $W_N(x)$, respectively. There are 16 types of Chebyshev transforms, corresponding to the four choices of nodes in (4.1) and the four choices of Chebyshev polynomials (Püschel and Moura 2003).

For example, if $T_N(x)$ is selected and we are given coefficients $\mathbf{f} \in \mathbb{C}^N$ so that

$$f(x) = \mathbf{T}_N(x)\mathbf{f},$$

where $\mathbf{T}_N = (T_0, \dots, T_{N-1})$, then the four synthesis transforms are

$$\mathbf{T}_N(\mathbf{x}^{\text{I}})\mathbf{f}, \quad \mathbf{T}_N(\mathbf{x}^{\text{II}})\mathbf{f}, \quad \mathbf{T}_N(\mathbf{x}^{\text{III}})\mathbf{f}, \quad \mathbf{T}_N(\mathbf{x}^{\text{IV}})\mathbf{f},$$

or, equivalently,

$$\sum_{n=0}^{N-1} f_n T_n(x_k^{\text{I}}), \quad \sum_{n=0}^{N-1} f_n T_n(x_k^{\text{II}}), \quad \sum_{n=0}^{N-1} f_n T_n(x_k^{\text{III}}), \quad \sum_{n=0}^{N-1} f_n T_n(x_k^{\text{IV}}).$$

Since $T_n(x) = \cos(n \cos^{-1} x)$ for $x \in [-1, 1]$, these transformations can all be related to discrete cosine transforms (DCTs), which can be computed using an FFT. Table 4.1 enumerates all the different discrete sine and cosine transforms and their relations with the four kinds of Chebyshev polynomials and their respective Gauss–Chebyshev nodes. More precise details on the connections can be found in Püschel and Moura (2003). When combined with non-uniform fast Fourier transforms (Dutt and Rokhlin 1993, Potts and Steidl 2003, Greengard and Lee 2004, Ruiz-Antolín and Townsend 2018), fast algorithms for synthesis and analysis at sample locations that deviate from Gauss–Chebyshev points are also possible.

More generally, given the three-term recurrence relationships for OPs, there are fast algorithms for synthesis and analysis at the associated Gauss quadrature nodes (Tygert 2010b). However, the direct connection to FFTs is no longer present. Instead, our general strategy is to convert (using the connection problem) into one of the Chebyshev bases. The fast algorithms in Tygert (2010b) are only possible when the recurrence relationships for the OPs are either known explicitly or fast to compute. For several multivariate orthogonal polynomials and semiclassical orthogonal polynomials, the recurrence relationships can be costly to compute (Snowball and Olver 2020).

4.2. The connection problem

The connection problem concerns the question of how to change between different basis expansions, and in particular how to do so fast. We can state this in terms of *connection coefficients*.

Definition 4.1. Let $\{\phi_n(x)\}_{n=0}^\infty$ be a family of orthogonal functions with respect to a measure μ and let $\{\psi_n(x)\}_{n=0}^\infty$ be another family of orthogonal functions with respect to a measure $\hat{\mu}$. Assuming that $\{\psi_n(x)\}_{n=0}^\infty \subset L^2[\hat{\mu}]$, we have

$$\phi_n(x) = \sum_{\ell=0}^{\infty} c_{\ell,n} \psi_\ell(x), \quad c_{\ell,n} = \frac{\langle \psi_\ell, \phi_n \rangle_{d\mu}}{\langle \psi_\ell, \psi_\ell \rangle_{d\mu}}, \quad (4.2)$$

where $\{c_{\ell,n}\}$ are referred to as the connection coefficients.

In quasimatrix notation, (4.2) can be conveniently expressed as

$$\Phi = \Psi \underbrace{\begin{pmatrix} c_{0,0} & c_{0,1} & \cdots \\ c_{1,0} & c_{1,1} & \cdots \\ \vdots & \vdots & \ddots \end{pmatrix}}_{C_\Phi^\Psi},$$

where $\Phi := (\phi_0 | \phi_1 | \dots)$ and $\Psi := (\psi_0 | \psi_1 | \dots)$. This means that connection coefficients are the entries of a basis conversion matrix. In particular, given a polynomial expansion $u(x) = \Phi(x)\mathbf{u}$, we have that $u(x) = \Psi(x)C_\Phi^\Psi \mathbf{u} = \Psi(x)\mathbf{v}$, where $\mathbf{v} = C_\Phi^\Psi \mathbf{u}$ are the expansion coefficients in Ψ .

We call connection problems where Φ and Ψ are orthogonal with respect to the same measure *measure-preserving*. If Φ and Ψ are orthogonal with respect to different measures, then we call the connection problem *measure-perturbing*. There is an interesting connection between measure-perturbing connection problems and sparse recurrence relationships (Narayan and Hesthaven 2012). In the measure-perturbing case, we wish to apply $N \times N$

truncations of C_Φ^Ψ to vectors in $O(N \text{polylog}(N))$ operations.³ As we shall see later, the measure-preserving case requires the application of rectangular truncations of C_Φ^Ψ , though both dimensions are still asymptotically governed by N .

In the special case of connection problems between two complete degree-graded orthogonal polynomial bases, the connection coefficient matrix is upper-triangular. To emphasize this we use the notation

$$R_\Phi^\Psi := C_\Phi^\Psi.$$

If Φ and Ψ are orthonormal with respect to the same measure, we use the notation

$$Q_\Phi^\Psi := C_\Phi^\Psi$$

to emphasize that it is a matrix with orthonormal columns. Moreover, if $\deg(\phi_k) = \deg(\psi_k)$ for $k \geq 0$, then Q_Φ^Ψ is a diagonal matrix with ± 1 on the diagonal due to the uniqueness of univariate OPs. Therefore, we focus on the cases (1) $\deg(\phi_n) = \deg(\psi_n) + 1$, (2) $\deg(\phi_n) = \deg(\psi_n) + 2$ with a certain symmetry in the domain and weight function, and (3) $\deg(\phi_n) \gg \deg(\psi_n)$.

Univariate OP transforms are an essential building block in the multivariate setting. For tensor product multivariable OPs (see Section 2.6.1), the transforms are Kronecker products of univariate transforms. For OPs from non-degenerate versions of Koornwinder's construction (see Section 2.6.2), both the measure-preserving and measure-perturbing transforms are required. We illustrate this via spherical harmonics, due to their prevalence in applications.

4.3. Chebyshev polynomials

There are sparse relationships between the four types of Chebyshev polynomials, coming from trigonometric identities, which lead to sparse connection problems. For example, we have (Proposition A.13)

$$\mathbf{T} = \mathbf{U} \underbrace{\begin{pmatrix} 1 & 0 & -1/2 \\ & 1/2 & 0 & \ddots \\ & & 1/2 & \ddots \\ & & & \ddots \end{pmatrix}}_{R_T^U} = \mathbf{V} \underbrace{\begin{pmatrix} 1 & 1/2 & & \\ & 1/2 & 1/2 & \\ & & 1/2 & 1/2 \\ & & & \ddots & \ddots \end{pmatrix}}_{R_T^V}.$$

³ The term $\text{polylog}(N)$ refers to any polynomial (usually of low degree) in terms of $\log(N)$.

For a polynomial of degree $\leq N$, one can efficiently convert between Chebyshev bases in $O(N)$ operations by multiplying or inverting via back-substitution with these connection coefficient matrices. This is useful if one already has access to Chebyshev coefficients of a particular type, especially if the coefficients are known to a good relative accuracy, which is preserved under these conversions provided they decay reasonably quickly, or as an alternative to implementing all 16 types of DCTs and DSTs.

4.4. Ultraspherical polynomials

For any $\lambda > 0$, we have (Proposition A.9)

$$\mathbf{C}^{(\lambda)} \begin{pmatrix} \lambda & & & \\ & \lambda+1 & & \\ & & \lambda+2 & \\ & & & \ddots \end{pmatrix} = \mathbf{C}^{(\lambda+1)} \begin{pmatrix} \lambda & 0 & -\lambda & & \\ & \lambda & 0 & -\lambda & \\ & & \ddots & \ddots & \ddots \end{pmatrix},$$

which can also be written as

$$\mathbf{C}^{(\lambda)} = \mathbf{C}^{(\lambda+1)} R_{(\lambda)}^{(\lambda+1)}.$$

For a polynomial of degree $\leq N$, one can convert between $\mathbf{C}^{(\lambda)}$ and $\mathbf{C}^{(\lambda+m)}$ bases, where m is an integer, by applying or inverting m banded operators in $O(mN)$ operations. Moreover, one can also efficiently convert between Jacobi bases $\mathbf{P}^{(a,b)}$ and $\mathbf{P}^{(a+i,b+j)}$, where i and j are integers, using the bidiagonal conversion operators in Proposition A.3.

4.5. The Chebyshev–Legendre transform and its inverse

Many different algorithms accelerate the Chebyshev–Legendre transform and its inverse, that is, applying the conversions R_T^P and $R_P^T = (R_T^P)^{-1}$, where

$$\mathbf{T} = \mathbf{P} R_T^P, \quad \mathbf{P} = \mathbf{T} R_P^T.$$

These include an adaptation of the fast multipole method (Alpert and Rokhlin 1991), algorithms based on semi-separable matrices (Keiner 2008, Keiner 2009, Keiner 2011) and analysis of the Toeplitz–Hankel structure of the connection coefficients (Townsend, Webb and Olver 2018). There are also methods for direct synthesis and analysis using the asymptotics of Legendre polynomials (Orszag 1986, Mori, Suda and Sugihara 1999, Hale and Townsend 2014, Slevinsky 2018b) as well as ideas related to non-oscillatory phase functions (Bremer and Yang 2019).

The Chebyshev–Legendre transform has the beneficial feature that the conversion matrices are known explicitly, albeit they are not banded.

Lemma 4.2 (Alpert and Rokhlin 1991). The connection operators for Chebyshev–Legendre are given by

$$\pi R_P^T = \begin{pmatrix} \Lambda(0)^2 & 0 & \Lambda(1)^2 & 0 & \Lambda(2)^2 & \cdots \\ & 2\Lambda(0)\Lambda(1) & 0 & 2\Lambda(1)\Lambda(2) & 0 & \ddots \\ & & 2\Lambda(0)\Lambda(2) & 0 & 2\Lambda(1)\Lambda(3) & \ddots \\ & & & \ddots & \ddots & \ddots \end{pmatrix},$$

$$R_T^P = \begin{pmatrix} 1 & 0 & -\frac{\Lambda(0)\Lambda(1/2)}{6} & 0 & -\frac{\Lambda(1)\Lambda(3/2)}{10} & \cdots \\ & \frac{\sqrt{\pi}}{2\Lambda(1)} & 0 & -\frac{9\Lambda(0)\Lambda(3/2)}{20} & 0 & \cdots \\ & & \frac{\sqrt{\pi}}{2\Lambda(2)} & 0 & -\frac{5\Lambda(0)\Lambda(3/2)}{7} & \cdots \\ & & & \ddots & \ddots & \ddots \end{pmatrix},$$

where $\Lambda(k) = \Gamma(k + 1/2)/\Gamma(k + 1)$. More explicitly, for $0 < i < j$ and $i + j$ even, we have $(R_P^T)_{ij} = \mathcal{M}(i, j)$ and $(R_T^P)_{ij} = \mathcal{L}(i, j)$, where

$$\mathcal{M}(x, y) = \frac{2}{\pi} \Lambda\left(\frac{y-x}{2}\right) \Lambda\left(\frac{y+x}{2}\right),$$

$$\mathcal{L}(x, y) = \frac{-y(x+1/2)}{(y+x+1)(y-x)} \Lambda\left(\frac{y-x-2}{2}\right) \Lambda\left(\frac{y+x-1}{2}\right).$$

Proof. There are a number of ways to prove this, but we introduce a non-standard way based on expressing R_P^T as the solution to a Sylvester equation of the form

$$AV = BVA\Lambda, \quad (4.3)$$

where A and B are upper-triangular banded matrices and Λ is diagonal. This version of the proof will lead to generalization of the construction in a way that can be exploited for other families of OPs via the algorithm in Section 4.6.4. As the columns of V are discrete eigenvectors, the solution is unique up to diagonal column scaling. Thus, if we specify the diagonal entries of V , then we can recover the remaining entries by back-substitution.

The differential equation satisfied by the Legendre quasimatrix, \mathbf{P} , is

$$[(x^2 - 1)\mathcal{D}^2 + 2x\mathcal{D}]\mathbf{P} = \mathbf{P}\Lambda,$$

where Λ is a diagonal matrix satisfying $\Lambda_{n,n} = n(n+1)$. Substituting in $\mathbf{P} = \mathbf{T}R_P^T$, we find that the Sylvester equation is

$$AR_P^T = BR_P^T\Lambda.$$

Here, we used the relationship

$$[(x^2 - 1)\mathcal{D}^2 + 2x\mathcal{D}]\mathbf{T} = \mathbf{U} \underbrace{[-L_{(2)}^U D_T^{(2)} + 2X_U D_T^U]}_A, \quad \mathbf{T} = \mathbf{U} \underbrace{R_T^U}_B,$$

where D_T^U is given in Proposition A.15, $D_T^{(2)}$ is given in Proposition A.11, $L_{(2)}^U$ is given in Proposition A.10, and X_U is the multiplication-by- x operator for \mathbf{U} . Direct inspection shows that the proposed formula for R_P^T satisfies this (banded) Sylvester equation, and it is the unique solution due to the normalizations of the polynomial families. \square

The existing methods for computing the Chebyshev–Legendre transform build on having an explicit formula for the entries. Here we outline various approaches.

4.5.1. The low-rank approximation approach

Alpert and Rokhlin (1991) exploit the fact that the entries of R_T^P and R_P^T can be written in terms of the bivariate smooth functions \mathcal{M} and \mathcal{L} . Then the conversion matrices are hierarchically subdivided using bivariate Chebyshev approximations of \mathcal{M} and \mathcal{L} .

Definition 4.3 (Keiner 2009). A square $[x_0, x_0 + c] \times [y_0, y_0 + c] \subset \mathbb{R}^2$ with $c > 0$ is said to be *well-separated* if $y_0 - x_0 \geq 2c$.

Alpert and Rokhlin (1991) prove rigorous error bounds for bivariate Chebyshev approximations to \mathcal{M} and \mathcal{L} on well-separated squares. The error bounds are derived by keeping one of the variables of \mathcal{M} and \mathcal{L} fixed and letting the other variable freely range over the possible values it may take in S . For example, keeping y fixed, error bounds are derived for Chebyshev interpolants to $\mathcal{M}(x(t), y)$ on well-separated squares, where

$$x(t) = x_0 + c(t+1)/2, \quad t \in [-1, 1].$$

Theorem 4.4 (Keiner 2009). Let \mathcal{M}_k denote the Chebyshev interpolant at k points to $\mathcal{M}(x(t), y)$. Then, for $c > 1$, we have

$$\sup_{t \in [-1, 1]} |\mathcal{M}(x(t), y) - \mathcal{M}_k(x(t), y)| \leq \frac{4M\rho^{-k}}{\rho - 1},$$

where $M = (2\sqrt{2}\exp(5/3))/\pi$ and $\rho = 3 + \sqrt{8}$.

Thus a Chebyshev polynomial interpolant of degree at most

$$\left\lceil \log\left(\frac{4\sqrt{2}e^{5/3}}{\pi(1+\sqrt{2})\varepsilon}\right) \middle/ \log(3 + \sqrt{8}) \right\rceil$$

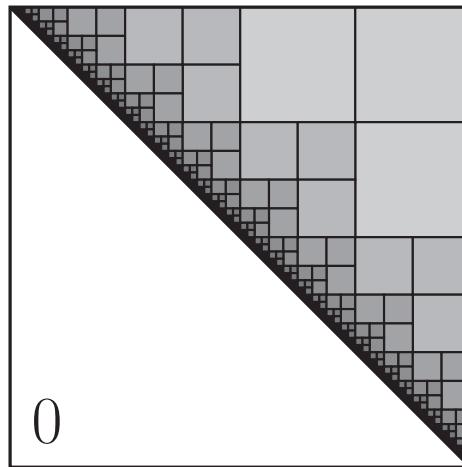


Figure 4.1. An illustration of the hierarchical decomposition of the Chebyshev–Legendre connection coefficients. The opacity illustrates the data-sparsity in the hierarchical decomposition.

is sufficient to approximate $\mathcal{M}(x(t), y)$ to within $\varepsilon > 0$ on any well-separated square. This naturally leads to a hierarchical partitioning of the connection coefficient matrix, as depicted in Figure 4.1. The matrix \mathcal{L} has a similar structure.

This approximation to \mathcal{M} can be applied to a vector in $O(N \log N)$ operations. If the function \mathcal{M} is interpolated in *both* variables, and *just* the right information is translated up and down the hierarchy, then the fast multipole method (FMM) (Greengard and Rokhlin 1987, Martinsson and Rokhlin 2007) accelerates the same matrix–vector product in $O(N)$ operations. This approach has a relatively large pre-computation cost, where by pre-computation we mean any computation that can be completed before knowing the vector. Still, once the data structure is created, it is the fastest approach for computing the Chebyshev–Legendre transform, making it ideal for applications in which the repeated action of a Chebyshev–Legendre transform is required.

4.5.2. The Toeplitz–Hankel approach

The method of Townsend *et al.* (2018) takes an alternative approach with a lower pre-computational cost. Observe that R_T^P can be expressed as the upper-triangular component of a diagonally scaled Hadamard (entry-wise) product of a Toeplitz matrix and a Hankel matrix. That is,

$$R_P^T = D[T \circ H],$$

where $D = \text{diag}(1/\pi, 2/\pi, 2/\pi, \dots)$,

$$T = \begin{pmatrix} \Lambda(0) & 0 & \Lambda(1) & 0 & \Lambda(2) & 0 & \cdots \\ & \Lambda(0) & 0 & \Lambda(1) & 0 & \Lambda(2) & \ddots \\ & & \Lambda(0) & 0 & \Lambda(1) & 0 & \ddots \\ & & & \ddots & \ddots & \ddots & \ddots \end{pmatrix},$$

$$H = \begin{pmatrix} \Lambda(0) & \Lambda(1/2) & \Lambda(1) & \cdots \\ \Lambda(1/2) & \Lambda(1) & \Lambda(3/2) & \cdots \\ \Lambda(1) & \Lambda(3/2) & \Lambda(2) & \ddots \\ \vdots & \vdots & \ddots & \ddots \end{pmatrix},$$

and ‘ \circ ’ denotes the Hadamard product, *i.e.* $(A \circ B)_{kj} = A_{kj}B_{kj}$. Now, Townsend *et al.* (2018) prove that H is a positive definite matrix and can be well-approximated by a low-rank matrix. In particular, it can be well-approximated as $H \approx \sum_{s=1}^r a_s \ell_s \ell_s^\top$, where $r = O(\log N)$, using a pivoted Cholesky algorithm. It then follows that

$$R_P^T \approx \sum_{s=1}^r a_s D \text{diag}(\ell_s) T \text{diag}(\ell_s).$$

Since applying an $N \times N$ Toeplitz matrix to a vector costs $O(N \log N)$ operations, we can compute $R_P^T v$ in $O(rN \log N) = O(N \log^2 N)$ operations. This extends naturally to applying R_T^P and some Jacobi–Jacobi transforms, where the explicit representation is also known.

4.6. The measure-perturbing connection problem

The algorithms mentioned above for Chebyshev–Legendre transforms are based on the explicit formula for the connection coefficients as the ratio of gamma functions. Beyond classical orthogonal polynomials, we cannot expect to have such precise information on the connection coefficients.⁴ Here we generalize the above construction based on the following observation: it suffices to know that the conversion matrix satisfies a triangular banded Sylvester equation (4.3) to apply it fast using a divide-and-conquer algorithm. The approach is formally equivalent to Keiner’s divide-and-conquer algorithms for eigenvectors of semi-separable matrices but expresses the problems in the more natural format of triangular-banded generalized eigenvalue problems (GEPs). The GEP form can be deduced via recurrence

⁴ Even the general Jacob–Jacobi connection coefficients are difficult to analyse, as they are proportional to the generalized hypergeometric functions ${}_3F_2$ (Andrews, Askey and Roy 1999, Lemma 7.1.1).

relationships and differential equations in an analogous way to the proof of Lemma 4.2.

Many classes of orthogonal polynomials satisfy a linear homogeneous differential equation. For example, classical orthogonal polynomials are the polynomial eigenfunctions of the second-order differential operators of the form

$$\sigma p_n'' + \tau p_n' = \lambda_n p_n, \quad (4.4)$$

where σ is a polynomial of degree at most two, τ is a polynomial of degree one, and $\lambda_n = (n/2)[(n-1)\sigma'' + 2\tau']$. This can be extended beyond classical orthogonal polynomials; see Appendix C. Upon closer inspection, both sides of (4.4) are degree-preserving. Since σ and τ are also polynomial, it follows that identities for differentiation of classical orthogonal polynomials allow one to obtain an upper-triangular and banded discretization of the differential equation. We now demonstrate this for connection problems involving Jacobi, Laguerre and ultraspherical polynomials.

4.6.1. The Jacobi–Jacobi connection problem

The Jacobi–Jacobi connection problem relates two families of Jacobi polynomials. We have

$$\mathbf{P}^{(\alpha, \beta)} = \mathbf{P}^{(\gamma, \delta)} R_{(\alpha, \beta)}^{(\gamma, \delta)}$$

for some $\alpha, \beta, \gamma, \delta > -1$. By considering the differential operator in (4.4) and using the recurrences from Appendix A.1, we find that

$$[(x^2 - 1)\mathcal{D}^2 + (\alpha - \beta + (\alpha + \beta + 2)x)\mathcal{D}] \mathbf{P}^{(\gamma, \delta)} = \mathbf{P}^{(\gamma+1, \delta+1)} A,$$

where

$$A = -L_{(\gamma+2, \delta+2)}^{(\gamma+1, \delta+1)} D_{(\gamma, \delta)}^{(\gamma+2, \delta+2)} + ((\alpha - \beta)I + (\alpha + \beta + 2)X_{(\gamma+1, \delta+1)}) D_{(\gamma, \delta)}^{(\gamma+1, \delta+1)},$$

while the right-hand side is

$$\mathbf{P}^{(\gamma, \delta)} = \underbrace{\mathbf{P}^{(\gamma+1, \delta+1)} R_{(\gamma, \delta)}^{(\gamma+1, \delta+1)}}_B.$$

Thus we have the following generalized eigenvalue equation:

$$A R_{(\alpha, \beta)}^{(\gamma, \delta)} = B R_{(\alpha, \beta)}^{(\gamma, \delta)} \Lambda,$$

where Λ is the diagonal eigenvalue matrix with $\Lambda_{n,n} = n(n + \alpha + \beta + 1)$.

A normalization of the generalized eigenvectors is necessary since (4.4) determines the eigenvectors only up to a scaling. This can be obtained by examining the leading coefficients of both families. For Jacobi polynomials, these are a ratio of Pochhammer symbols:

$$(R_{(\alpha, \beta)}^{(\gamma, \delta)})_{n,n} = \frac{(n + \alpha + \beta + 1)_n}{(n + \gamma + \delta + 1)_n},$$

which may be evaluated recursively.

4.6.2. The Laguerre–Laguerre connection problem

The Laguerre–Laguerre connection problem relates two families of Laguerre polynomials, that is,

$$\mathbf{L}^{(\alpha)} = \mathbf{L}^{(\beta)} \tilde{R}_{(\alpha)}^{(\beta)},$$

for some $\alpha, \beta > -1$. We have

$$\begin{aligned} & [-x\mathcal{D}^2 + (x - \alpha - 1)\mathcal{D}] \mathbf{L}^{(\beta)} \\ &= \mathbf{L}^{(\beta+1)} \underbrace{[-\tilde{L}_{(\beta+2)}^{(\beta+1)} \tilde{D}_{(\beta)}^{(\beta+2)} + (\tilde{X}_{(\beta+1)} - (\alpha + 1)I) \tilde{D}_{(\beta)}^{(\beta+1)}]}_A, \end{aligned}$$

while the right-hand side is

$$\mathbf{L}^{(\beta)} = \mathbf{L}^{(\beta+1)} \underbrace{\tilde{R}_{(\beta)}^{(\beta+1)}}_B$$

and $\Lambda_{n,n} = n$. For Laguerre polynomials, the normalization of the connection coefficients is trivial as $(\tilde{R}_{(\alpha)}^{(\beta)})_{n,n} = 1$.

4.6.3. The ultraspherical–ultraspherical connection problem

For the ultraspherical–ultraspherical connection problem, that is,

$$\mathbf{C}^{(\lambda)} = \mathbf{C}^{(\mu)} R_{(\lambda)}^{(\mu)}, \quad \text{for } \lambda, \mu > -1/2, \lambda, \mu \neq 0,$$

we note that (4.4) becomes

$$\begin{aligned} & [(x^2 - 1)\mathcal{D}^2 + (2\lambda + 1)x\mathcal{D}] \mathbf{C}^{(\mu)} \\ &= \mathbf{C}^{(\mu+1)} [-L_{(\mu+2)}^{(\mu+1)} D_{(\mu)}^{(\mu+2)} + (2\lambda + 1) X_{(\mu+1)} D_{(\mu)}^{(\mu+1)}] \\ &= \mathbf{C}^{(\mu+1)} R_{(\mu)}^{(\mu+1)} \Lambda, \end{aligned}$$

where $\Lambda_{n,n} = n(n + 2\lambda)$. For ultraspherical polynomials, the normalization of the connection coefficients is given by the ratio of Pochhammer symbols:

$$(R_{(\lambda)}^{(\mu)})_{n,n} = (\lambda)_n / (\mu)_n,$$

which again may be evaluated recursively.

4.6.4. The triangular-banded divide-and-conquer scheme

We consider a divide-and-conquer scheme to obtain a structured form of the generalized eigenvectors of a banded upper-triangular Sylvester equation

$$AV = BV\Lambda.$$

Since A and B are upper-triangular matrices, it is clear that the generalized eigenvalues are the ratios of the diagonal entries of the matrices. That is, $\Lambda_{i,i} = A_{i,i}/B_{i,i}$ for $i \geq 1$.

In what follows, we consider the $N \times N$ truncation of this system. That is, we consider $A, B \in \mathbb{R}^{N \times N}$. Let $s = \lfloor N/2 \rfloor$ and let b be the upper bandwidth of both A and B . We suppose that V has the form

$$V = \begin{pmatrix} V_{1,1} & \\ & V_{2,2} \end{pmatrix} \begin{pmatrix} I & V_{1,2} \\ & I \end{pmatrix},$$

where $V_{1,1} \in \mathbb{R}^{s \times s}$, $V_{1,2} \in \mathbb{R}^{s \times (N-s)}$ and $V_{2,2} \in \mathbb{R}^{(N-s) \times (N-s)}$ are all roughly the same size. Block-dividing A , B and Λ , we have

$$\begin{pmatrix} A_{1,1} & A_{1,2} \\ & A_{2,2} \end{pmatrix} \begin{pmatrix} V_{1,1} & V_{1,1}V_{1,2} \\ & V_{2,2} \end{pmatrix} = \begin{pmatrix} B_{1,1} & B_{1,2} \\ & B_{2,2} \end{pmatrix} \begin{pmatrix} V_{1,1}\Lambda_1 & V_{1,1}V_{1,2}\Lambda_2 \\ & V_{2,2}\Lambda_2 \end{pmatrix}.$$

This decomposes into three subproblems $A_{1,1}V_{1,1} = B_{1,1}V_{1,1}\Lambda_1$, $A_{2,2}V_{2,2} = B_{2,2}V_{2,2}\Lambda_2$ and

$$A_{1,1}V_{1,1}V_{1,2} + A_{1,2}V_{2,2} = B_{1,1}V_{1,1}V_{1,2}\Lambda_2 + B_{1,2}V_{2,2}\Lambda_2.$$

This last equation is a Sylvester matrix equation for the block $V_{1,2}$.

Supposing that $V_{1,1}^{-1}B_{1,1}^{-1}A_{1,1}V_{1,1} = \Lambda_1$, the matrix equation can be cast into the following form:

$$\Lambda_1 V_{1,2} - V_{1,2}\Lambda_2 = V_{1,1}^{-1}B_{1,1}^{-1}(B_{1,2}V_{2,2}\Lambda_2 - A_{1,2}V_{2,2}).$$

Since A and B are banded with an upper bandwidth of b , it follows that $A_{1,2}$ and $B_{1,2}$ are matrices with only a small number of non-zero entries, arranged in a lower-triangular fashion in their bottom left corners. Thus the matrix-matrix product $A_{1,2}V_{2,2}$ results in a rank- b matrix with non-zero entries only in the last b rows. Similarly, due to the zero pattern in the matrix $A_{1,2}V_{2,2} - B_{1,2}V_{2,2}\Lambda_2$, only the last b columns of $B_{1,1}^{-1}$ are needed, which are calculated with back-substitution using the last b columns of the $s \times s$ identity. We represent this block as the outer product of an $s \times b$ matrix X and a $b \times (N-s)$ matrix Y ,

$$V_{1,1}^{-1}B_{1,1}^{-1}(A_{1,2}V_{2,2} - B_{1,2}V_{2,2}\Lambda_2) =: XY^\top.$$

To conquer this problem, we examine the solution component-wise,

$$(V_{1,2})_{i,j} = \frac{(-XY^\top)_{i,j}}{(\Lambda_1)_{i,i} - (\Lambda_2)_{j,j}}.$$

If we use the following notation to denote a Cauchy matrix generated by two vectors $x, y \in \mathbb{R}^N$,

$$(C(x,y))_{i,j} = \frac{1}{x_i - y_j},$$

then $V_{1,2}$ is the Hadamard product of a finite-rank matrix $-XY^\top$ and a Cauchy matrix $C(\text{diag}(\Lambda_1), \text{diag}(\Lambda_2))$. This can be expressed as a sum of diagonal scalings of the same Cauchy matrix,

$$V_{1,2} = -\sum_{k=1}^b \text{diag}(X_{:,k}) C(\text{diag}(\Lambda_1), \text{diag}(\Lambda_2)) \text{diag}(Y_{:,k}).$$

Due to the separation of eigenvalues, the Cauchy matrices can be approximated, to within $0 < \epsilon < 1$, by a matrix of rank $O(b \log(\varepsilon^{-1}))$.

The structured form of the generalized eigenvectors permits inversion and transposition with the same computational complexity. We have

$$V^{-1} = \begin{pmatrix} I & -V_{1,2} \\ & I \end{pmatrix} \begin{pmatrix} V_{1,1}^{-1} \\ & V_{2,2}^{-1} \end{pmatrix},$$

$$V^\top = \begin{pmatrix} I & \\ V_{1,2}^\top & I \end{pmatrix} \begin{pmatrix} V_{1,1}^\top & \\ & V_{2,2}^\top \end{pmatrix},$$

$$V^{-\top} = \begin{pmatrix} V_{1,1}^{-\top} & \\ & V_{2,2}^{-\top} \end{pmatrix} \begin{pmatrix} I & \\ -V_{1,2}^\top & I \end{pmatrix}.$$

Tables 4.2 and 4.3 show the timings and memory usage in single and double precision using the general-purpose algorithm introduced above. The timings demonstrate the quasi-linear complexity of the algorithm as well as its high accuracy.

4.6.5. Words of caution

Not every measure-perturbing connection problem is well-conditioned. The structured form of the upper-triangular generalized eigenvectors allows for the following lower bound on the 2-norm condition number:

$$\text{cond}_2(V) := \|V\|_2 \|V^{-1}\|_2 \geq \rho(V)\rho(V^{-1}),$$

where $\rho(V)$ is the spectral radius of V . This lower bound is equal to the ratio of largest to smallest absolute diagonal entries. An upper bound of the condition number can be given in terms of the Frobenius norm:

$$\|V\|_2 \leq \|V\|_F \leq \sqrt{(\|V_{1,1}\|_F^2 + \|V_{2,2}\|_F^2)(N + \|V_{1,2}\|_F^2)}.$$

There are analogous upper and lower bounds for $\|V^{-1}\|_2$.

Keiner (2008) uses the lower bound on the condition number to determine when to subdivide the ultraspherical–ultraspherical connection problem into integer substeps between ultraspherical parameters λ and μ and the fractional remainder. This analysis also translates to analogous results for Jacobi–Jacobi and Laguerre–Laguerre connection problems (Keiner 2011). In all cases, it is best to use the banded connection problems for integer

Table 4.2. Calculation times, memory requirements and error of the Chebyshev–Legendre transform (and back), the archetype of the measure-perturbing connection problem, in single precision. s_d stands for the memory requirements of a dense $N \times N$ matrix, s_p stands for the memory requirements of the pre-computations, t_p stands for the pre-computation time, t_e stands for the execution time, averaged over ten trials, and ϵ_2 stands for the 2-norm relative error in transforming the coefficients $u_\ell = (\ell+1)^{-1}$, for $\ell = 0, \dots, N-1$, forward and backward over the ten trials.

N	s_d (bytes)	s_p (bytes)	t_p (seconds)	t_e (seconds)	ϵ_2
512	1.05e+06	4.10e+05	0.00069	0.000048	2.50e−07
2048	1.68e+07	2.39e+06	0.0059	0.00045	2.44e−07
8192	2.68e+08	1.26e+07	0.035	0.0026	2.32e−07
32768	4.29e+09	6.23e+07	0.23	0.015	2.27e−07
131072	6.87e+10	2.98e+08	1.36	0.071	3.18e−07
524288	1.10e+12	1.39e+09	7.48	0.34	3.64e−07
2097152	1.76e+13	6.35e+09	38.55	1.61	3.95e−07
8388608	2.81e+14	2.85e+10	201.26	7.51	3.89e−07

Table 4.3. Calculation times, memory requirements and error of the Chebyshev–Legendre transform (and back), the archetype of the measure-perturbing connection problem, in double precision. s_d stands for the memory requirements of a dense $N \times N$ matrix, s_p stands for the memory requirements of the pre-computations, t_p stands for the pre-computation time, t_e stands for the execution time, averaged over ten trials, and ϵ_2 stands for the 2-norm relative error in transforming the coefficients $u_\ell = 1/(\ell+1)$, for $\ell = 0, \dots, N-1$, forward and backward over the ten trials.

N	s_d (bytes)	s_p (bytes)	t_p (seconds)	t_e (seconds)	ϵ_2
512	2.10e+06	1.12e+06	0.00073	0.000057	2.60e−16
2048	3.36e+07	7.58e+06	0.0068	0.00065	4.54e−16
8192	5.37e+08	4.26e+07	0.059	0.0057	3.70e−16
32768	8.59e+09	2.18e+08	0.48	0.035	5.21e−16
131072	1.37e+11	1.07e+09	3.04	0.19	5.05e−16
524288	2.20e+12	5.06e+09	18.13	0.89	4.73e−16
2097152	3.52e+13	2.34e+10	94.52	4.02	5.60e−16
8388608	5.63e+14	1.06e+11	507.21	19.03	5.60e−16

increments/decrements in the parameters until the remaining difference in any parameter is less than 1.

For semiclassical orthogonal polynomials, the degree dependence in the variable coefficients in the differential equations enters into the linear algebra as diagonal scalings of the columns of the matrices of generalized eigenvectors. For example, it can be shown that the associated Jacobi–Jacobi connection coefficients satisfy

$$AV + BV\Lambda = CV\Omega,$$

where $\|A\|_2, \|\Omega\|_2 = O(N^4)$, $\|B\|_2, \|\Lambda\|_2 = O(N^2)$ and $\|C\|_2 = O(1)$. Thus each member of this three-term matrix equation is equally as important as the others as $N \rightarrow \infty$. While numerical experiments suggest the off-diagonal blocks in the resulting matrix equations remain low-rank, further analysis is required to show that they can be rapidly approximated.

4.7. The measure-preserving connection problem

Measure-preserving connection problems are those between different bases that are orthogonal with respect to the same measure μ . We focus our attention on weighted orthogonal polynomial expansions. For example, the weighted orthonormal Jacobi polynomials $\hat{\mathbf{P}}^{(2,2)} := (1-x^2)\mathbf{P}^{(2,2)}M_{(2,2)}^{-1/2}$ are orthonormal with respect to $w(x) = 1$ on $[-1,1]$, as are the normalized Legendre polynomials $\hat{\mathbf{P}} := \mathbf{P}M_P^{-1/2}$. That is,

$$\begin{aligned}\hat{\mathbf{P}}^\top \hat{\mathbf{P}} &= M_P^{-1/2} M_P M_P^{-1/2} = I, \\ (\hat{\mathbf{P}}^{(2,2)})^\top \hat{\mathbf{P}}^{(2,2)} &= M_{(2,2)}^{-1/2} (\mathbf{P}^{(2,2)})^\top (1-x^2)^2 \mathbf{P}^{(2,2)} M_{(2,2)}^{-1/2} = I.\end{aligned}$$

Thus we aim to understand the connection problem

$$\hat{\mathbf{P}}^{(2,2)} = \hat{\mathbf{P}} Q_{(2,2)}^P,$$

where we use the notation Q in place of C to emphasize that it is an orthogonal transformation. Therefore, a measure-preserving transform is equivalent to a matrix–vector product involving an orthogonal matrix. It is natural to view $Q_{(2,2)}^P$ (and other measure-preserving transforms) as a product of Givens rotations.

The measure-preserving connection problem is an important pathway to spherical harmonic synthesis and analysis as spherical harmonics involve the associated Legendre polynomials

$$Y_\ell^m(\theta, \varphi) = \sqrt{\frac{\ell+1/2}{2\pi} \frac{(\ell-m)!}{(\ell+m)!}} P_\ell^m(\cos\theta) e^{im\varphi}.$$

Over the past twenty years, there has been substantial progress in the construction and implementation of a fast spherical harmonic transform. Driscoll, Healy Jr and Rockmore (1997) introduced a procedure based on the so-called split-Legendre functions for constructing fast transforms for synthesis and analysis of spherical harmonics. In infinite-precision arithmetic, their algorithm is exact. However, their floating-point error analysis of the apparatus suggests the possibility of exponential growth in the error as m increases.

A few years after Driscoll *et al.* (1997), several groups including Potts, Steidl and Tasche (1998), Suda and Takami (2002), Healy Jr, Rockmore, Kostelec and Moore (2003) and Kunis and Potts (2003) sought more stable strategies and efficiency improvements on the original scheme.

In an approach that generalizes the use of asymptotic formulas to accelerate the synthesis and analysis of Legendre polynomials, Mohlenkamp (1999) used WKBJ asymptotic expansions of associated Legendre functions to accelerate the spherical harmonic transform. In addition to an asymptotically fast algorithm with predicted complexity $O(N^2 \log^2 N)$, Mohlenkamp also develops an intermediate transform with complexity $O(N^5/2 \log N)$.

In the same spirit as the measure-perturbing problems of Section 4.6, asymptotically fast algorithms with $O(N^2 \log N)$ complexity have been developed by Rokhlin and Tygert (2006) based on a divide-and-conquer approach applied to a discretization of the differential equation satisfied by the associated Legendre functions. This is then accelerated further by the FMM. While this was the first asymptotically fast algorithm with asymptotically low pre-computation, it is only predicted to outperform other algorithms for absurdly large degrees.

Tygert continued to develop alternative spherical harmonic transforms by accelerating the direct synthesis and analysis (Tygert 2008) via divide-and-conquer of the Jacobi operators combined with a re-interpolation from Gauss–Jacobi nodes to the Gauss–Legendre grid for every associated Legendre order. He also developed an alternative approach based on compressing synthesis and analysis by the butterfly algorithm (Tygert 2010a).

Notwithstanding this impressive body of literature, there are still improvements to be made. As a testament to this, some of the best software currently available for spherical harmonics is based on direct synthesis and analysis at Gaussian or equiangular grids (Schaeffer 2013, Reinecke and Seljebotn 2013, Ishioka 2018) with $O(N^3)$ complexity. One main bottleneck in making asymptotically fast transforms practical is the significant memory requirements in the storage of pre-computed data structures.

In light of these difficulties, Slevinsky (2019) shows that (1) the neighbouring associated Legendre connection problem may be factored as a product of Givens rotations, and (2) the butterfly algorithm may also be applied to the connection problem. These two features offer improved backward stability

of the connection problem, enable skeletonization of the pre-computation and provide insight on *a rank property* by relating connection coefficients to Fourier integral operators. A more vigorous skeletonization procedure is proposed in Slevinsky (2017), which offers the complexity of $O(N^{3/2} \log N)$ as a pre-computation cost in exchange for a modest increase in the execution complexity over the Rokhlin and Tygert (2006) scheme. With these results in hand, the question of making asymptotically fast transforms accessible at reasonably low degrees is no longer mathematical, but rather one of software. Our current progress in this endeavour is an open-source implementation, freely available in C (Slevinsky 2018a) and Julia (Slevinsky 2016).

Remark 4.5. If an orthogonal polynomial appears with a tilde on top, then this implies standard orthonormalization with respect to its Hilbert space. The so-called neighbouring connection problems are based on a data-sparse factorization of connection coefficients. That is, rather than using recurrence relations to evaluate orthogonal polynomials, we convert their expansions using sequences of Givens rotations. Heuristically speaking, this procedure is more accurate because Givens rotations preserve Euclidean lengths, while recurrence relations do not.

Let $I_{m \times n}$ denote the $m \times n$ rectangular identity matrix with ones on the main diagonal and zeros everywhere else.

4.7.1. The neighbouring connection problem

For any $x_0 \in \mathbb{R}$, let $(x - x_0)\varphi_n(x)$ and $\psi_n(x)$ be two families of orthonormal polynomials with respect to μ . The connection coefficients allow us to compute the expansion

$$(x - x_0)\varphi_n(x) = \sum_{\ell=0}^{n+1} q_{\ell,n} \psi_\ell(x).$$

As columns of the connection coefficients are mutually orthonormal, we may factor the principal rectangular finite sections as a product of $n + 1$ Givens rotations, the proof of which is by the *QR* factorization of a rectangular matrix with zeros below the first sub-diagonal.

Let G_n denote the real Givens rotation:

$$G_n = \begin{pmatrix} 1 & \cdots & 0 & 0 & \cdots & 0 \\ \vdots & \ddots & \vdots & \vdots & & \vdots \\ 0 & \cdots & c_n & s_n & \cdots & 0 \\ 0 & \cdots & -s_n & c_n & \cdots & 0 \\ \vdots & & \vdots & \vdots & \ddots & \vdots \\ 0 & \cdots & 0 & 0 & \cdots & 1 \end{pmatrix},$$

where $s_n = \sin \theta_n$ and $c_n = \cos \theta_n$, for some $\theta_n \in [0, 2\pi)$, are in the intersections of the n th and $(n+1)$ st rows and columns, embedded in the identity of a conformable size.

Theorem 4.6. The principal rectangular section of connection coefficients $Q \in \mathbb{R}^{(n+2) \times (n+1)}$ between $(x - x_0)\varphi_n(x)$ and $\psi_\ell(x)$ may be represented via the product of $n+1$ Givens rotations,

$$Q = G_0 \cdots G_n I_{(n+2) \times (n+1)},$$

where the sines and cosines for the Givens rotations are given by

$$s_n = -q_{n+1,n} \quad \text{and} \quad c_n = \sqrt{1 - q_{n+1,n}^2}.$$

By relating the sines and the cosines to the highest connection coefficient of each column, they may be readily determined by comparing the coefficients of the highest powers of x .

For Jacobi polynomials, two special values of x_0 are ± 1 , as both families may be identified as Jacobi polynomials with different but related parameters.

Corollary 4.7. The matrix of connection coefficients $Q^{(\alpha, \beta)} \in \mathbb{R}^{(n+2) \times (n+1)}$ between $(1-x)\tilde{P}_n^{(\alpha+2, \beta)}(x)$ and $\tilde{P}_\ell^{(\alpha, \beta)}(x)$ may be represented via the product of $n+1$ Givens rotations where the sines and cosines for the Givens rotations are given by

$$s_n^{(\alpha, \beta)} = \sqrt{\frac{(n+1)(n+\beta+1)}{(n+\alpha+2)(n+\alpha+\beta+2)}},$$

$$c_n^{(\alpha, \beta)} = \sqrt{\frac{(\alpha+1)(2n+\alpha+\beta+3)}{(n+\alpha+2)(n+\alpha+\beta+2)}}.$$

Since $P_n^{(\alpha, \beta)}(-x) = (-1)^n P_n^{(\beta, \alpha)}(x)$, the connection coefficients between $(1+x)\tilde{P}_n^{(\alpha, \beta+2)}(x)$ and $\tilde{P}_\ell^{(\alpha, \beta)}(x)$ are easily obtained by reversing the roles of α and β above and negating the sines.

For Laguerre polynomials, a special value of x_0 is 0, for the same reason.

Corollary 4.8. The matrix of connection coefficients $Q^{(\alpha)} \in \mathbb{R}^{(n+2) \times (n+1)}$ between $x\tilde{L}_n^{(\alpha+2)}(x)$ and $\tilde{L}_\ell^{(\alpha)}(x)$ may be represented via the product of $n+1$ Givens rotations where the sines and cosines for the Givens rotations are given by

$$s_n^{(\alpha)} = \sqrt{\frac{n+1}{n+\alpha+2}}, \quad c_n^{(\alpha)} = \sqrt{\frac{\alpha+1}{n+\alpha+2}}.$$

If μ admits families of orthogonal polynomials that have even–odd symmetry, i.e. $\psi_n(-x) = (-1)^n \psi_n(x)$, then the connection coefficients between

$(x^2 - x_0^2)\varphi_n(x)$ and $\psi_\ell(x)$ respect this symmetry through a chessboard pattern of zeros which is also observed in the Givens rotations.

Let G_n denote the real Givens rotation

$$G_n = \begin{pmatrix} 1 & \cdots & 0 & 0 & 0 & \cdots & 0 \\ \vdots & \ddots & \vdots & \vdots & \vdots & & \vdots \\ 0 & \cdots & c_n & 0 & s_n & \cdots & 0 \\ 0 & \cdots & 0 & 1 & 0 & \cdots & 0 \\ 0 & \cdots & -s_n & 0 & c_n & \cdots & 0 \\ \vdots & & \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & \cdots & 0 & 0 & 0 & \cdots & 1 \end{pmatrix},$$

where the sines $s_n = \sin \theta_n$ and the cosines $c_n = \cos \theta_n$, for some $\theta_n \in [0, 2\pi)$, are now in the intersections of the n th and $(n+2)$ nd rows and columns, embedded in the identity of a conformable size.

Theorem 4.9. The principal rectangular section of connection coefficients $Q \in \mathbb{R}^{(n+3) \times (n+1)}$ between $(x^2 - x_0^2)\varphi_n(x)$ and $\psi_\ell(x)$ may be represented via the product of $n+1$ Givens rotations,

$$Q = G_0 \cdots G_n I_{(n+3) \times (n+1)},$$

where the sines and cosines for the Givens rotations are given by

$$s_n = -q_{n+2,n} \quad \text{and} \quad c_n = \sqrt{1 - q_{n+2,n}^2}.$$

For Jacobi polynomials with equal parameters, a special value of x_0 is 1.

Corollary 4.10. The matrix of connection coefficients $Q^{(\alpha)} \in \mathbb{R}^{(n+3) \times (n+1)}$ between $(1 - x^2)\tilde{P}_n^{(\alpha+2, \alpha+2)}(x)$ and $\tilde{P}_\ell^{(\alpha, \alpha)}(x)$ may be represented via the product of $n+1$ Givens rotations, where the sines and cosines for the Givens rotations are given by

$$s_n^{(\alpha)} = \sqrt{\frac{(n+1)(n+2)}{(n+2\alpha+3)(n+2\alpha+4)}},$$

$$c_n^{(\alpha)} = \sqrt{\frac{(2\alpha+2)(2n+2\alpha+5)}{(n+2\alpha+3)(n+2\alpha+4)}}.$$

The representations for the Givens rotations of the neighbouring connection problems along with the representations of the sines and cosines as square roots of rationals *ipso facto* enables the design and implementation of backward stable direct algorithms (Slevinsky 2019, Algorithm 2.5).

Table 4.4. The 2-norm and ∞ -norm relative errors in a full spherical harmonic synthesis and analysis on an equiangular grid with pseudo-random expansion coefficients drawn from $U(-1,1)$.

$N+1$	SHTns		ISPACK		Fast Transforms	
	ϵ_2	ϵ_∞	ϵ_2	ϵ_∞	ϵ_2	ϵ_∞
1024	5.4e-14	7.9e-13	4.6e-14	6.8e-13	2.0e-15	8.2e-15
2048	1.3e-13	3.4e-12	9.4e-14	1.2e-12	2.8e-15	1.2e-14
4096	2.6e-13	9.4e-12	2.0e-13	5.5e-12	4.0e-15	2.0e-14
8192	7.0e-13	5.5e-11	4.5e-13	1.6e-11	5.9e-15	4.2e-14
16384	1.0e-12	6.4e-11	8.3e-13	3.9e-11	7.9e-15	6.7e-14

Table 4.4 shows a full spherical harmonic synthesis and analysis compared with SHTns and ISPACK, two high-performance spherical harmonic libraries based on direct synthesis and analysis. Compare with Ishioka (2018, Table 2).

4.7.2. Associated Legendre functions and symmetric definite and banded generalized eigenvalue problems

For asymptotically fast transforms, we must consider more than the neighbouring connection problem. In this case, we would again like to exploit the differential eigenvalue problems satisfied by weighted orthonormal polynomials. A hat on top of a symbol denotes that the orthonormal polynomials are weighted by the square root of the orthonormality weight. For example, the weighted Jacobi polynomials satisfy

$$\hat{P}_n^{(\alpha, \beta)}(x) := (1-x)^{\alpha/2}(1+x)^{\beta/2} \tilde{P}_n^{(\alpha, \beta)}(x).$$

If $\beta = \alpha = m \in \mathbb{N}_0$, then it is well known that these weighted Jacobi polynomials are proportional to associated Legendre functions, which are the latitudinal component of spherical harmonics.

Normalized associated Legendre functions of degree ℓ and order m are eigenfunctions of the following linear differential equation:

$$(1-x^2)(-\mathcal{D})[(1-x^2)\mathcal{D}\tilde{P}_\ell^m(x)] + m^2\tilde{P}_\ell^m(x) = \ell(\ell+1)(1-x^2)\tilde{P}_\ell^m(x), |m| \leq \ell, \quad (4.5)$$

which we write in quasimatrix notation as

$$(1-x^2)(-\mathcal{D})(1-x^2)\mathcal{D}\mathbf{P}^m + m^2\mathbf{P}^m = \mathbf{P}^m \underbrace{(I - X_m^2)}_M \Lambda_m$$

for X_m defined in (A.1). This means that M is a pentadiagonal, symmetric

and positive definite matrix. If we write $\mathbf{P}^\nu = \mathbf{P}^m Q_\nu^m$, then we find that

$$\begin{aligned} (1-x^2)(-\mathcal{D})(1-x^2)\mathcal{D}\mathbf{P}^\nu &= (1-x^2)(-\mathcal{D})(1-x^2)\mathcal{D}\mathbf{P}^m Q_\nu^m \\ &\Downarrow \\ \mathbf{P}^\nu[-\nu^2 I + (I - X_\nu^2)\Lambda_\nu] &= \mathbf{P}^m[-m^2 I + (I - X_m^2)\Lambda_m] \\ &\Downarrow \\ \mathbf{P}^m[-\nu^2 Q_\nu^m + (I - X_m^2)Q_\nu^m \Lambda_\nu] &= \mathbf{P}^m[-m^2 I + (I - X_m^2)\Lambda_m]Q_\nu^m \\ &\Downarrow \\ MQ_\nu^m \Lambda_\nu &= [(\nu^2 - m^2)I + M\Lambda_m]Q_\nu^m. \end{aligned}$$

That is, we have

$$[M\Lambda_m + (\nu^2 - m^2)I]Q_\nu^m = MQ_\nu^m \Lambda_\nu.$$

Unfortunately, the symmetry in M is destroyed by the column scaling in Λ . Rokhlin and Tygert (2006) realized that the generalized eigenvalue problem is symmetrized when multiplied by M^{-1} from the left, resulting in

$$[\Lambda_m + (\nu^2 - m^2)M^{-1}]Q_\nu^m = Q_\nu^m \Lambda_\nu.$$

This symmetric diagonal-plus-semiseparable regular eigenvalue problem is the starting point for their fast spherical harmonic transform (Rokhlin and Tygert 2006). The FMM accelerates the divide-and-conquer algorithm for symmetric diagonal-plus-semiseparable eigenvalue problems developed by Chandrasekaran and Gu (2004).

An alternative to multiplication by M^{-1} is to use the Cholesky factorization and a change of bases. If we take the Cholesky factorization of $M = R^\top R$, consider new generalized eigenvectors $V = R^{-\top} Q_\nu^m$ and multiply by $R^{-\top}$ from the left, then we transform the problem into the following:

$$[R\Lambda_m R^\top + (\nu^2 - m^2)I]V = RR^\top V \Lambda_\nu.$$

This generalized eigenvalue problem is symmetric definite and pentadiagonal⁵ allowing for a fast spectral decomposition (Borges and Gragg 1993, Gu and Eisenstat 1995). Furthermore, the generalized eigenvectors V are (RR^\top) -orthogonal, in the sense that

$$V^\top RR^\top V = I.$$

In Appendix A of Slevinsky (2017), explicit formulas are given for the Cholesky factor R and the products RR^\top and $R\Lambda_m R^\top$.

In this case, the problem is pentadiagonal, but the zero pattern ensures that it may be decoupled into two symmetric definite and tridiagonal problems.

⁵ It is in fact tridiagonal if the even–odd symmetry is highlighted by a perfect shuffle.

4.7.3. Weighted normalized Jacobi polynomials

The construction for associated Legendre polynomials is also possible for other weighted Jacobi transforms, arising in fast transforms for OPs. The weighted normalized Jacobi polynomials $\hat{P}_n^{(\alpha,\beta)}(x)$ are eigenfunctions of the following linear differential equation:

$$\begin{aligned} & (1-x^2)(-\mathcal{D})[(1-x^2)\mathcal{D}\hat{P}_n^{(\alpha,\beta)}(x)] \\ & + \left(\left(\frac{\alpha}{2}\right)^2(1+x)^2 + \left(\frac{\beta}{2}\right)^2(1-x)^2 - \frac{(\alpha\beta+\alpha+\beta)}{2}(1-x^2) \right) \hat{P}_n^{(\alpha,\beta)}(x) \\ & = n(n+\alpha+\beta+1)(1-x^2)\hat{P}_n^{(\alpha,\beta)}(x). \end{aligned} \quad (4.6)$$

This Sturm–Liouville problem is the two-parameter generalization of (4.5) for associated Legendre functions. When we view multiplication by $1 \pm x$ as an operator acting on the basis of weighted normalized Jacobi polynomials, we can identify that certain terms in (4.6) are symmetric positive definite and pentadiagonal matrices.

If we expand the weighted normalized Jacobi polynomials of parameters γ and δ in the basis of weighted normalized Jacobi polynomials of parameters α and β , that is,

$$\hat{P}^{(\gamma,\delta)}(x) = \hat{P}^{(\alpha,\beta)}(x)Q_{(\gamma,\delta)}^{(\alpha,\beta)},$$

then we may rewrite the Sturm–Liouville problem in (4.6) symbolically as

$$\left\{ M\Lambda_{(\alpha,\beta)} + \left[\left(\frac{\gamma}{2}\right)^2 - \left(\frac{\alpha}{2}\right)^2 \right] M^+ + \left[\left(\frac{\delta}{2}\right)^2 - \left(\frac{\beta}{2}\right)^2 \right] M^- \right. \\ \left. - \left(\frac{\gamma\delta + \gamma + \delta - \alpha\beta - \alpha - \beta}{2} \right) M \right\} Q_{(\gamma,\delta)}^{(\alpha,\beta)} = MQ_{(\gamma,\delta)}^{(\alpha,\beta)}\Lambda_{(\gamma,\delta)}. \quad (4.7)$$

Here we have $M := I - X_{(\alpha,\beta)}^2$, $M^+ := (I + X_{(\alpha,\beta)})^2$, $M^- := (I - X_{(\alpha,\beta)})^2$, and $\Lambda_{(\alpha,\beta)}$ is the diagonal scaling of the basis by $n(n+\alpha+\beta+1)$ for $n \geq 0$.

Due to the orthonormal polynomial recurrence relation, all the multiplication operators may be derived analytically, and algorithms for their component-wise computation to high relative accuracy are described in Appendix B in Slevinsky (2017).

If we let S denote the weighted sum of symmetric positive definite multiplication operators on the left-hand side of (4.7), then we have

$$(M\Lambda_{(\alpha,\beta)} + S)Q_{(\gamma,\delta)}^{(\alpha,\beta)} = MQ_{(\gamma,\delta)}^{(\alpha,\beta)}\Lambda_{(\gamma,\delta)},$$

where M is a symmetric positive definite and banded matrix, $\Lambda_{(\alpha,\beta)}$ is diagonal and S is symmetric and banded. If we use the same strategy as

for spherical harmonics, then we would take the Cholesky factorization of $M = R^\top R$ and rearrange to obtain

$$(R\Lambda_{(\alpha,\beta)}R^\top + R^{-\top}SR^\top)V = RR^\top V\Lambda_{(\gamma,\delta)}. \quad (4.8)$$

As before, $R\Lambda_{(\alpha,\beta)}R^\top$ is symmetric and banded and RR^\top is symmetric positive definite and banded, but the new term $R^{-\top}SR^\top$ is no longer a scalar multiple of the identity matrix (as for spherical harmonics). In fact, not much of its structure is apparent by this formulation.

To reveal the structure of $R^{-\top}SR^\top$, we make use of another property of our eigenproblem. Since M and S are multiplication operators with the same separable Hilbert spaces attached to the domain and range, they commute. That is, $MS = SM$, or in terms of the commutator $[M, S] = 0$. We use the commutator and the (formal) invertibility of M to write $S = MSM^{-1}$, leading to the equivalent representation

$$(R\Lambda_{(\alpha,\beta)}R^\top + RSR^{-1})V = RR^\top V\Lambda_{(\gamma,\delta)}.$$

By the Cholesky factorization, R is an upper-triangular and banded matrix. Thus $R^{-\top}SR^\top$ has non-zero entries up to and including the second superdiagonal. Furthermore

$$RSR^{-1} = (R^{-\top}S^\top R^\top)^\top = (R^{-\top}SR^\top)^\top = R^{-\top}SR^\top.$$

It turns out that $R^{-\top}SR^\top$ in (4.8) is in fact a symmetric and banded matrix.

4.7.4. Weighted Laguerre polynomials

The weighted normalized Laguerre polynomials $\hat{L}_n^{(\alpha)}(x)$ are eigenfunctions of the linear differential equation

$$x(-D)[x\mathcal{D}\hat{L}_n^{(\alpha)}(x)] + \frac{1}{4}(\alpha - x)^2\hat{L}_n^{(\alpha)}(x) = \left(n + \frac{1}{2}\right)x\hat{L}_n^{(\alpha)}(x).$$

We may view M as multiplication by x , and Λ as a diagonal scaling. When $\hat{\mathbf{L}}^{(\alpha)}(x) = \hat{\mathbf{L}}^{(\beta)}Q_{(\alpha)}^{(\beta)}$, we again arrive at

$$(M\Lambda + S)Q_{(\alpha)}^{(\beta)} = MQ_{(\alpha)}^{(\beta)}\Lambda,$$

where $[M, S] = 0$. Since M is a symmetric and positive definite matrix, this equation is symmetrizable by the same technique as above.

4.8. Symmetric definite and banded divide-and-conquer

We have shown that measure-preserving connection problems for weighted and orthonormalized classical orthogonal polynomials result in the symmetric definite and banded generalized eigenproblem of the form

$$AV = BV\Lambda, \quad V^\top BV = I. \quad (4.9)$$

There are several significant contributions to the literature on divide-and-conquer for special types of eigenproblems, with special attention to rank-one symmetric modifications of known eigenproblems in the conquer step. This includes the seminal works of Bunch, Nielsen and Sorensen (1978) and Cuppen (1981) and improvements to stability in Gu and Eisenstat (1994). For the related symmetric arrowhead conquer step, this is also described in Gu and Eisenstat (1995) and includes the observation that the FMM may accelerate an eigenmatrix–vector product. For the generalized eigenproblem, the tridiagonal case was described in Borges and Gragg (1993) with arrowhead divide-and-conquer and Elsner, Fasse and Langmann (1997) with a symmetric rank-one modification. These algorithms have been adapted quite recently to the hierarchical semi-separable format (Vogel, Xia, Cauley and Balakrishnan 2016). Here, we extend divide-and-conquer to the symmetric definite and banded generalized eigenvalue problem in (4.9).

Let $A, B \in \mathbb{R}^{N \times N}$, and let $s = \lfloor N/2 \rfloor$ and $b \in \mathbb{N}$ be the (symmetric) bandwidth of both A and B . We divide A and B conformably,

$$A = \begin{pmatrix} A_1 & A_3^\top \\ A_3 & A_2 \end{pmatrix}, \quad B = \begin{pmatrix} B_1 & B_3^\top \\ B_3 & B_2 \end{pmatrix},$$

where

$$A_1, B_1 \in \mathbb{R}^{s \times s}, \quad A_2, B_2 \in \mathbb{R}^{(N-s) \times (N-s)} \quad \text{and} \quad A_3, B_3 \in \mathbb{R}^{(N-s) \times s}.$$

Since A and B are banded, A_3 and B_3 have the canonical decompositions

$$\begin{aligned} A_3 &= \sum_{j=1}^b \sum_{i=1}^j (A_3)_{i,s-b+j} e_i e_{s-b+j}^\top, \\ B_3 &= \sum_{j=1}^b \sum_{i=1}^j (B_3)_{i,s-b+j} e_i e_{s-b+j}^\top. \end{aligned}$$

Thus we may represent A and B as finite-rank perturbations of two divided problems of nearly half the original size. While there are a few choices to the finite-rank perturbations, in the context of an eigensolver, it is important to choose perturbations that are symmetric and favourable to the root-finder. We will now show that this is always possible.

If we define

$$w_{i,j} = e_{i+s} + \widetilde{\operatorname{sgn}}[(B_3)_{i,s-b+j}] e_{s-b+j}, \quad \widetilde{\operatorname{sgn}}(x) = \begin{cases} 1 & x \geq 0, \\ -1 & x < 0, \end{cases}$$

then A and B may be decomposed as

$$A = \begin{pmatrix} \tilde{A}_1 \\ \tilde{A}_2 \end{pmatrix} + \sum_{j=1}^b \sum_{i=1}^j \widetilde{\text{sgn}}[(B_3)_{i,s-b+j}] (A_3)_{i,s-b+j} w_{i,j} w_{i,j}^\top, \quad (4.10)$$

$$B = \begin{pmatrix} \tilde{B}_1 \\ \tilde{B}_2 \end{pmatrix} + \sum_{j=1}^b \sum_{i=1}^j |(B_3)_{i,s-b+j}| w_{i,j} w_{i,j}^\top, \quad (4.11)$$

where

$$\begin{aligned} \tilde{A}_1 &= A_1 - \sum_{j=1}^b \sum_{i=1}^j \widetilde{\text{sgn}}[(B_3)_{i,s-b+j}] (A_3)_{i,s-b+j} e_{s-b+j} e_{s-b+j}^\top, \\ \tilde{A}_2 &= A_2 - \sum_{j=1}^b \sum_{i=1}^j \widetilde{\text{sgn}}[(B_3)_{i,s-b+j}] (A_3)_{i,s-b+j} e_i e_i^\top, \\ \tilde{B}_1 &= B_1 - \sum_{j=1}^b \sum_{i=1}^j |(B_3)_{i,s-b+j}| e_{s-b+j} e_{s-b+j}^\top, \\ \tilde{B}_2 &= B_2 - \sum_{j=1}^b \sum_{i=1}^j |(B_3)_{i,s-b+j}| e_i e_i^\top. \end{aligned}$$

Next, assume that the two subproblems have been solved, so that

$$V_i^\top (\tilde{A}_i - \lambda \tilde{B}_i) V_i = D_i - \lambda I \quad \text{for } i = 1, 2.$$

Then (4.10) and (4.11) are congruent to the $b(b+1)/2$ -rank symmetric perturbation of the diagonal generalized eigenvalue problem. That is, there exist vectors

$$z_{i,j} = \begin{pmatrix} V_1^\top \\ V_2^\top \end{pmatrix} w_{i,j},$$

so that

$$\begin{aligned} &\begin{pmatrix} V_1^\top \\ V_2^\top \end{pmatrix} (A - \lambda B) \begin{pmatrix} V_1 \\ V_2 \end{pmatrix} \\ &= \begin{pmatrix} D_1 \\ D_2 \end{pmatrix} - \lambda I + \sum_{j=1}^b \sum_{i=1}^j (\rho_{i,j} - \lambda \sigma_{i,j}) z_{i,j} z_{i,j}^\top, \end{aligned} \quad (4.12)$$

where $\rho_{i,j} = \widetilde{\text{sgn}}[(B_3)_{i,s-b+j}] (A_3)_{i,s-b+j}$ and $\sigma_{i,j} = |(B_3)_{i,s-b+j}| \geq 0$.

In the conquer step, we must now solve the finite-rank symmetric perturbation of the diagonal generalized eigenvalue problem. This leads us to consider, in turn, the rank-one symmetric perturbation of the diagonal generalized eigenvalue problem, as the finite-rank case may be broken into a sequence of rank-one updates.

Based on (4.12), we consider the following generalized eigenvalue problem:

$$(D + \rho z z^\top) v = \lambda(I + \sigma z z^\top) v, \quad (4.13)$$

where $\sigma \geq 0$. The generalized eigenvalue problem may be deflated in the following three cases.

- (1) If $z_i = 0$, then $\lambda_i = d_i$ is an eigenvalue with the corresponding eigenvector $v = e_i$.
- (2) If $d_i = d_j$ for some pair of indices $i \neq j$, then a Givens rotation in the (i,j) -plane may be elected to introduce a zero in either z_i or z_j , while leaving the entries of D unchanged. Thus $\lambda_i = d_i$ and we deflate as above.
- (3) If $\tau := \rho/\sigma = d_i$ for some index i , then $v = e_i$ is also an eigenvector with eigenvalue $\lambda_i = \tau$.

Therefore, by elementary deflation and permutation, we assume that the diagonal entries of D are increasing, *i.e.* $d_1 < \dots < d_N$, $z_i \neq 0$, $\tau \neq d_i$, and that $z^\top z = 1$. The scaling of z may be absorbed by the constants ρ and σ .

The following theorem was proved substantially by Elsner *et al.* (1997), although they did not present the case where $\sigma \geq 0$, which allows for an exact localization of all N eigenvalues.

Theorem 4.11. With all aforementioned assumptions, the generalized eigenvalues of (4.13) are roots of the secular equation

$$f(\lambda) = \frac{1}{\rho - \sigma\lambda} + \sum_{i=1}^N \frac{z_i^2}{d_i - \lambda}.$$

Let $d_0 = -\infty$ and $d_{N+1} = +\infty$. Let $d_j < \tau (= \rho/\sigma) < d_{j+1}$ for some $j = 0, \dots, N$. Then there exists an eigenvalue in each interval (d_i, d_{i+1}) for $1 \leq i < j$, one in the interval (d_j, τ) , another in the interval (τ, d_{j+1}) and one more in each interval (d_i, d_{i+1}) for $j < i < N$. Furthermore, the corresponding generalized eigenvector v_j to λ_j is given by

$$v_j = \left(\frac{z_1}{d_1 - \lambda_j}, \frac{z_2}{d_2 - \lambda_j}, \dots, \frac{z_N}{d_N - \lambda_j} \right)^\top. \quad (4.14)$$

The proof of this theorem relies on the observation that f is monotonically increasing, given that $\sigma \geq 0$. In our choice of the finite-rank decomposition of A and B in (4.10) and (4.11), this is always true and we need never consider the monotonicity-breaking case of $\sigma < 0$.

By inspecting the equation

$$v_j^\top (I + \sigma z z^\top) v_j = \sum_{i=1}^N \frac{z_i^2}{(d_i - \lambda_j)^2} + \sigma \left(\sum_{i=1}^N \frac{z_i^2}{d_i - \lambda_j} \right)^2,$$

Table 4.5. Calculation times, memory requirements and error of the associated Legendre transform $\tilde{P}_{2n}^{2n}(x)$ to $\tilde{P}_{2n}^0(x)$ (and back), in single precision. s_d stands for the memory requirements of a dense $n \times n$ matrix, s_p stands for the memory requirements of the pre-computations, t_p stands for the pre-computation time, t_e stands for the execution time, averaged over ten trials, and ϵ_2 stands for the 2-norm relative error in transforming the coefficient $u_0 = 1$ forward and backward over the ten trials.

n	s_d (bytes)	s_p (bytes)	t_p (seconds)	t_e (seconds)	ϵ_2
256	2.62e+05	3.91e+05	0.0044	0.000048	1.46e-07
1024	4.19e+06	5.93e+06	0.056	0.00088	3.00e-07
4096	6.71e+07	6.24e+07	0.72	0.0089	1.09e-05
16384	1.07e+09	4.79e+08	8.67	0.069	2.31e-04
65536	1.72e+10	3.17e+09	106.23	0.45	1.39e-03

Table 4.6. Calculation times, memory requirements and error of the associated Legendre transform $\tilde{P}_{2n}^{2n}(x)$ to $\tilde{P}_{2n}^0(x)$ (and back), in double precision. s_d stands for the memory requirements of a dense $n \times n$ matrix, s_p stands for the memory requirements of the pre-computations, t_p stands for the pre-computation time, t_e stands for the execution time, averaged over ten trials, and ϵ_2 stands for the 2-norm relative error in transforming the coefficient $u_0 = 1$ forward and backward over the ten trials.

n	s_d (bytes)	s_p (bytes)	t_p (seconds)	t_e (seconds)	ϵ_2
256	5.24e+05	7.95e+05	0.011	0.000047	5.78e-16
1024	8.39e+06	1.26e+07	0.11	0.0011	7.98e-16
4096	1.34e+08	1.66e+08	1.42	0.016	1.49e-14
16384	2.15e+09	1.47e+09	17.64	0.14	2.16e-13
65536	3.44e+10	1.05e+10	211.40	1.02	2.38e-12

we are provided with a natural choice for the normalization

$$\tilde{v}_j = v_j / \sqrt{v_j^\top (I + \sigma z z^\top) v_j}. \quad (4.15)$$

Both (4.14) and (4.15) show the connection to diagonally scaled Cauchy and Coulomb matrices whose matrix–vector products may be accelerated by the FMM.

4.8.1. Computational issues

To find the roots of the secular equations, we use the cubically convergent algorithm described by Borges and Gragg (1993). This uses inverse rational interpolation for the function

$$\Phi_i(\lambda) := p + \frac{q}{d_i - \lambda} + \frac{r}{d_{i+1} - \lambda},$$

with suitable modifications for the eigenvalues in (d_j, τ) and (τ, d_{j+1}) . The constants p , q and r are found by equating $\Phi_i^{(\nu)}(\lambda^{(k)}) = f^{(\nu)}(\lambda^{(k)})$, for $\nu = 0, 1, 2$, and for every approximate eigenvalue $\lambda^{(k)}$, for $k = 0, 1, \dots$. The one root of Φ_i in the interval (d_i, d_{i+1}) is a solution of a quadratic equation.

It may occur that λ_i is near one of the poles, d_i or d_{i+1} . In this case, it is advantageous to simulate an extended precision by representing λ_i as the sum of λ_i^{hi} , defined by

$$\lambda_i^{\text{hi}} = \begin{cases} d_i & \text{for } f\left(\frac{d_i + d_{i+1}}{2}\right) > 0, \\ d_{i+1} & \text{otherwise,} \end{cases}$$

and λ_i^{lo} , defined by the difference $\lambda_i - \lambda_i^{\text{hi}}$, though the difference is not explicitly computed. Instead, there is a shift of origin in the secular equation where difference is transferred to the diagonal entries d_i . The numerical instability as the result of dividing by the difference of two nearly equal numbers is thereby eliminated.

Tables 4.5 and 4.6 show the timings and memory usage in single and double precision using the general-purpose algorithm introduced above. The timings help to demonstrate the quasi-linear complexity. The symmetric definite and banded divide-and-conquer approaches have a much higher pre-computation than the upper-triangular and banded divide-and-conquer for the measure-perturbing connection problem.

5. Ordinary differential equations

Orthogonal polynomials have been used in the numerical solution of ordinary differential equations for a long time, going back to Lanczos (1938). A key contribution was the integral reformulation method of Clenshaw (1957), which observed that general linear differential equations with polynomial coefficients could be reduced to a banded, infinite-dimensional linear system using recurrence relationships of Chebyshev polynomials, via reformulation as an integral equation. Since then integral reformulation using Chebyshev polynomials has been rediscovered in a variety of contexts (Coutsias, Hagstrom and Torres 1996, Greengard 1991, Zebib 1984, Viswanath 2015).

A related scheme from the finite element method (FEM) community is the p -FEM of Shen (1994), which uses a similar basis as Heinrichs (1989) applied to the weak formulation. These approaches lead naturally to the p -FEM on tetrahedra of Beuchler and Schöberl (2006) and Li and Shen (2010) discussed in Section 6.6. There are two major benefits of this approach due to it being applied to the weak formulation of a differential equation: (1) the discretizations are automatically symmetric if the differential equation is self-adjoint, and (2) only continuity (as opposed to continuity of the derivative) is required if an interval is subdivided into multiple elements (though this is less relevant for ODEs than for PDEs). However, p -FEMs are less versatile for higher-order equations and general boundary conditions, and they also require a problem to be posed in weak formulation before being discretized.

The idea of using recurrence relationships to represent differential operators using banded linear systems is a continuing theme in the ultraspherical spectral method of Olver and Townsend (2013), where sparse recurrence relationships between ultraspherical polynomials are employed. The ultraspherical spectral method is similar in spirit to Clenshaw's integral reformulation and p -FEMs, with the benefit of increased versatility. It automatically produces banded discretizations for general differential operators, including operators with both differentiation and multiplication, leading to optimal $O(N)$ complexity computations. In particular, it is applicable for Sturm–Liouville equations and high-order equations and is extensible to partial differential equations and singular integral equations, as discussed in Sections 6 and 7.

Each of the aforementioned methods produces a banded infinite-dimensional system which can be solved in multiple ways, as noted originally by Clenshaw (1957):

The relations between the coefficients, obtained by application of the fundamental equations, constitute an infinite set of linear algebraic equations in the unknowns a_r and $a_r^{(s)}$, which can be solved to any required accuracy by recurrence or iteration. We shall usually use the former method ...

In particular, there are various approaches.

- (1) *Shooting method.* Guess the tail coefficients and do back-substitution (Clenshaw 1957).
- (2) *Finite section method.* Truncate the operator and solve the resulting finite linear system.
- (3) *τ -method.* Perform a rectangular truncation that captures the operators exactly with extra columns added to make the problem square (Lanczos 1938).

- (4) (*F. W. J. Olver's algorithm.*) Perform adaptive Gaussian elimination until a convergence criterion is achieved, at which point perform back-substitution (Olver 1967, Lozier 1980).
- (5) *Adaptive QR method.* Perform adaptive Givens rotations or Householder reflections until a convergence criterion is achieved, at which point perform back-substitution (Olver and Townsend 2013).

A detailed description of the benefits and trade-offs of these techniques is outside the scope of this paper. Still, we will advocate for Olver's algorithm for diagonally dominant discretizations and the adaptive QR algorithm otherwise, as they achieve adaptive error control and stability, which leads to faster algorithms as the same equation need not be solved multiple times.

Clenshaw's integral reformulation, ultraspherical spectral methods and p -finite element methods are fast in the sense that they achieve $O(N)$ complexity for N degrees of freedom, due to the sparsity of the discretization, provided any variable coefficients are smooth and therefore well-approximated by polynomials. There are a slew of non-fast methods based on collocation, Galerkin approximations and the Chebyshev-tau method that represent differentiation with a dense matrix. We will not discuss these methods further.

We also do not discuss in detail the convergence proofs of the various methods. For non-singular differential equations, Clenshaw's integral reformulation, ultraspherical spectral methods and p -finite element methods represent them as operators that can be pre- or post-conditioned to be a compact perturbation of the identity. This gives us the machinery of second-kind Fredholm operators to rely on (Olver and Townsend 2013).

5.1. Clenshaw's integral reformulation

The idea behind Chebyshev-based integral reformulation is to multiply the differential equation by indefinite integration operators on either the left or right and use the relationship given in Proposition A.17. Clenshaw (1957) applied the relationship in Proposition A.17 on the right. For the Airy equation we find that

$$u = \mathcal{Q}^2 \mathbf{T} \mathbf{v} = \mathbf{T} Q^2 \mathbf{v},$$

where \mathcal{Q} is the indefinite integration operator so that the zeroth Chebyshev coefficient is zero and $Q \equiv Q_T$ is the associated matrix from Proposition A.17.⁶ For the Airy example we have

$$(\mathcal{D}^2 - x) \iint \mathbf{T} \mathbf{v} = \mathbf{T} (I - X Q^2) \mathbf{v} = \mathbf{T} \mathbf{f}.$$

⁶ In reality, it is not clear what Clenshaw used, as the relationship (2.4) in Clenshaw (1957) divides by zero for the integral of $T_1(x)$.

Cancelling out the (formally invertible) quasimatrix \mathbf{T} we get an infinite-dimensional linear system:

$$(I - XQ^2)\mathbf{v} = \mathbf{f}.$$

Note that this system is a compact perturbation of the identity, and so is readily solvable via, say, finite section.

A missing component so far is boundary conditions: we can find a particular solution using the above linear system, but not necessarily the solution we were after, and indeed with the standard Airy equation where $f = 0$, we recover the rather uninteresting solution $\mathbf{v} = 0$. Clenshaw, therefore, proposed a shooting method: initiate the recurrence from a large N using two different final conditions and iterate backwards, then combine with the constants of integration to match the prescribed boundary conditions.

In linear algebra terms, this is equivalent to appending additional unknowns corresponding to the low-order Chebyshev coefficients, and additional rows that impose the boundary conditions. That is, we write

$$u(x) = (T_0 \mid T_1 \mid \mathbf{T}Q^2)\mathbf{v}.$$

So the Airy equation with boundary conditions becomes

$$\begin{pmatrix} 1 & -1 & (1, -1, 1, \dots)Q^2 \\ 1 & 1 & (1, 1, 1, \dots)Q^2 \\ -X_{:,0} & -X_{:,1} & I - XQ^2 \end{pmatrix} \begin{pmatrix} c_0 \\ c_1 \\ \mathbf{v} \end{pmatrix} = \begin{pmatrix} a \\ b \\ \mathbf{f} \end{pmatrix},$$

where $X_{:,j}$ is the j th column of the multiplication operator X .

Integral reformulation is straightforward to construct and leads to sparse, almost banded systems; see Figure 5.1 for the sparsity structure for the Airy equation (which is the same as the ultraspherical spectral method). This system can be efficiently solved, using either shooting techniques, as in Clenshaw (1957), or linear algebra techniques such as adaptive QR of Olver and Townsend (2013). It is generalizable to higher-order differential equations. However, it will not immediately work with Sturm–Liouville-type equations: in an operator like $\mathcal{D}a(x)\mathcal{D}$ we cannot remove the leftmost derivative by multiplying on the right by integration operators without reformulating using integration by parts. Finally, it is not easily generalizable to higher-dimensional domains such as triangles.

Remark 5.1. Note that Clenshaw did not truncate X and Q^2 and then multiply: because both are banded operators, it is possible to deduce the entries directly. We follow this approach: we have access to the true infinite-dimensional operators, which limits the source of errors.

5.2. Ultraspherical spectral method

The ultraspherical spectral method of Olver and Townsend (2013) is an alternative approach for arriving at sparse discretizations, building on the observation that the domain and range of the equation need not be in the same basis. This gives more flexibility, and indeed enabled extensions to singular integral equations and partial differential equations outlined below.

We expand in Chebyshev basis $u = \mathbf{T}u$. By changing continuous operators to discrete ones we arrive at

$$\begin{aligned} (\mathcal{D}^2 - x)\mathbf{T} &= \mathcal{D}\mathbf{U}D_T^U + \mathbf{T}X \\ &= \mathbf{C}^{(2)}D_U^{(2)}D_T^U - \mathbf{U}R_T^UX_T \\ &= \mathbf{C}^{(2)}(D_T^{(2)} - R_T^{(2)}X_T). \end{aligned}$$

Thus the Airy equation becomes

$$\begin{pmatrix} \mathbf{T}(-1) \\ \mathbf{T}(1) \\ \mathbf{C}^{(2)}(D_T^{(2)} - R_T^{(2)}X) \end{pmatrix} \mathbf{u} = \begin{pmatrix} a \\ b \\ \mathbf{C}^{(2)}R_T^{(2)}\mathbf{f} \end{pmatrix}.$$

We can then cancel out the (formally invertible) quasimatrix $\mathbf{C}^{(2)}$ to arrive at a discrete system:

$$\begin{pmatrix} \mathbf{T}(-1) \\ \mathbf{T}(1) \\ D_T^{(2)} + R_T^{(2)}X \end{pmatrix} \mathbf{u} = \begin{pmatrix} a \\ b \\ R_T^{(2)}\mathbf{f} \end{pmatrix}.$$

Note that X is tridiagonal, $R_T^{(2)}$ has bandwidths $(0,4)$ and $D_T^{(2)}$ has bandwidths $(-2,2)$, and as the product of banded matrices adds the bandwidths, the operator is banded with bandwidths $(1,5)$. Thus the resulting system with the boundary is an *almost banded* linear system: it can be written as

$$\text{banded} + \text{triu}(UV^\top),$$

where UV^\top is a rank-2 matrix and the banded part has bandwidth $(3,3)$. A spy plot is shown in Figure 5.1. This structure is preserved under linear algebra operations such as Householder reflections, with the bandwidth of the banded part uniformly bounded. Thus it is amenable to fast linear algebra.

What makes the ultraspherical spectral method is its versatility. We can handle arbitrary functional constraints: Neumann conditions, Robin conditions, integral constraints, *etc.* For example, one can convert any linear ordinary differential operator with variable coefficients such as

$$a_\nu(x) \frac{d^\nu}{dx^\nu} + \cdots + a_2(x) \frac{d^2}{dx^2} + a_1(x) \frac{d}{dx} + a_0(x)$$

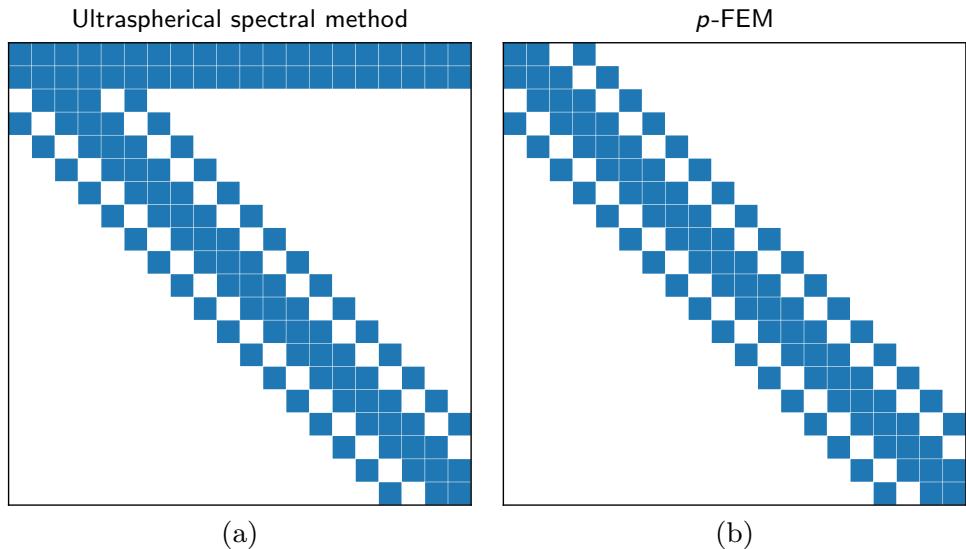


Figure 5.1. A spy plot of the sparsity structure for Clenshaw’s integral reformulation and the ultraspherical spectral method (a) compared to the p -finite element method (b), for a 20×20 discretization for the Airy operator $\mathcal{D}^2 - x$.

to an (infinite) banded matrix,

$$a_\nu(X_{(\nu)})D_T^{(\nu)} + \cdots + R_{(2)}^{(\nu)}a_2(X_{(2)})D_T^{(2)} + R_U^{(\nu)}a_1(X_U)D_T^U + R_T^{(\nu)}.$$

We do not require the equation to be given in reduced form: for example, a Sturm–Liouville operator

$$a(x)\frac{d}{dx}b(x)\frac{d}{dx} + c(x)$$

also leads to a banded matrix

$$a(X_{(2)})D_U^{(2)}b(X_U)D_T^U + R_T^{(2)}c(X_T)$$

without having to use the product rule to simplify the differential equations. This is a fairly minor issue for second-order equations but can become significant for higher-order differential equations.

We can also handle general boundary conditions. For example, an arbitrary differential condition at a point x_0 of the form

$$\beta_\mu u^{(\mu)}(x_0) + \cdots + \beta_1 u'(x_0) + \beta_0 u(x_0)$$

is represented by the row-vector

$$\mathbf{C}^{(\mu)}(x_0) [\beta_\mu D_T^{(\mu)} + \cdots + \beta_1 R_U^{(\mu)} D_T^U + \beta_0 R_T^{(\mu)}]$$

and an integral constraint like $\int_{-1}^1 u(x) dx$ is represented by

$$(\mathbf{T}(1) - \mathbf{T}(-1)) Q_T.$$

5.2.1. Example: Airy equation

As a simple example, we consider the Airy equation

$$\epsilon^2 u'' = xu \quad u(-1) = 1, u(1) = 0,$$

which is discretized as

$$\begin{pmatrix} \mathbf{T}(-1) \\ \mathbf{T}(1) \\ \epsilon^2 D_T^{(2)} - R_T^{(2)} X \end{pmatrix} \mathbf{u} = \begin{pmatrix} 1 \\ 0 \\ \mathbf{0} \end{pmatrix}.$$

In Figure 5.2(a) we plot the computed coefficients for four choices of ϵ , highlighting the need for fast algorithms. Note that the calculated coefficients converge well below machine precision. The fact that the tail is below machine precision is a consequence of the diagonal dominance of the discretization (it can be shown to converge in any space of algebraically decaying coefficients (Olver and Townsend 2013)). This makes the approach particularly appealing for adaptivity, as we can robustly choose the convergence criteria. Here we calculated the coefficients until they underflowed. In Figure 5.2(b) we show the timing for various ϵ using the adaptive criteria. This demonstrates the $O(N)$ complexity, as the N required to resolve the solution scales linearly with ϵ^{-1} .

5.2.2. Example: catastrophes and integrals with coalescing saddles

Integrals with coalescing saddles arise in a number applications such as modelling tsunamis (Berry 2007), and are defined by (Olver *et al.* 2010, 36.2.4)

$$\Psi_K(\mathbf{x}) := \int_{-\infty}^{\infty} \exp i\Phi_K(t; \mathbf{x}) dt, \quad \Phi_K(t; \mathbf{x}) := t^{K+2} + \sum_{m=1}^K x_m t^m.$$

We focus our discussion on $K = 2$, which is an integral due to Pearcey (1946), and given by

$$\Psi(x, y) = \Psi_2(x, y) = \int_{-\infty}^{\infty} \exp i(t^4 + yt^2 + xt) dt.$$

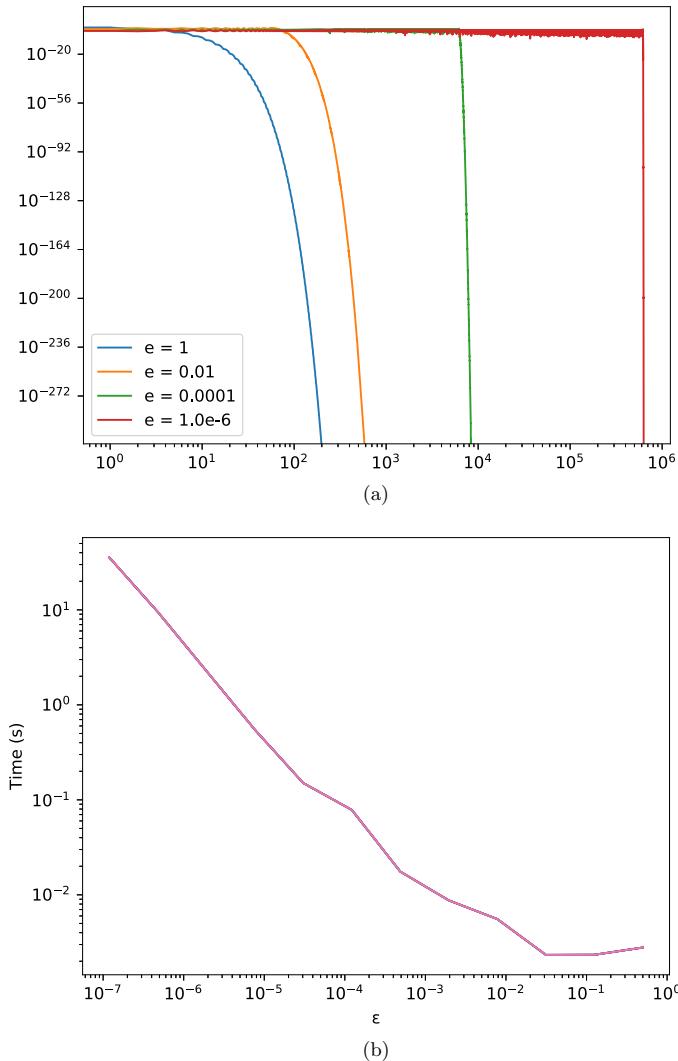


Figure 5.2. The computed coefficients for the Airy equation (a) and the time taken to solve the Airy equation using the ultraspherical spectral methods with the adaptive QR method (b). The degrees of freedom needed to resolve the solution are proportional to ϵ^{-1} .

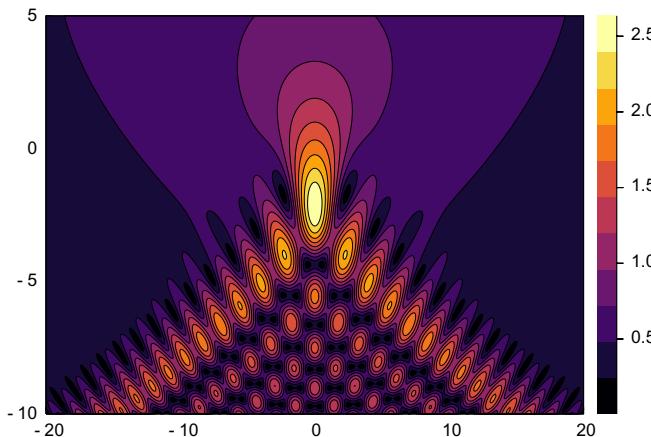


Figure 5.3. A depiction of the Pearcey integral, computed using the ultraspherical spectral method applied to a third-order differential equation in x for each y value.

Catastrophic integrals satisfy differential equations with respect to $x_1 \equiv x$ of order $K+1$; in the case of the Pearcey integral we have (Olver *et al.* 2010, 32.10.4)

$$\frac{d^3\Psi}{dx^3} - \frac{y}{2} \frac{d\Psi}{dx} - \frac{ix}{4} \Psi = 0.$$

The ultraspherical spectral method gives a convenient way of discretizing this equation: if we write

$$\Psi(x,y) = \mathbf{T}(x)\mathbf{u}(y),$$

we have for each choice of y

$$\left(D_T^{(3)} - \frac{y}{2} R_U^{(3)} D_T^U - \frac{i}{4} R_T^{(3)} X_T \right) \mathbf{u}(y) = 0.$$

This is a third-order equation; however, we can evaluate $\Psi(0,y)$ using a convergent series expansion (Olver *et al.* 2010, 36.8.1) and asymptotically accurate expansions for $\Psi(\pm b,y)$ where b is large, using Olver *et al.* (2010, 36.11.2), thereby imposing three conditions. Thus each slice of y can be computed in $O(N)$ operations. Putting everything together, we display the solution in Figure 5.3.

5.3. *p*-finite element methods

Finite element methods (FEM) have the attractive feature that they preserve self-adjointness of problems, and require only a weak formulation of a differential equation. The idea behind *p*-FEM as used by Beuchler and

Schöberl (2006) is to use orthogonal polynomials to enable high polynomial order approximations for the test-and-trial space, maintaining sparsity in the process. This is readily explained in quasimatrix notation. Take for example the weak formulation of the Airy equation with Dirichlet conditions:

$$-\langle v', u' \rangle - \langle v, xu \rangle = \langle v, f \rangle.$$

Write u and v in weighted Jacobi expansions and f in Legendre expansion:

$$u = (1-x^2)\mathbf{C}^{(3/2)}\mathbf{u}, \quad v = (1-x^2)\mathbf{C}^{(3/2)}\mathbf{v}, \quad f = \mathbf{P}\mathbf{f}.$$

Then we have

$$\begin{aligned} \langle v, xu \rangle &= \mathbf{v}^\top (\mathbf{C}^{(3/2)})^\top (1-x^2)x(1-x^2)\mathbf{C}^{(3/2)}\mathbf{u} \\ &= \mathbf{v}^\top (L_{(3/2)}^P)^\top \mathbf{P}^\top \mathbf{P} X_P L_{(3/2)}^P \mathbf{u} \end{aligned}$$

and similarly

$$\begin{aligned} \langle v', u' \rangle &= \mathbf{v}^\top (\mathcal{D}(1-x^2)\mathbf{C}^{(3/2)})^\top \mathcal{D}(1-x^2)\mathbf{C}^{(3/2)}\mathbf{u} \\ &= \mathbf{v}^\top (W_{(3/2)}^P)^\top \mathbf{P}^\top \mathbf{P} W_{(3/2)}^P \mathbf{u}. \end{aligned}$$

Thus the differential equation becomes

$$(-(W_{(3/2)}^P)^\top M_P W_{(3/2)}^P - (L_{(3/2)}^P)^{(3/2)} M_P X_P L_{(3/2)}^P) \mathbf{u} = (L_{(3/2)}^P)^\top M \mathbf{f},$$

where $M_P = \mathbf{P}^\top \mathbf{P}$ is the diagonal mass matrix.

The system is symmetric: since $\mathbf{Q} := \mathbf{P} M_P^{-1/2}$ are the orthonormalized Legendre polynomials, we have

$$x \mathbf{P} M_P^{-1/2} = \mathbf{P} X_P M_P^{-1/2} = \mathbf{P} M_P^{-1/2} \tilde{X} = \mathbf{Q} \tilde{X}$$

so that $M_P X_P = M_P^{1/2} \tilde{X} M_P^{1/2}$ is symmetric. Thus square truncations are banded symmetric matrices. For positive definite problems we can therefore reliably use Cholesky decompositions, as implemented in LAPACK.

5.3.1. Example: a stochastic Airy equation

We consider a stochastic Airy equation⁷ with Dirichlet conditions

$$\epsilon u'' - xu = \frac{dW}{dx}, \tag{5.1}$$

where $W(x)$ is a Wiener process. The weak formulation of this equation is

$$-\epsilon \langle v', u' \rangle - \langle v, xu \rangle = -\langle v', W \rangle,$$

⁷ This should not be confused with the stochastic Airy operator $\epsilon u'' - xu + (dW/dx)u$, which has recently seen significant interest in the random matrix community (Krishnapur, Rider and Virág 2016, Edelman and Sutton 2007).

which results in a banded discretization as above. This problem requires very large N to solve: as ϵ becomes small it becomes increasingly oscillatory, and we also need large N to approximate the Wiener process. We use the QR factorization (as opposed to the Cholesky factorization) as the Airy operator is not positive definite.

In Figure 5.4 we compute the solutions to this equation. We include the timings, which consist of (1) calculating the Legendre coefficients of W using transforms, (2) building the operator, (3) computing the QR factorization, and (4) solving the equation using the factorization. This results in a quasi-optimal complexity algorithm.

5.3.2. Example: beam equation

We now consider the clamped beam equation

$$\begin{aligned} u_{tt} &= u_{xxxx}, \\ u(0,x) &= u_0(x), u_t(0,x) = v_0(x), \\ u(t,-1) &= u_x(t,-1) = u(t,1) = u_x(t,1) = 0, \end{aligned}$$

where we use the p -FEM to discretize in space. We solve the weak formulation

$$\langle v, u_t \rangle = \langle \Delta v, \Delta u \rangle$$

using the space $u(t,x) = (1-x^2)\mathbf{P}^{(2,2)}\mathbf{u}(t)$ to arrive at the discretization

$$\underbrace{(L_{(2,2)}^{(0,0)})^\top M_P L_{(2,2)}^{(0,0)}}_M \mathbf{u}_t = \underbrace{(W_{(2,2)}^{(0,0)})^\top M_P W_{(2,2)}^{(0,0)}}_A \mathbf{u}.$$

This is a standard ODE with a mass matrix, where M is symmetric with bandwidths (4,4) while A is diagonal.

In Figure 5.5 we demonstrate the solution of the beam equation. Since M and A are banded, we can perform implicit time-steps in $O(N)$ operations, where we truncate the discretization to the first $N = 1000$ terms in x . We use the Rodas 5 time-stepping routine with time-step 0.01. As M is symmetric positive definite, we can use the banded Cholesky factorization.

5.4. Domain decomposition

The p -finite element method naturally extends to multiple elements. For example, suppose we have a mesh $x_1 < x_2 < \dots < x_M$. Within each element we can use the *bubble functions*

$$\begin{cases} (1-m_j(x)^2)\mathbf{C}^{(3/2)}(m_j(x)) & x_j \leq x \leq x_{j+1}, \\ 0 & \text{otherwise,} \end{cases}$$

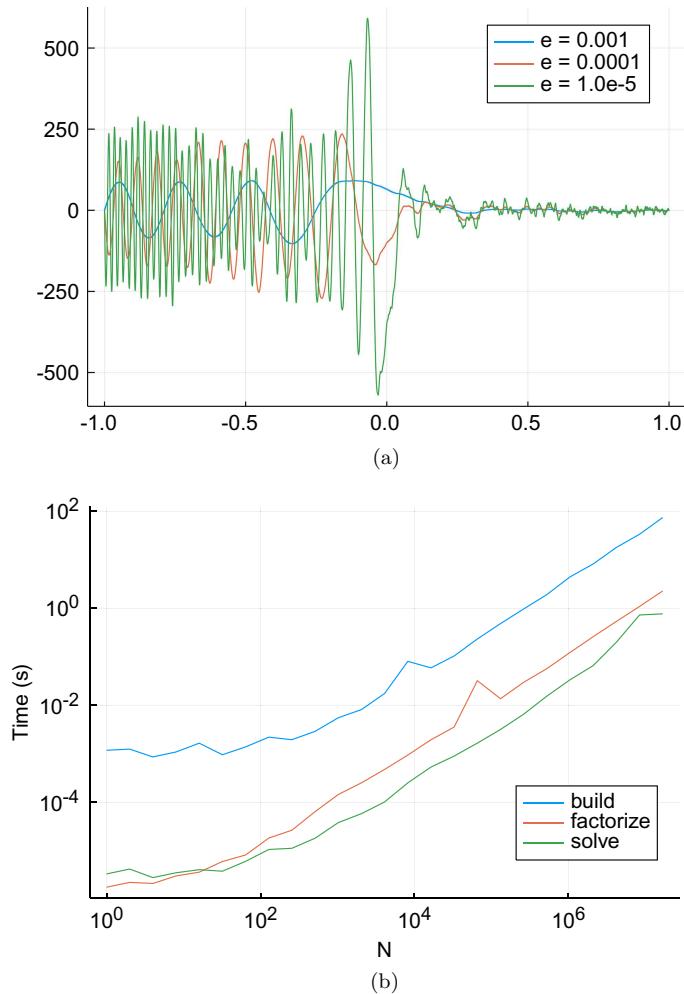


Figure 5.4. (a) Solutions to the stochastic Airy equation (5.1) with three choices of ϵ , computed using the p -FEM. (b) The timings in seconds for building, computing the QR factorization, and inverting using the pre-computed QR factorization.

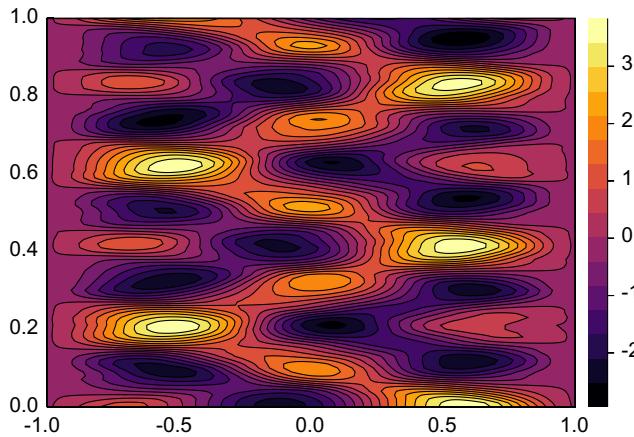


Figure 5.5. A solution of the beam equation using the p -finite element method. We use $N = 1000$ spatial discretization with a time-step of 0.01 with Rodas 5.

where $m_j: [x_j, x_{j+1}] \rightarrow [-1, 1]$ is the affine map. This is then combined with the low-order finite element basis

$$\begin{cases} \frac{x - x_{j-1}}{x_j - x_{j-1}} & x_{j-1} \leq x \leq x_j, \\ \frac{x_{j+1} - x}{x_{j+1} - x_j} & x_j \leq x \leq x_{j+1}. \end{cases}$$

As both the bubble functions and the additional low-order bases can both be transformed to piecewise Legendre polynomials, we can systematically deduce the resulting discretization using recurrence relationships. This discretization is extremely sparse, with low-rank communication between elements, due to the low-order terms. Interlacing the coefficients reduces this to a banded linear system whose bandwidth is dictated by the number of elements.

Ultraspherical spectral methods can also be incorporated into domain decomposition methods, where continuity conditions are enforced as boundary conditions. Alternatively, one may use integral reformulation, which Lee and Greengard (1997) used to construct a very effective adaptive mesh refinement strategy.

5.5. Eigenvalue problems

The p -FEM and ultraspherical spectral method can be used to reduce differential eigenvalue problems to generalized eigenvalue problems. We demonstrate this on the Schrödinger equation:

$$-u'' + |x|u = \lambda u, \quad -20 < x < 20$$

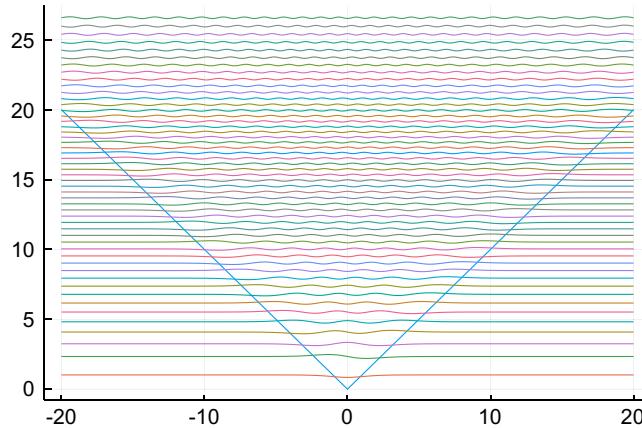


Figure 5.6. The first 50 eigenstates of the Schrödinger equation with potential $V(x) = |x|$ computed using the p -FEM, with two elements.

with Dirichlet conditions. The eigenstates are plotted in Figure 5.6. Since the potential is piecewise polynomial, the discretization is symmetric and banded, and the eigenvalues can be calculated in $O(N^2)$ complexity using a specialized solver (Schwarz 1968, Crawford 1973, Kaufman 1993, Kaufman 2000). Here we decompose the interval into two elements.

As an alternative to the Beucher–Schöberl p -FEM, Aurentz and Slevinsky (2020) describe how basis recombination (Shen 1994, Doha and Abd-Elhameed 2002, Julien and Watson 2009, Livermore 2010) can be leveraged to symmetrize the ultraspherical spectral method for self-adjoint problems while preserving the banded property that we have come to expect from spectral discretizations of linear problems. This method allows for more general self-adjoint boundary conditions, does not require the problem to be posed in the weak or self-adjoint form for the symmetry to transcend from problem to discretization, and may solve weakly defined problems by specifying relative discontinuity conditions in the basis.

In the self-adjoint setting, a (2ν) th-order linear differential eigenvalue problem with self-adjoint boundary conditions takes the form

$$\mathcal{L}u = [(-\mathcal{D})^\nu(a_\nu \mathcal{D}^\nu) + (-\mathcal{D})^{\nu-1}(a_{\nu-1} \mathcal{D}^{\nu-1}) + \cdots + a_0]u = \lambda w u, \quad (5.2)$$

$$\mathcal{B}u = 0, \quad (5.3)$$

where the usual conditions on the coefficients a_i and w are required for a sensible solution; at the very least, $a_\nu \neq 0$, and $w \geq 0$ and $w = 0$ only on a set of Lebesgue measure zero. In this form, \mathcal{L} is a self-adjoint linear differential operator on a domain $D \subset \mathbb{R}$ in the Hilbert space $H = L^2(D, w(x) dx)$ provided it is accompanied by $\mathcal{B}: H \rightarrow \mathbb{C}^{2\nu}$, a set of 2ν self-adjoint boundary conditions. We will also assume that all the variable

coefficients are at most degree- m polynomials, that is, for general smooth variable coefficients we initially approximate them by polynomials.

Given the quotient Hilbert space $H_{\mathcal{B}} := \{u \in H : \mathcal{B}u = 0\}$, a classical approach calls for a Ritz–Galerkin discretization for the trial and test functions taken from an n -dimensional subspace $V_n \subset H_{\mathcal{B}}$. Indeed, when $V_n = \mathbb{P}_{n-1}$ and an orthonormal polynomial basis satisfying the boundary conditions is utilized, the discrete realization of $\langle u, \mathcal{L}v \rangle$ is automatically symmetric (Hermitian). Unfortunately, the resulting discretizations are dense. Thus data-sparsity in problem formulations, such as (5.2) with finite-degree polynomial coefficients, is ignored when transformed into a computational problem, resulting in $O(n^3)$ complexity algorithms (Golub and Loan 2013) for approximate eigenvalues and eigenfunctions.

Let $\{\phi_n\}_{n=0}^\infty$ be an orthonormal polynomial basis for H with $\deg(\phi_n) = n$. For linearly independent boundary conditions in \mathcal{B} , let $\{\rho_n\}_{n=0}^\infty$ be the recombinations of the orthonormal basis $\{\phi_n\}_{n=0}^\infty$ that annihilate \mathcal{B} , guaranteeing $\mathcal{B}\rho_n = 0$. This assumption suggests the existence of an infinite-dimensional lower-triangular and banded conversion such that

$$\boldsymbol{\rho} = \boldsymbol{\phi}A. \quad (5.4)$$

For a problem on $[-1,1]$ with Dirichlet boundary conditions, the recombinations are the same basis used in Section 5.3, up to normalization. With the conversion in hand, a banded QR factorization reveals the unitary operator mapping the orthonormal polynomial basis for H to the orthonormal polynomial basis $\{\psi_n\}_{n=0}^\infty$ for the quotient Hilbert space $H_{\mathcal{B}}$, $\boldsymbol{\psi} = \boldsymbol{\phi}Q$, and the auxiliary resulting conversion, $\boldsymbol{\sigma} = \boldsymbol{\psi}R$. Finally, define the polynomials $\{\sigma_n\}_{n=0}^\infty$ by

$$\boldsymbol{\sigma}R^* = \boldsymbol{\psi}. \quad (5.5)$$

Since we have $\boldsymbol{\psi}Q^* = \boldsymbol{\phi}$, we can gather all three adjoint conversions together to obtain

$$\boldsymbol{\sigma}A^* = \boldsymbol{\phi}. \quad (5.6)$$

Aurentz and Slevinsky (2020) used these four bases to prove that the Petrov–Galerkin scheme for (5.2) is symmetric definite and banded, where the solution is formally represented by $u(x) = \boldsymbol{\rho}\mathbf{v} = \boldsymbol{\psi}R\mathbf{v} = \boldsymbol{\psi}\mathbf{u} = \boldsymbol{\phi}A\mathbf{v} = \boldsymbol{\phi}\mathbf{w}$ and the residuals are estimated by $\boldsymbol{\sigma}$.

5.5.1. The Petrov–Galerkin scheme is self-adjoint

It follows by classical arguments that the Ritz–Galerkin scheme is self-adjoint:

$$\mathcal{L}\boldsymbol{\psi}\mathbf{u} = \boldsymbol{\psi}L\mathbf{u} = \lambda w \boldsymbol{\psi}\mathbf{u} = \lambda \boldsymbol{\psi}M\mathbf{u},$$

where $L = L^* \in \mathbb{C}^{\infty \times \infty}$ and $M = M^* \in \mathbb{C}^{\infty \times \infty}$ is positive definite. If we use $\mathbf{u} = R\mathbf{v}$ and (5.5) to represent $\boldsymbol{\psi}$ in terms of $\boldsymbol{\sigma}$, the symmetry may be

imported into the Petrov–Galerkin scheme

$$\boldsymbol{\sigma}R^*LR\mathbf{v} = \lambda\boldsymbol{\sigma}R^*MR\mathbf{v}.$$

Both R^*LR and R^*MR are self-adjoint and the latter is also positive definite.

5.5.2. The Petrov–Galerkin scheme is banded

Considering the Petrov–Galerkin scheme

$$\mathcal{L}\boldsymbol{\rho}\mathbf{v} = \lambda w\boldsymbol{\rho}\mathbf{v},$$

the combination of (5.4) and (5.6) informs us that $w\boldsymbol{\rho} = w\phi A = \phi M_B A = \boldsymbol{\sigma}A^*M_B A$, where $M_B = M_B^*$ is also positive definite and banded by executing the Clenshaw–Smith algorithm with the orthonormal polynomials’ Jacobi operator. Similarly, $\mathcal{L}\boldsymbol{\rho} = \mathcal{L}\phi A$, and by the ultraspherical spectral method (Olver and Townsend 2013), there exists a banded discretization of \mathcal{L} , namely L_B , mapping normalized Legendre polynomials to normalized Jacobi polynomials with integer parameters in proportion to the highest derivatives. Applying the inverse banded conversion operator, and further utilizing $\phi = \boldsymbol{\sigma}A^*$, we arrive at

$$\boldsymbol{\sigma}A^*C^{-1}L_B A\mathbf{v} = \lambda\boldsymbol{\sigma}A^*M_B A\mathbf{v}. \quad (5.7)$$

At first glance, this formula does not appear to guarantee a banded discretization, due to the appearance of the inverse of a banded matrix. To show that the discretization is indeed banded, we note that the conversion matrix C is upper-triangular and banded, which means that its inverse will *not* extend the bandwidth of $L_B A$ any lower. Since A^* is also upper-triangular and since the Petrov–Galerkin scheme is self-adjoint, we conclude that it must be banded above as well.

5.5.3. A sixth-order generalized eigenvalue problem

Aurentz and Slevinsky (2020) show through several numerical and theoretical applications that the method is general, versatile and easy to use. They show that piecewise defined problems on bounded and unbounded domains may be stitched together with continuity conditions. They also show that the method applies to higher-order problems with ease.

The following sixth-order generalized eigenvalue problem of Gutierrez and Laura (1995) models the vibration of non-uniform rings:

$$\begin{aligned} & \{(-\mathcal{D})^3[g\mathcal{D}^3] - (-\mathcal{D})^2[2\pi^2g\mathcal{D}^2] + (-\mathcal{D})[(\pi^4g + \pi^2g'')\mathcal{D}]\}u \\ &= \lambda\{(-\mathcal{D})[\pi^4f\mathcal{D}] + \pi^6f\}u, \end{aligned} \quad (5.8)$$

where $f(x) = 1 + 4(1 - r)[(x - 1/2)^2 - 1/4]$, $r \in (0, \infty)$ is a thickness parameter, $g = f^3$, and the boundary conditions are $u = u' = u''' = 0$ at

Table 5.1. The first four computed eigenvalues of (5.8) with $r = 1$ in 256-bit extended precision, computed using the method of Aurentz and Slevinsky (2020) and rounded to fit this page.

n	λ_n
0	5.138 119 642 550 555 649 415 412 286 723 460 005 700 131 ...
1	47.932 044 726 525 070 683 741 246 754 121 353 113 209 603 ...
2	195.375 225 987 311 298 618 562 101 588 390 534 216 626 466 ...
3	520.732 437 610 822 993 012 167 361 990 447 367 561 584 024 ...

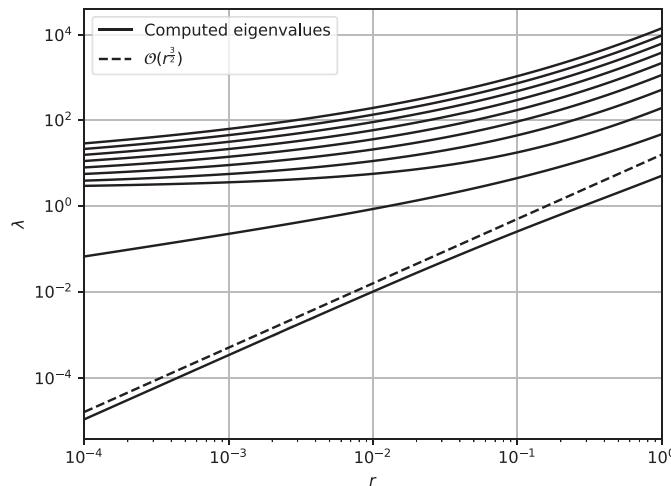


Figure 5.7. The first ten generalized eigenvalues of (5.8) as functions of the thickness parameter r . Figure reproduced from Aurentz and Slevinsky (2020) with permission. Copyright © 2020 Elsevier.

both $x = 0$ and $x = 1$. Due to the third-order derivatives in the boundary conditions, they may not be implemented by multiplying an orthogonal polynomial expansion by a polynomial weight, and thus require basis recombination.

With sixth-order derivatives in the first operator in (5.8), this is by all accounts an ill-conditioned generalized eigenvalue problem. And as $r \rightarrow 0$, this ill-conditioning is compounded by the fact that both linear operators become singular as the smallest generalized eigenvalue tends to zero.

Using the general Rayleigh quotient iteration (Rayleigh 1937, Parlett 1974) in 256-bit extended precision, the first four computed eigenvalues are tabulated in Table 5.1. Figure 5.7 tracks the first ten eigenvalues as $r \rightarrow 0$.

6. Partial differential equations

Multivariate OPs on rectangles/hypercubes and triangles/tetrahedra have played a central role in spectral element methods for partial differential equations. For triangles, see Dubiner (1991) and Karniadakis and Sherwin (2013) for a comprehensive overview. These methods result in dense matrices for large p .

Recently, approaches for sparse discretizations for large p have been introduced in the p -finite element method on triangles (Beuchler and Schöberl 2006, Li and Shen 2010) and the ultraspherical spectral method on rectangles (Townsend and Olver 2015), disks (Vasil *et al.* 2016) and triangles (Olver *et al.* 2019).

Our model problem is the multivariate analogue of the Airy equation: a variable coefficient inhomogeneous Helmholtz equation

$$\Delta u + k^2 xy^2 u = f(x, y).$$

We will consider both squares and triangles in two dimensions. The type of boundary conditions allowed depends on the method under discussion.

6.1. Ultraspherical spectral methods on rectangles

On rectangular domains it is natural to use tensor product bases, and thus we consider the expansion of the form

$$u(x, y) = \mathbf{T}^2(x, y)U = \mathbf{T}(x)U\mathbf{T}(y)^\top,$$

where U is an infinite matrix of coefficients and $\mathbf{T}^2 = \mathbf{T} \otimes \mathbf{T}$. The separable nature of the decomposition makes it straightforward to construct partial differential operators from ordinary differential operators, for example

$$\frac{\partial u}{\partial x} = \mathbf{U}(x)D_T^U U\mathbf{T}(y)^\top = (\mathbf{U} \otimes \mathbf{T})(x, y)(D_T^U \otimes I)U.$$

Thus we can immediately construct the partial derivatives

$$\begin{aligned}\partial_x^2 \mathbf{T}^2 &= (\mathbf{C}^{(2)} D_T^{(2)}) \otimes \mathbf{T} = (\mathbf{C}^{(2)} \otimes \mathbf{T})(D_T^{(2)} \otimes I) = \mathbf{C}^{(2,2)}(D_T^{(2)} \otimes R_T^{(2)}), \\ \partial_y^2 \mathbf{T}^2 &= (\mathbf{T} \otimes \mathbf{C}^{(2)})(I \otimes D_T^{(2)}) = \mathbf{C}^{(2,2)}(R_T^{(2)} \otimes D_T^{(2)}),\end{aligned}$$

where $\mathbf{C}^{(a,b)} = \mathbf{C}^{(a)} \otimes \mathbf{C}^{(b)}$. Thus we have

$$\Delta \mathbf{T}^2 = \mathbf{C}^{(2,2)} \underbrace{[D_T^{(2)} \otimes R_T^{(2)} + R_T^{(2)} \otimes D_T^{(2)}]}_{\Delta_{T,T}^{(2,2)}}.$$

Similarly, we have

$$x\mathbf{T} = \mathbf{C}^{(2,2)} \underbrace{[(R_T^{(2)} X_T) \otimes I]}_X \quad \text{and} \quad y\mathbf{T} = \mathbf{C}^{(2,2)} \underbrace{[I \otimes (R_T^{(2)} X_T)]}_Y.$$

Putting everything together, we have

$$[\Delta + xy^2]\mathbf{T}^2 = \mathbf{C}^{(2,2)} [\Delta_{T,T}^{(2,2)} + R_{T,T}^{(2,2)} XY^2],$$

where $R_{T,T}^{(2,2)} = R_T^{(2)} \otimes R_T^{(2)}$. This is a sparse discretization.

We can similarly specify boundary conditions which are particularly nice for separable boundary conditions. For example, a Neumann condition on the top of an interval can be written as

$$u_y(x, 1) = (\mathbf{T} \otimes \mathbf{U}(1))(I \otimes D_T^U)U = \mathbf{T}(x)U(D_T^U)^\top \mathbf{U}(1)^\top.$$

Putting everything together, if we write the right-hand side as

$$f(x, y) = \mathbf{C}^{(2,2)} F$$

and the boundary data given in Chebyshev expansions as

$$\begin{aligned} u(x, -1) &= g_{-1}(x) = \mathbf{T}\mathbf{g}_{-1}, \\ u(x, 1) &= g_1(x) = \mathbf{T}\mathbf{g}_1, \\ u(-1, y) &= h_{-1}(y) = \mathbf{T}\mathbf{h}_{-1}, \\ u(1, y) &= h_1(y) = \mathbf{T}\mathbf{h}_1, \end{aligned}$$

then we can express the variable coefficient Helmholtz problem with Dirichlet conditions as

$$\begin{aligned} (\mathbf{T}(-1) \otimes I)U &= \mathbf{g}_{-1}, \\ (\mathbf{T}(1) \otimes I)U &= \mathbf{g}_1, \\ (I \otimes \mathbf{T}(-1))U &= \mathbf{h}_{-1}, \\ (I \otimes \mathbf{T}(1))U &= \mathbf{h}_1, \\ (\Delta_{T,T}^{(2,2)} + R_{T,T}^{(2,2)} XY^2)U &= F. \end{aligned}$$

Alternatively, we can view this as the following matrix problem:

$$\begin{aligned} \mathbf{T}(-1)U &= \mathbf{g}_{-1}, \\ \mathbf{T}(1)U &= \mathbf{g}_1, \\ U\mathbf{T}(-1)^\top &= \mathbf{h}_{-1}^\top, \\ U\mathbf{T}(1)^\top &= \mathbf{h}_1^\top, \\ D_T^{(2)}U(R_T^{(2)})^\top + R_T^{(2)}U(D_T^{(2)})^\top + R_T^{(2)}X_TU(X_T^2)^\top(R_T^{(2)})^\top &= F. \end{aligned}$$

This can be converted into a sparse linear system by choosing an ordering for the coefficients of U .

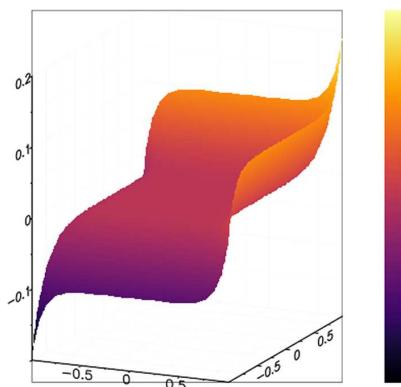


Figure 6.1. The solution to the screened Poisson equation using the ultraspherical spectral method.

6.2. Sylvester equations

Townsend and Olver (2015) observed that the boundary conditions in the matrix formulation of the problem can be used to eliminate degrees of freedom in the differential equation, thereby reducing the equation to a generalized Sylvester equation, that is, we wish to solve an equation of the form

$$A_1 U B_1^\top + \cdots + A_r U B_r^\top = F,$$

where all A_j and B_j are sparse matrices, and r is known as the separability rank of the partial differential operator. For the special case of $r = 2$ this can be accomplished using a Schur factorization: find Q and Z orthogonal so that $A_1 = QR_1Z$ and $A_2 = QR_2Z$, thus reducing the problem to

$$R_1 V B_1^\top + R_2 V B_2^\top = Q^\top F$$

so that $V = ZU$. This can be solved via back-substitutions as a sequence of sparse linear solves in $O(N^{3/2})$ complexity for N degrees of freedom.

As an example, consider the ‘good’ Helmholtz equation with Neumann conditions:

$$\Delta u - k^2 u = 0, \quad u_x(\pm 1, y) = u_y(x, \pm 1) = 1.$$

We can write this in rank-2 form as

$$(D_T^{(2)} - k^2 S_T^{(2)}) U (S_T^{(2)})^\top + S_T^{(2)} U (D_T^{(2)})^\top = 0.$$

The solution appears in Figure 6.1.

6.3. Fast Poisson solvers

In the particular situation of Poisson's equation with zero homogeneous Dirichlet conditions, that is,

$$u_{xx} + u_{yy} = f, \quad (x, y) \in [-1, 1]^2, \quad u(\pm 1, \cdot) = u(\cdot, \pm 1) = 0, \quad (6.1)$$

Fortunato and Townsend (2019) take the idea a step further and develop an $O(N \log N \log(1/\epsilon))$ complexity solver. First, it is noted that since (6.1) has homogeneous Dirichlet conditions, the solution can be written as $u(x, y) = (1 - x^2)(1 - y^2)v(x, y)$ for some function $v(x, y)$. So the function $u(x, y)$ is expanded in a carefully selected ultraspherical polynomial basis

$$u(x, y) = (1 - x^2)\tilde{\mathbf{C}}^{(3/2)}(x)U\tilde{\mathbf{C}}^{(3/2)}(y)^\top(1 - y^2), \quad (x, y) \in [-1, 1]^2, \quad (6.2)$$

where $\tilde{\mathbf{C}}^{(3/2)}$ are the orthonormalized ultraspherical polynomials. These OPs are selected because $\tilde{C}_j^{(3/2)}(x)$ is an eigenfunction of the differential operator

$$u \mapsto \frac{d^2}{dx^2}[(1 - x^2)u],$$

that is,

$$\mathcal{D}^2[(1 - x^2)\tilde{\mathbf{C}}^{(3/2)}] = \tilde{\mathbf{C}}^{(3/2)} \underbrace{\text{diag}(-(j+1)(j+2))}_{\Lambda}_{j=0,1,\dots}.$$

This means that (6.1) can be discretized as a generalized Sylvester matrix equation of the form

$$\Lambda U(I - \tilde{X}_{(3/2)}^2) + (I - \tilde{X}_{(3/2)}^2)U\Lambda^\top = F, \quad (6.3)$$

where U is the matrix of expansion coefficients in (6.2), F is the matrix of bivariate $\tilde{\mathbf{C}}^{(3/2)}$ expansion coefficients for f , that is,

$$f(x, y) = \tilde{\mathbf{C}}^{(3/2)}(x)F\tilde{\mathbf{C}}^{(3/2)}(y)^\top,$$

and $\tilde{X}_{(3/2)}$ is the $N \times N$ matrix that represents multiplication by $1 - x^2$ in the $\tilde{\mathbf{C}}^{(3/2)}$ basis (which is symmetric). The Sylvester equation in (6.3) is rearranged and solved by an alternating direction implicit method (Fortunato and Townsend 2019).

6.4. Ultraspherical spectral methods on disks

There are fundamentally two approaches to spectral methods on disks: treat them as separable geometries using polar coordinates, or use non-separable bases. It is possible to construct ultraspherical spectral methods using either approach, with various trade-offs between the techniques. For separable

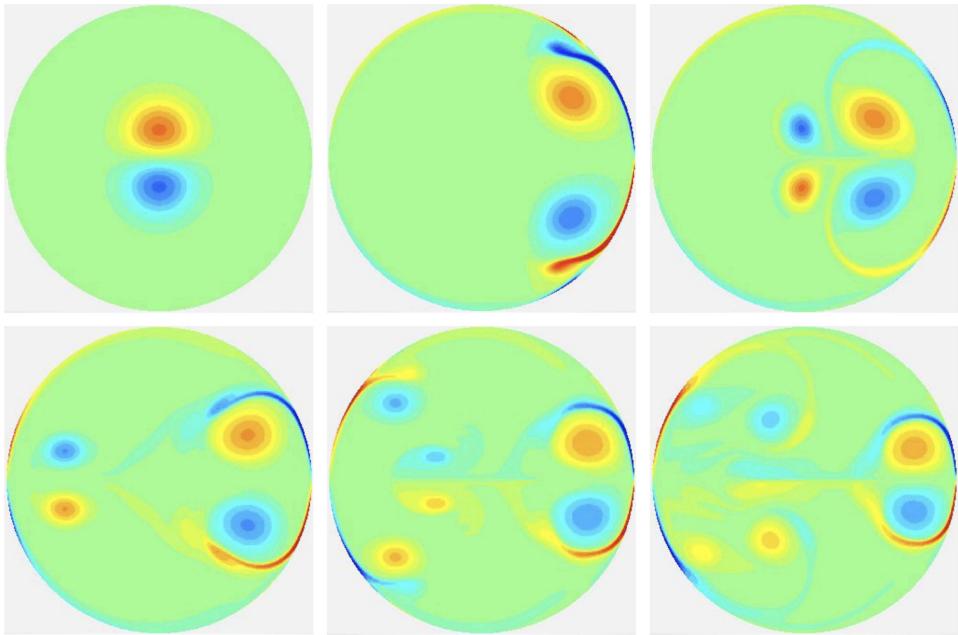


Figure 6.2. Navier–Stokes simulation using spectral methods at time slices $t = 0, 30, 60, 90, 120, 150$.

techniques, see Wilber, Townsend and Wright (2017). Figure 6.2 repeats a simulation from Torres and Coutsias (1999) showing the successful use of spectral methods in solving Navier–Stokes in a disk.

Alternatively, we can use orthogonal polynomials on a disk as in Vasil *et al.* (2016), using generalized Zernike polynomials, which have the benefit of not having any (weak) singularities at the origin: they are polynomials in Cartesian coordinates x and y . Consider a function decomposed into Fourier modes, that is,

$$u(r, \theta) = \sum_{m=-\infty}^{\infty} e^{im\theta} u_m(r).$$

We define two linear operators \mathcal{D}_{\pm} via their action on each Fourier mode:

$$\begin{aligned}\mathcal{D}_+ e^{im\theta} u_m(r) &:= e^{im\theta} \left[\partial_r - \frac{m}{r} \right] u_m = e^{im\theta} r^m \partial_r r^{-m} u_m, \\ \mathcal{D}_- e^{im\theta} u_m(r) &:= e^{im\theta} \left[\partial_r + \frac{m}{r} \right] u_m = e^{im\theta} r^{-m} \partial_r r^m u_m.\end{aligned}$$

For the generalized Zernike polynomials these operators are diagonal with an appropriate change of basis, and hence can be combined to construct the Laplacian and other differential operators arising in applications. Fast transforms between Zernike polynomials and tensor product bases exist (Slevinsky 2017).

6.5. Ultraspherical spectral methods on triangles

Olver *et al.* (2019) introduced the ultraspherical spectral method using orthogonal polynomials on the triangle. As in the disk, we no longer have a tensor product structure without a change of coordinates, but the *sparsity* of discretizations is preserved. For Jacobi polynomials on the triangle $\mathbf{P}^{(a,b,c)}$ there exist sparse matrices

$$D_{x,(a,b,c)}^{(a+1,b,c+1)} \quad \text{and} \quad D_{y,(a,b,c)}^{(a,b+1,c+1)}$$

satisfying (Corollary B.2)

$$\begin{aligned} \partial_x \mathbf{P}^{(a,b,c)} &= \mathbf{P}^{(a+1,b,c+1)} D_{x,(a,b,c)}^{(a+1,b,c+1)}, \\ \partial_y \mathbf{P}^{(a,b,c)} &= \mathbf{P}^{(a,b+1,c+1)} D_{y,(a,b,c)}^{(a,b+1,c+1)}. \end{aligned}$$

Likewise, we have sparse conversion matrices satisfying

$$\mathbf{P}^{(a,b,c)} = \mathbf{P}^{(a+i,b+j,c+k)} R_{(a,b,c)}^{(a+i,b+j,c+k)},$$

for any integers $i,j,k \geq 0$. We can therefore compose these to construct sparse partial derivatives of the form

$$\partial_x \mathbf{P}^{(a,b,c)} = \underbrace{\mathbf{P}^{(a+i+1,b+j,c+k+1)} R_{(a+1,b,c+1)}^{(a+i+1,b+j,c+k+1)}}_{D_{x,(a,b,c)}^{(a+i+1,b+j,c+k+1)}} D_{x,(a,b,c)}^{(a+1,b,c+1)}.$$

In particular, we can construct a sparse representation of the Laplacian via

$$\Delta \mathbf{P}^{(0,0,0)} = \mathbf{P}^{(2,2,2)} \underbrace{(D_{x,(0,0,0)}^{(2,2,2)} + D_{y,(0,0,0)}^{(2,2,2)})}_{\Delta_{(0,0,0)}^{(2,2,2)}}.$$

The sparsity structure is depicted in Figure 6.3(a).

6.5.1. Dirichlet boundary conditions

Boundary conditions are more complicated, so let us consider the case of zero Dirichlet conditions by incorporating it into the basis. That is, we use the *bubble functions* $xyz \mathbf{P}^{(1,1,1)}$, where $z := 1 - x - y$ as in the p -FEM discussed below. Using the weighted differentiation recurrences (Corollary B.3), we

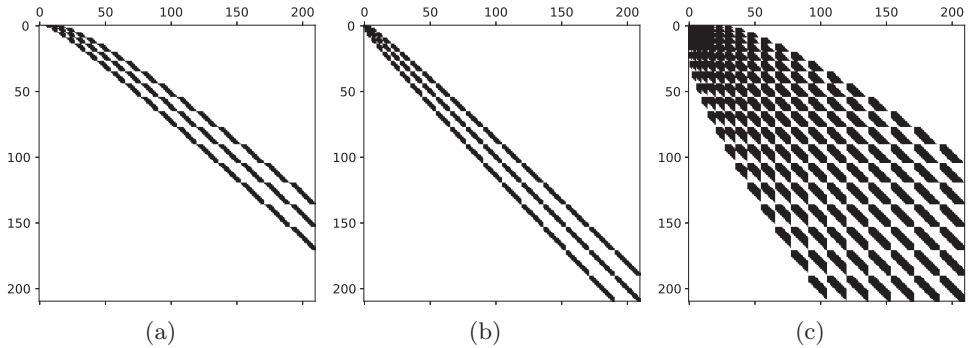


Figure 6.3. The sparsity pattern of the Laplacian $\Delta_{(0,0,0)}^{(2,2,2)}$ (a), the weighted Laplacian Δ_W (b) and the weighted variable coefficient Helmholtz operator $\Delta_W + S_{(0,0,0)}^{(1,1,1)}v(X,Y)L_{(1,1,1)}^{(0,0,0)}$ with $V(x,y) = xy^2$ (c). The weak Laplacian in p -FEM has the same sparsity structure as (b) and (c), with the added benefit of symmetry in the coefficients. Figure reproduced from Olver *et al.* (2019) with permission. Copyright © 2019 Society for Industrial and Applied Mathematics. All rights reserved.

arrive at

$$\Delta xyz \mathbf{P}^{(1,1,1)} = \mathbf{P}^{(1,1,1)} \underbrace{\left(D_{x,(0,0,0)}^{(1,1,1)} W_{x,(1,1,1)}^{(0,0,0)} + D_{y,(0,0,0)}^{(1,1,1)} W_{y,(1,1,1)}^{(0,0,0)} \right)}_{\Delta_W}.$$

Thus we arrive at a discretization of the variable coefficient Helmholtz equation:

$$(\Delta_W + R_{(0,0,0)}^{(1,1,1)} XY^2 L_{(1,1,1)}^{(0,0,0)}) \mathbf{u} = \mathbf{f},$$

where $X = X_{(0,0,0)}$, $Y = Y_{(0,0,0)}$ and

$$f(x,y) = \mathbf{P}^{(1,1,1)}(x,y) \mathbf{f}.$$

The sparsity structure of Δ_W and the variable coefficient Helmholtz equation is depicted in Figure 6.3(a).

To handle non-zero Dirichlet conditions we use the basis $\mathbf{Q}^{(1,1,1)}$ defined in Appendix B.2, which supplement $xyz \mathbf{P}^{(1,1,1)}$ to include the missing polynomials. We have weighted partial derivatives

$$\begin{aligned} \mathcal{D}_x \mathbf{Q}^{(1,1,1)} &= \mathbf{P}^{(0,0,0)} \tilde{W}_{x,(1,1,1)}^{(0,0,0)}, \\ \mathcal{D}_y \mathbf{Q}^{(1,1,1)} &= \mathbf{P}^{(0,0,0)} \tilde{W}_{y,(1,1,1)}^{(0,0,0)}, \end{aligned}$$

and thus we have

$$\Delta \mathbf{Q}^{(1,1,1)} = \mathbf{P}^{(1,1,1)} \underbrace{\left(D_{x,(0,0,0)}^{(1,1,1)} \tilde{W}_{x,(1,1,1)}^{(0,0,0)} + D_{y,(0,0,0)}^{(1,1,1)} \tilde{W}_{y,(1,1,1)}^{(0,0,0)} \right)}_{\Delta \tilde{W}}.$$

The basis also has a sparse block-banded restriction to a Legendre expansion along each edge (Proposition B.9) and is given by

$$\begin{aligned} \mathbf{Q}^{(1,1,1)}(x,0) &= \mathbf{P}(2x-1) R_{(1,1,1)}^{P,x=0}, \\ \mathbf{Q}^{(1,1,1)}(0,y) &= \mathbf{P}(2y-1) R_{(1,1,1)}^{P,y=0}, \\ \mathbf{Q}^{(1,1,1)}(x,1-x) &= \mathbf{P}(2x-1) R_{(1,1,1)}^{P,z=0}. \end{aligned}$$

Thus we arrive at the following discretization:

$$\begin{pmatrix} R_{(1,1,1)}^{P,x=0} \\ R_{(1,1,1)}^{P,y=0} \\ R_{(1,1,1)}^{P,z=0} \\ \Delta \tilde{W} + R_{(0,0,0)}^{(1,1,1)} X Y^2 \tilde{R}_{(1,1,1)}^{(0,0,0)} \end{pmatrix} \mathbf{u} = \begin{pmatrix} \mathbf{q} \\ \mathbf{r} \\ \mathbf{s} \\ \mathbf{f} \end{pmatrix}.$$

Neumann and Robin conditions can be incorporated by writing the PDE as a system of first-order PDEs, reducing these conditions to Dirichlet.

6.6. p -finite element methods

The Beuchler and Schöberl (2006) p -FEM on a triangle was originally constructed directly using Jacobi polynomial relationships to evaluate the corresponding inner products. Instead, as in Section 5.3, we use the quasi-matrix notation to manipulate recurrence relationships and thereby construct the p -FEM discretization. While we present the construction on a triangle, it is straightforward to derive a similar method on rectangles and disks.

Following the one-dimensional construction in Section 5.3, we write u and v in terms of weighted OPs on the right-angled triangle as

$$u = xyz \mathbf{P}^{(1,1,1)} \mathbf{u} \quad \text{and} \quad v = xyz \mathbf{P}^{(1,1,1)} \mathbf{v},$$

where $z = 1 - x - y$. Then we have

$$\langle \nabla v, \nabla u \rangle = \mathbf{v}^\top (\nabla w \mathbf{P}^{(1,1,1)})^\top \nabla xyz \mathbf{P}^{(1,1,1)} \mathbf{u}.$$

But we also have

$$\nabla xyz \mathbf{P}^{(1,1,1)} = \begin{pmatrix} \partial_x xyz \mathbf{P}^{(1,1,1)} \\ \partial_y xyz \mathbf{P}^{(1,1,1)} \end{pmatrix} = \begin{pmatrix} \mathbf{P}^{(0,0,0)} W_{x,(1,1,1)}^{(0,0,0)} \\ \mathbf{P}^{(0,0,0)} W_{y,(1,1,1)}^{(0,0,0)} \end{pmatrix}.$$

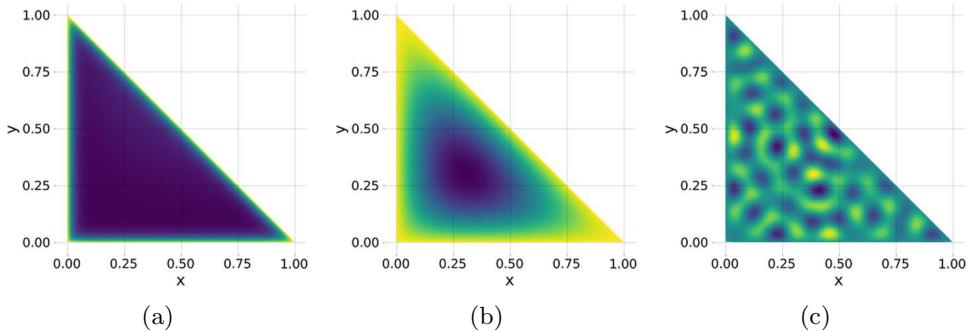


Figure 6.4. The ‘good’ with $k = -50^2$ (a), Poisson with $k = 0$ (b) and ‘bad’ with $k = 50^2$ (c) solutions to the inhomogeneous Helmholtz equation (6.4).

Lowering down to $\mathbf{P}^{(0,0,0)}$, we arrive at a sparse discretization of the weak Laplacian:

$$\begin{aligned} & (\nabla xyz \mathbf{P}^{(1,1,1)})^\top \nabla xyz \mathbf{P}^{(1,1,1)} \\ &= (W_{x,(1,1,1)}^{(0,0,0)})^\top M W_{x,(1,1,1)}^{(0,0,0)} + (W_{y,(1,1,1)}^{(0,0,0)})^\top M W_{y,(1,1,1)}^{(0,0,0)}, \end{aligned}$$

where

$$M = M_{(0,0,0)} = (\mathbf{P}^{(0,0,0)})^\top \mathbf{P}^{(0,0,0)}$$

is the diagonal mass matrix corresponding to Legendre polynomials on the triangle; see Proposition B.1. This discretization is symmetric and banded-block-banded with the same sparsity structure as the ultraspherical spectral method, as depicted in Figure 6.3.

6.6.1. Example: ‘good’ and ‘bad’ Helmholtz equations

A nice feature of p -FEM is that it preserves symmetry and positive definiteness, with the latter allowing for efficient Cholesky factorizations. To demonstrate this, we consider the inhomogeneous, variable coefficient Helmholtz equation

$$u'' + k(1 + xy^2)u = 1 \quad (6.4)$$

with zero Dirichlet conditions. When $k < 0$, we call this the ‘good’ Helmholtz equation as it is symmetric positive definite (SPD) and forms boundary layers. In contrast, for $k > 0$, the solution is highly oscillatory, and the operator is not positive definite but is still self-adjoint. For $k = 0$, this reduces to the Poisson equation.

In Figure 6.4 we depict solutions for $k = -50^2, 0, 50^2$ showing the transition from boundary layers to oscillations. In Figure 6.5 we show the coefficients for $k = -200^2, 0, 200^2$, which show that the corner singularities drive

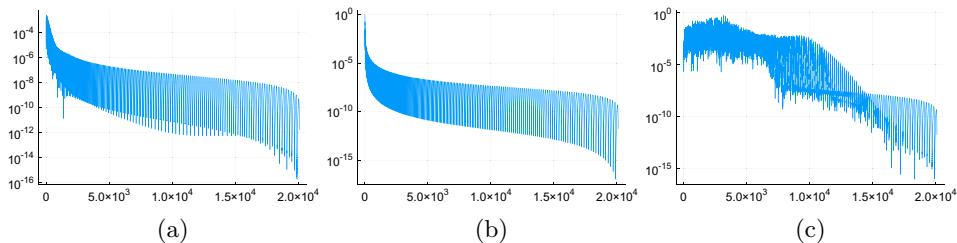


Figure 6.5. The coefficients for the ‘good’ with $k = -200^2$ (a), Poisson with $k = 0$ (b) and ‘good’ with $k = 200^2$ (c) solutions to the inhomogeneous Helmholtz equation (6.4).

algebraic decay. At the same time, the boundary layers and oscillations generate slow decay of a different nature. Thus, for moderate k , it is the corner singularities that dictate convergence. In Figure 6.6 we show the timings in building, factorizing and inverting the system. For $k \leq 0$ the operator is symmetric negative definite, and a Cholesky factorization is used, while for $k > 0$ the operator is only symmetric and an LDL^\top factorization is used. We use SuiteSparse to compute the factorizations, which demonstrates an apparent better complexity for Cholesky factorizations, though this may be lost for higher N .

6.7. Domain decomposition

The tools developed lead naturally to a method for solving PDEs on triangle meshes. For the ultraspherical spectral method, this is achieved by imposing continuity in value and derivative over the boundary and reformulating the problem to a system of first-order PDEs. We refer the reader to Olver *et al.* (2019) for further details. Figure 6.7 demonstrates these techniques on a Helmholtz equation on a polygon, by subdividing into triangles and imposing continuity up to normal derivatives between each element.

For the p -FEM, this is achieved by using the basis $xyz\mathbf{P}^{(1,1,1)}$ in each element as the *bubble functions*. *Face bubble functions* are then added between adjoining elements: for example, if we have two triangles, one on top of the other with adjoining edge from $(0,0)$ to $(1,0)$, on the interface we have

$$x(1-x)\mathbf{P}^{(1,1)},$$

which can be extended to the triangles as $Q_{n,1}^{(1,0,1)}(x,y)$, which vanishes on the other two edges. This extension can then be differentiated and converted as appropriate. We refer the reader to Beuchler and Schöberl (2006) for more details.

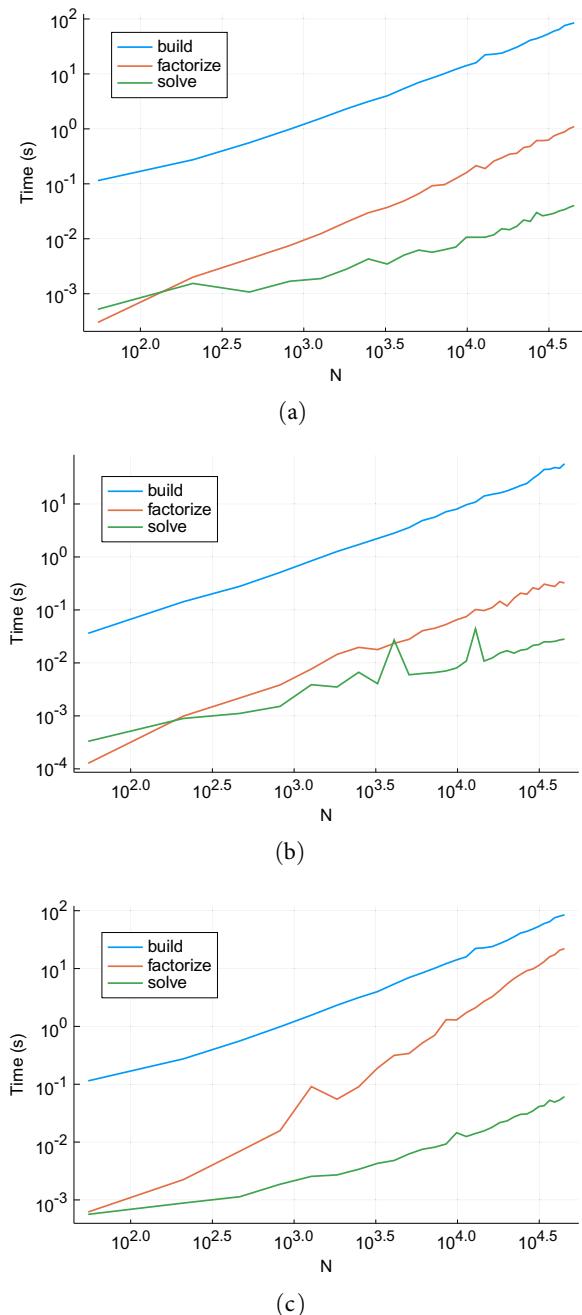


Figure 6.6. The timings for the ‘good’ with $k = -200^2$ (a), Poisson with $k = 0$ (b) and ‘bad’ with $k = 200^2$ (c) solutions to build, factorize and invert the inhomogeneous Helmholtz equation (6.4). For $k \leq 0$ the operator is symmetric negative definite and a Cholesky factorization is used, while for $k > 0$ the operator is only symmetric and an LDL^\top factorization is used.

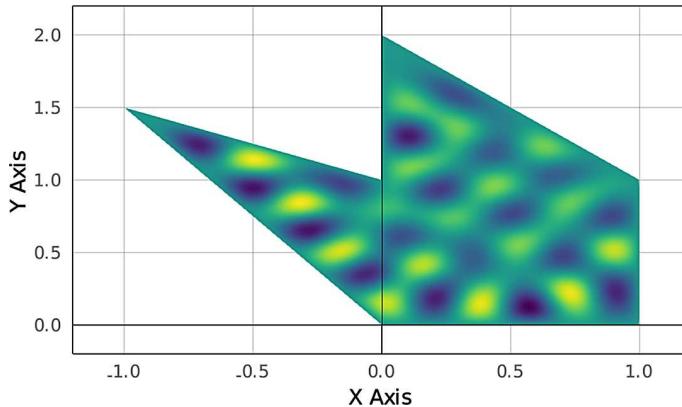


Figure 6.7. A solution to the Helmholtz equation on a polygon. Figure reproduced from Olver *et al.* (2019) with permission. Copyright © 2019 Society for Industrial and Applied Mathematics. All rights reserved.

7. Singular integral equations

Singular integral equations arise in potential theory, boundary integral reformulations of partial differential equations, and elsewhere. Conventional numerical methods typically involve discretization of the geometry using either quadrature rules (*à la* Nyström method) or collocation points, though these methods must deal with singularities in the integrals. A recent success story along these lines is the quadrature-by-expansion (QBX) method of Klöckner, Barnett, Greengard and O’Neil (2013).

As an alternative in two dimensions (*i.e.* one-dimensional singular integral equations corresponding to two-dimensional PDEs), we advocate using weighted OPs to calculate the singular integrals *exactly* via recurrence relationships. Cauchy transforms of weighted OPs satisfy the same recurrence relationship as the OPs themselves, which has been used successfully in numerical computing going back to Elliott (1982) and Gautschi and Wimp (1987). Less well known is that this is applicable for logarithmic singularities as well. Slevinsky and Olver (2017) used these relationships, which show that singular integral operators are banded on a single interval, to construct an ultraspherical spectral method for solving singular integral equations. Multiple scattered plates can also be tackled using low-rank techniques, and more complicated smooth arcs using suitable change-of-variables formulas.

7.1. Cauchy and Hilbert transforms

The *Cauchy transform* over a contour Γ is denoted by

$$\mathcal{C}_\Gamma f(z) := \frac{1}{2\pi i} \int_\Gamma \frac{f(t)}{t-z} dt, \quad (7.1)$$

which is analytic for $z \notin \Gamma$. Provided it has well-defined left and right limits (which is the case for reasonable functions f , *e.g.* smooth functions times Jacobi weights on $[-1,1]$), we use the notation

$$\mathcal{C}_\Gamma^+ f(x) = \lim_{\substack{z \rightarrow x \\ \text{from the left}}} \mathcal{C}_\Gamma f(z),$$

$$\mathcal{C}_\Gamma^- f(x) = \lim_{\substack{z \rightarrow x \\ \text{from the right}}} \mathcal{C}_\Gamma f(z).$$

We omit the subscript Γ when it is implied by the context.

Since Cauchy transforms are themselves just integrals, they can be discretized via a quadrature rule. That is,

$$\mathcal{C}_\Gamma f(z) \approx \frac{1}{2\pi i} \sum_{j=1}^n w_j \frac{f(t_j)}{t_j - z}, \quad (7.2)$$

where $\{t_j\}$ and $\{w_j\}$ are quadrature nodes and weights for the contour Γ . However, this is not a good idea for Cauchy transforms as they are singular integrals. In particular, the quantity on the left-hand side of (7.2) has a branch cut while the right-hand side of (7.2) has poles. We may need to evaluate the singular integral on or near Γ , where the quadrature rule is not accurate. We visualize this using a phase portrait⁸ in the style of Wegert (2012) in Figure 7.1. The true Cauchy transform has a branch cut on the interval, whereas Gauss quadrature results in poles along the branch cut.

As an alternative, we propose first expanding the integrand f in (7.1) in a weighted OP basis. Once an expansion is known, we have a natural approximation to the Cauchy transform:

$$\mathcal{C}_{(a,b)} f(z) \approx \sum_{n=0}^{N-1} f_n \mathcal{C}_{(a,b)}[wp_n](x).$$

The approximation on the right has a jump on the interval of orthogonality $[a,b]$ matching the true Cauchy transform. In fact, in a certain sense, we achieve uniform convergence (though we omit details), and therefore we can approximate the Cauchy transform up to and on the jump.

The calculation of a general Cauchy transform is reduced to one involving the generation of weighted OPs. Fortunately, Cauchy transforms satisfy a recurrence relationship that makes them straightforward to evaluate. Let

$$C_n(z) := \mathcal{C}_{(a,b)}[p_n w](z).$$

⁸ A phase portrait of a complex-analytical function plots only the phase of the function. Here red corresponds to $\theta = 0$, blue corresponds to $\theta = 2\pi/3$ and green corresponds to $\theta = -2\pi/3$.

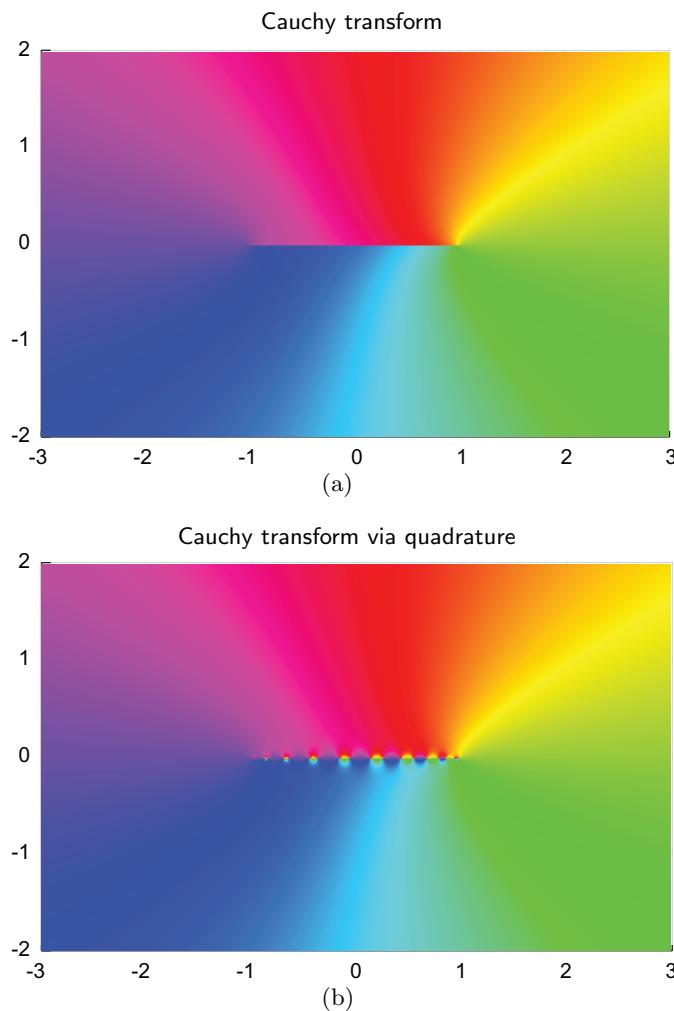


Figure 7.1. The Cauchy transform of e^x over $[-1, 1]$ computed using Legendre polynomials (a) compared to Gauss–Legendre quadrature (b). The Legendre polynomial-based approximation is uniformly accurate to roughly machine precision up to and on the interval. In contrast, the quadrature-based approach is not accurate on the contour itself, and has introduced artificial poles.

Then we have

$$\begin{aligned} zC_0(z) &= a_0C_0(z) + b_0C_1(z) - \frac{1}{2\pi i} \int_a^b w(x)dx, \\ zC_n(z) &= c_nC_{n-1}(z) + a_nC_n(z) + b_nC_{n+1}(z), \quad n \geq 1. \end{aligned}$$

This well-known formula (see *e.g.* Gautschi 1981b) is derived by the following simple calculation:

$$\begin{aligned} zC_n(z) &= \frac{1}{2\pi i} \int_a^b \frac{(z-x)p_n(x)w(x)}{x-z} dx + \int_a^b \frac{xp_n(x)w(x)}{x-z} dx \\ &= -\frac{1}{2\pi i} \int_a^b p_n(x)w(x)dx + c_nC_{n-1}(z) + a_nC_n(z) + b_nC_{n+1}(z). \end{aligned}$$

Note that the integral term is zero when $n > 0$.

The key feature of this recurrence is that if we know $C_0(z) = \mathcal{C}w(z)$ and $\int_a^b w(x)dx$, then we can determine $C_n(z)$ by solving the following lower-triangular system using forward substitution:

$$\left(\begin{array}{cccccc} 1 & & & & & \\ a_0 - z & b_0 & & & & \\ c_1 & a_1 - z & b_1 & & & \\ & c_2 & a_2 - z & b_2 & & \\ & & c_3 & a_3 - z & b_3 & \\ & & & \ddots & \ddots & \end{array} \right) \begin{pmatrix} C_0(z) \\ C_1(z) \\ C_2(z) \\ C_3(z) \\ C_4(z) \\ \vdots \end{pmatrix} = \begin{pmatrix} C_0(z) \\ \frac{1}{2\pi i} \int_a^b w(x)dx \\ 0 \\ 0 \\ 0 \\ \vdots \end{pmatrix}.$$

This recurrence-based approach also extends to the *Hilbert transform*⁹ given by

$$\mathcal{H}_{[a,b]}f(x) = \frac{1}{\pi} \int_a^b \frac{f(t)}{x-t} dt. \quad (7.3)$$

By the Plemelj theorem, the Hilbert and Cauchy transforms are related by (Muskhelishvili 1953, Trogdon and Olver 2016)

$$\mathcal{C}^+ - \mathcal{C}^- = I, \quad \mathcal{C}^+ + \mathcal{C}^- = i\mathcal{H}.$$

Therefore, the recurrence relationships for Cauchy transforms of OPs immediately imply one for Hilbert transforms of OPs. Let

$$H_n(z) := \mathcal{H}[p_n w](z).$$

⁹ The Hilbert transform is sometimes defined as minus that given in (7.3) (see *e.g.* Olver 2011, Trogdon and Olver 2016, Slevinsky and Olver 2017). Here, we prefer (7.3) as it has the convenient property that the Hilbert transform is the derivative of the log transform.

Then we have

$$\begin{aligned} xH_0(x) &= a_0H_0(x) + b_0H_1(x) + \frac{1}{\pi} \int_a^b w(x)dx, \\ xH_n(x) &= c_nH_{n-1}(x) + a_nH_n(x) + b_nH_{n+1}(x), \quad n \geq 1. \end{aligned}$$

7.1.1. Example: Hilbert transforms of Chebyshev polynomials

Consider the Hilbert transform of the Chebyshev polynomial T_n , that is,

$$H_n(x) = \frac{1}{\pi} \int_{-1}^1 \frac{T_n(t)}{(x-t)\sqrt{1-t^2}} dt.$$

The recurrence relationship satisfied by $H_n(x)$ is

$$xH_0(x) = H_1(x) + 1, \quad xH_n(x) = \frac{H_{n-1}(x)}{2} + \frac{H_n(x)}{2}, \quad n \geq 1.$$

In this case we have $H_0(x) = \mathcal{H}[(1-t^2)^{-1/2}](x) = 0$. Therefore we have

$$H_1(x) = -1, \quad H_2(x) = -2x, \quad H_{n+1}(x) = 2xH_n(x) - H_{n-1}(x).$$

This is precisely the three-term recurrence relationship satisfied by the Chebyshev polynomials (of the second kind). Due to the values for H_0 and H_1 , we find that $H_n(x) = -U_{n-1}(x)$ for $n \geq 1$. In quasimatrix notation, we have

$$\mathcal{H} \frac{\mathbf{T}}{\sqrt{1-\diamond^2}} = \mathbf{U} \underbrace{\begin{pmatrix} 0 & -1 & & \\ & -1 & & \\ & & \ddots & \\ & & & \ddots \end{pmatrix}}_{H_T^U},$$

where \diamond denotes the ‘dummy’ variable.

A related formula holds for the Hilbert transform of Chebyshev polynomials of the second kind, where the fact that $\mathcal{H}[\sqrt{1-\diamond^2}](x) = 1$ leads to $\mathcal{H}[\sqrt{1-\diamond^2}U_n] = T_{n+1}(x)$ for $n \geq 0$. In quasimatrix notation, we have

$$\mathcal{H} \sqrt{1-\diamond^2} \mathbf{U} = \mathbf{T} \underbrace{\begin{pmatrix} 0 & & & \\ 1 & & & \\ & 1 & & \\ & & \ddots & \\ & & & \ddots \end{pmatrix}}_{H_U^T}.$$

7.2. Inverting shifted Jacobi operators and Cauchy transforms

While forward recurrence is an effective approach for calculating Cauchy transforms, it can run into numerical stability issues if the evaluation point

is far from $[a,b]$: Cauchy transforms and OPs satisfy the same recurrence relationship (they are the *minimal* and *maximal* solutions, respectively) apart from the initial conditions. Any small perturbation to the initial condition for the minimal solution will pick up the maximal solution, which blows up as $n \rightarrow \infty$. On the other hand, $C_n(z) \rightarrow 0$ as $n \rightarrow \infty$ for z off the support of w .

There are effective methods for calculating minimal solutions to recurrence relationships. We have

$$\underbrace{\begin{pmatrix} a_0 - z & b_0 & & \\ c_1 & a_1 - z & b_1 & \\ & c_2 & a_2 - z & \ddots \\ & \ddots & \ddots & \ddots \end{pmatrix}}_{J-zI} \begin{pmatrix} C_0(z) \\ C_1(z) \\ C_2(z) \\ C_3(z) \\ \vdots \end{pmatrix} = \begin{pmatrix} \frac{1}{2\pi i} \int_a^b w(x) dx \\ 0 \\ 0 \\ \vdots \end{pmatrix}. \quad (7.4)$$

The operator $J - zI$ is invertible in ℓ^2 (assuming J is bounded) when z is not on the spectrum of J , which is precisely the support of the weight, the interval $[a,b]$. This means that the Cauchy transform is the unique solution to (7.4). As discussed in Section 5, there are multiple ways of solving an infinite-dimensional linear system such as (7.4), with Olver's algorithm being particularly attractive as it is fast, stable and adaptive.

There is a remaining question of when to use forward recurrence or directly inverting $J - zI$. As z approaches $[a,b]$, the decay rate of $C_n(z)$ degenerates, which requires computing an increasing number of terms to achieve convergence. In lockstep, the forward recurrence becomes increasingly stable to use. In practice we find that the following heuristic works well: if one requires the first $C_0(z), \dots, C_N(z)$ for a specified value N , use forward recurrence if N is in a Bernstein ellipse with radii proportional to $1/N$; otherwise, invert $J - zI$.

7.3. Log transforms

From the Green's function of the Laplacian, it is natural to consider integrals with a logarithmic kernel. If $\Gamma = [a,b]$, then the log transform is defined by

$$\mathcal{L}f(z) = \frac{1}{\pi} \int_a^b f(t) \log|z-t| dt,$$

where now we assume f is real-valued. It turns out that logarithmic singular integrals satisfy a simple recurrence relationship for weighted *classical* orthogonal polynomials, which we show by relating them to Cauchy transforms.

For $z \notin (-\infty, 1]$, we use the following formula:

$$v(z) = \operatorname{Re} \frac{1}{\pi} \int_a^b f(t) \log(z-t) dt.$$

(Note that the integrand avoids the branch cut of $\log z$.) To extend this to $z \in (-\infty, -1]$ (or more generally, $z \notin [-1, \infty)$), we can use the alternative expression given by

$$\begin{aligned} v(z) &= \operatorname{Re} \frac{1}{\pi} \int_a^b f(t) \log(t-z) dt = \operatorname{Re} \tilde{\mathcal{L}}f(z), \\ \tilde{\mathcal{L}}u(z) &:= \frac{1}{\pi} \int_a^b u(t) \log(z-t) dt, \end{aligned}$$

which follows from the fact that $\log|z-t| = \log|t-z|$.

To evaluate $\tilde{\mathcal{L}}u(z)$ in the complex plane we use the following expression¹⁰ in terms of Cauchy transforms:

$$\tilde{\mathcal{L}}u(z) = \frac{\log(z-a)}{\pi} \int_a^b u(x) dx + 2i\mathcal{C}U(z), \quad U(x) = \int_x^b u(t) dt. \quad (7.5)$$

As the indefinite integral of classical orthogonal polynomials is also orthogonal, we arrive at a closed-form expression in terms of Cauchy transforms of orthogonal polynomials.

7.3.1. Example: log kernel of Legendre, Chebyshev and Jacobi polynomials
Consider $\mathcal{L}P_n(z)$. For $n = 0$ we have

$$\tilde{\mathcal{L}}1(z) = \frac{2\log(z+1)}{\pi} + 2i\mathcal{C}[x-1](z).$$

For $n > 0$ the definite integral is zero due to orthogonality and we have an explicit expression for the indefinite integral from Proposition A.6, which gives us

$$\tilde{\mathcal{L}}P_n(z) = -\frac{i}{n}\mathcal{C}[(1-\diamond^2)P_{n-1}^{(1,1)}](z), \quad n \geq 1.$$

Taking the real part of $\tilde{\mathcal{L}}P_n(z)$ gives $\mathcal{L}P_n(z)$ for all n .

For Chebyshev polynomials, we have for $n > 0$, using Proposition A.16,

$$\int_x^1 \frac{T_n(t)}{\sqrt{1-t^2}} dt = -\frac{U_{n-1}(x)}{n} \sqrt{1-x^2}$$

so that

$$\tilde{\mathcal{L}}\left[\frac{T_n}{\sqrt{1-\diamond^2}}\right](z) = -\frac{2i}{n+1}\mathcal{C}[\sqrt{1-\diamond^2}U_{n-1}](z).$$

¹⁰ This follows from Liouville's theorem. Since the right-hand side of (7.5) matches the asymptotics and jumps of $\mathcal{L}u(z)$, they must be equal.

We are left with the $n = 0$ case, but this can be calculated explicitly as

$$\mathcal{L}[(1 - \diamond^2)^{-1/2}](z) = -\log|J_+^{-1}(z)| - \log 2$$

using the fact that

$$\int^z \frac{1}{\sqrt{z-1}\sqrt{z+1}} dz = -\log J_+^{-1}(z),$$

where $J_+^{-1}(z) = z - \sqrt{z-1}\sqrt{z+1}$ is a pre-image of the Joukowsky transform $z \mapsto (z + z^{-1})/2$.

A nice consequence of this is the following beautiful formula for evaluating on the interval $[-1, 1]$:

$$\begin{aligned} \mathcal{L}\left[\frac{T_n}{\sqrt{1-\diamond^2}}\right](x) &= -\frac{2i}{n}(\mathcal{C}^+ + \mathcal{C}^-)[\sqrt{1-\diamond^2}U_{n-1}](x) \\ &= -\frac{2}{n}\mathcal{H}[\sqrt{1-\diamond^2}U_{n-1}](x) \\ &= -\frac{T_n(x)}{n}. \end{aligned}$$

For the $n = 0$ case we have

$$\mathcal{L}[(1 - \diamond^2)^{-1/2}](x) = -\log 2.$$

Putting everything together, we arrive at¹¹

$$\mathcal{L}(1 - \diamond^2)^{-1/2}\mathbf{T} = \underbrace{\mathbf{T} \begin{pmatrix} -\log 2 & & & \\ & -1 & & \\ & & -1/2 & \\ & & & \ddots \end{pmatrix}}_{L_T}.$$

Finally, we find that the relationship $\mathcal{DL} = \mathcal{H}$ translates to

$$\mathcal{HL}(1 - \diamond^2)^{-1/2}\mathbf{T} = \mathcal{DL}(1 - \diamond^2)^{-1/2}\mathbf{T} = \mathcal{DL}L_T = \mathbf{U}D_T^U L_T = \mathbf{U}H_T^U.$$

Extension to other Jacobi polynomials is straightforward, though the $n = 0$ case often involves hypergeometric functions:

$$\begin{aligned} \mathcal{L}_{\mathbb{I}}w^{(\alpha, \beta)}(z) &= \frac{2^{\alpha+\beta+1}B(\alpha+1, \beta+1)}{\pi} \left\{ \log(z-1) \right. \\ &\quad \left. + \frac{\alpha+1}{\alpha+\beta+2} \left(\frac{2}{z-1} \right) {}_3F_2 \left(\begin{matrix} 1, 1, \alpha+2 \\ 2, \alpha+\beta+3 \end{matrix}; \frac{2}{1-z} \right) \right\}, \end{aligned} \tag{7.6}$$

which comes from integrating the Cauchy transforms of the Jacobi weight,

¹¹ Note that L_T should not be confused with the lowering operators, e.g. L_U^T , used before.

and where B is the Beta function (Olver *et al.* 2010, §5.12). This integration essentially follows from

$$\begin{aligned} \int z^{-1} {}_2F_1\left(\begin{array}{c} 1,b \\ c \end{array}; z\right) dz &= \int \sum_{k=0}^{\infty} \frac{(1)_k(b)_k}{(c)_k} \frac{z^{k-1}}{k!} dz, \\ &= \log z + C + \frac{b}{c} z {}_3F_2\left(\begin{array}{c} 1,1,b+1 \\ 2,c+1 \end{array}; z\right), \end{aligned}$$

where the constant is determined by the asymptotic behaviour at infinity.

Since generalized hypergeometric functions are challenging to compute near the branch cut, we identify this particular one with a partial derivative with respect to a parameter of a Gaussian hypergeometric function. Together with forward-mode automatic differentiation, this avoids numerical difficulties near the branch cut:

$$\begin{aligned} \mathcal{L}_{\mathbb{I}} w^{(\alpha, \beta)}(z) &= \frac{2^{\alpha+\beta+1} B(\alpha+1, \beta+1)}{\pi} \left\{ \log(z-1) \right. \\ &\quad \left. - \frac{\partial}{\partial \nu} {}_2F_1\left(\begin{array}{c} \nu, \alpha+1 \\ \alpha+\beta+2 \end{array}; \frac{2}{1-z}\right) \Big|_{\nu=0} \right\}. \end{aligned}$$

This formula is the result of taking the limit as $a \rightarrow 0$ below and using Olver *et al.* (2010, §15.5.3) and l'Hôpital's rule:

$$\begin{aligned} \lim_{a \rightarrow 0} \int z^{a-1} {}_2F_1\left(\begin{array}{c} a+1, b \\ c \end{array}; z\right) dz \\ = \lim_{a \rightarrow 0} \frac{1}{a} \left\{ z^a {}_2F_1\left(\begin{array}{c} a, b \\ c \end{array}; z\right) - 1 \right\} + C \\ = \frac{\partial}{\partial a} \left\{ z^a {}_2F_1\left(\begin{array}{c} a, b \\ c \end{array}; z\right) - 1 \right\} \Big|_{a=0} + C \\ = \log z + \frac{\partial}{\partial a} {}_2F_1\left(\begin{array}{c} a, b \\ c \end{array}; z\right) \Big|_{a=0} + C. \end{aligned}$$

7.4. Ultraspherical spectral methods for singular integral operators

There is a natural way of discretizing singular integral operators using sparse matrices (Slevinsky and Olver 2017). On a single interval this involves using the following quasimatrix relationships:

$$\mathcal{H}(1-\diamond^2)^{-1/2} \mathbf{T} = \mathbf{U} H_T^U \quad \text{and} \quad \mathcal{L}(1-\diamond^2)^{-1/2} \mathbf{T} = \mathbf{T} L_T.$$

Note that this also gives a way of calculating hyper-singular integrals. For example,

$$\mathcal{D}\mathcal{H}(1-\diamond^2)^{-1/2} \mathbf{T} = \mathbf{C}^{(2)} D_U^{(2)} H_T^U.$$

7.4.1. Example: the potential field of a metal sheet

As a simple example, consider the solution of the following Laplace equation, where $z = x + iy$:

- (1) $v_{xx} + v_{yy} = 0$ for $z \notin [-1,1] \cup \{i\epsilon\}$,
- (2) $v(x,y) = \log|z - i\epsilon| + O(1)$ as $z \rightarrow i\epsilon$,
- (3) $v(x,y) = o(1)$ as $z \rightarrow \infty$, and
- (4) $v(x,0) = \kappa$ for $-1 < x < 1$, where κ is an unknown constant.

This equation models the potential field of a unit charge located at $i\epsilon$ with a metal sheet that has no net charge placed on $[-1,1]$. As an ansatz for v , we write

$$v(z) = \mathcal{L}u(z) + \log|z - i\epsilon|,$$

which reduces the problem to one involving the following singular integral equation:

$$v(x,0) = \mathcal{L}u(x) = \kappa - \log|x - i\epsilon|, \quad -1 < x < 1, \quad \int_a^b u(x) dx = 0.$$

By writing $u = \mathbf{T}\mathbf{u}$, this becomes an almost diagonal linear system

$$\begin{pmatrix} 0 & \mathbf{e}_0^\top \\ -\mathbf{e}_0 & L_T \end{pmatrix} \begin{pmatrix} \kappa \\ \mathbf{u} \end{pmatrix} = \begin{pmatrix} 0 \\ \mathbf{f}_\epsilon \end{pmatrix}, \quad (7.7)$$

where $-\log|x - i\epsilon| = \mathbf{T}(x)\mathbf{f}_\epsilon$. Analyticity ensures that \mathbf{f}_ϵ has $N = O(1/\epsilon)$ coefficients above a prescribed tolerance, which can be calculated in $O(N \log N)$ time, using the Chebyshev transform described in [Section 4.1](#). The coefficients \mathbf{u} are then determined by solving (7.7), which costs $O(N)$ operations. Finally, we can evaluate $v(x,y)$ pointwise using the expressions involving the Cauchy transform, which in turn is achieved via the recurrence relationship. We demonstrate this example in [Figure 7.2](#), which depicts the solution with $\epsilon = 0.1$ and shows the number of coefficients needed to resolve the solution for different choices of ϵ , demonstrating the necessity of fast algorithms.

7.4.2. Other kernels

This framework also supports more general kernels such as those arising in elliptic PDEs, with the Helmholtz equation being the canonical case. In this setting, we are guaranteed that the kernel can be written as

$$K_1(x,y) \log|x - y| + \frac{K_2(x,y)}{x - y} + K_3(x,y),$$

where integration against this kernel is understood in the principal value

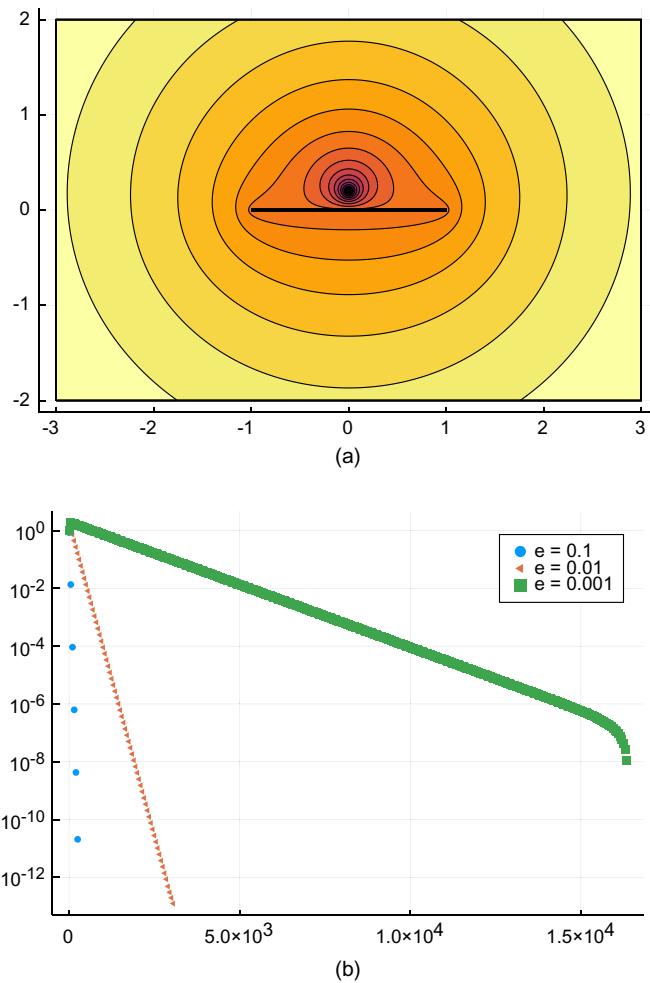


Figure 7.2. The solution to the Laplace equation with a single point source at $0.2i$ as described in Section 7.4.1 (a), and the coefficients of the solution for $\epsilon = 0.1, 0.01, 0.001$ (b). This figure demonstrates the benefits of an $O(N)$ algorithm when the forcing term is challenging to resolve.

sense. At this point, we approximate K_m by a low-rank function approximation

$$K_m(x,y) \approx \sum_{j=1}^r a_j^m(x) b_j^m(y), \quad m = 1, 2, 3,$$

where, for simplicity, we assume that the rank, r , does not depend on m . A low-rank approximation of K_m can be computed using a singular value decomposition or other low-rank techniques for functions (Townsend and Trefethen 2013). Setting $u = \mathbf{T}u$, we have the following expression for K_1 :

$$\begin{aligned} \int K_1(x,y) \log|x-y| \frac{\mathbf{T}(y)}{\sqrt{1-y^2}} dy &= \sum_{j=1}^r a_j^1(x) \mathcal{L} \left[\frac{b_j^1 \mathbf{T}}{\sqrt{1-\diamond^2}} \right] (x) \\ &= \mathbf{T}(x) \sum_{j=1}^r a_j^1(X_T) L_T b_j^1(X_T), \end{aligned}$$

where X_T is again the multiplication-by- x operator corresponding to the Chebyshev polynomials \mathbf{T} . The other kernels are handled similarly, leading to the discretization

$$\mathbf{U} \sum_{j=1}^r [R_T^U a_j^1(X_T) L_T b_j^1(X_T) + a_j^2(X_U) H_T^U b_j^1(X_T) + a_j^3(X_U) \mathbf{e}_0 \Sigma_T b_j^3(X_T)].$$

The bandwidth of the discretization depends on the polynomial degree of a_j^m and b_j^m for $m = 1, 2, 3$. If a_j^m and b_j^m for $m = 1, 2, 3$ are all of degree $\leq d$, then the number of sub/super-diagonals is precisely d .

7.4.3. Example: Helmholtz on an interval

In Figure 7.3(a) we demonstrate an acoustic scattering analogue of the example presented in Section 7.4.1, now with a large number of point scatterers on a Bernstein ellipse surrounding $[-1, 1]$. That is, we wish to solve

$$v_{xx} + v_{yy} + k^2 v = 0 \quad \text{for } z \notin [-1, 1] \cup \Gamma_\epsilon,$$

where Γ_ϵ consists of the set of point sources on the Bernstein ellipse. An increasing number of coefficients are needed to resolve the solution as the point sources approach the interval, but the resulting banded system can be solved in $O(N)$ complexity. This example originally appeared in Slevinsky and Olver (2017), which also includes timings showing the $O(N)$ complexity, where N is the total degrees of freedom needed to resolve the solution.

7.5. Multiple segments

So far we have described a systematic approach to solving singular integral equations on $[-1, 1]$ with a similar construction being viable on the unit

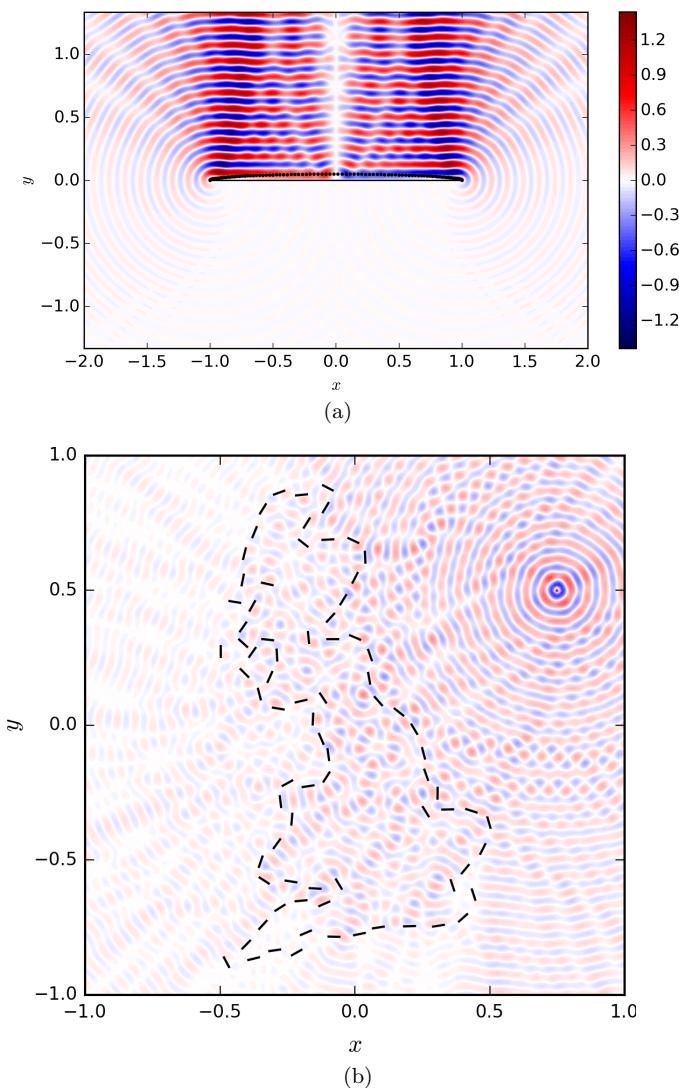


Figure 7.3. Acoustic scattering off an interval with many nearby point scatterers (a) and off a metal sheet model of Great Britain (b). Figure reproduced from Slevinsky and Olver (2017) with permission. Copyright © 2017 Elsevier.

circle with Laurent series. Similarly, any affine transformation of a line-segment embedded in \mathbb{R}^2 (or equivalently in \mathbb{C}) is straightforward. To solve SIEs posed on multiple intervals/circles, we discretize them in a block-wise manner.

As a simple example, suppose we want to solve Laplace's equation on the complement of two intervals $\Gamma_{-1} := [-2, -1]$ and $\Gamma_1 := [1, 2]$ and determine an expansion of the solution in the form

$$u(x) = \begin{cases} \sum_{k=0}^{\infty} u_k (1 - (2x+3)^2)^{-1/2} T_k(2x+3) & -1 < x < -2, \\ \sum_{k=0}^{\infty} v_k (1 - (2x-3)^2)^{-1/2} T_k(2x-3) & 1 < x < 2. \end{cases}$$

Define $\mathbf{T}_{-1}(x) := \mathbf{T}(2x+3)$ and $\mathbf{T}_1(x) := \mathbf{T}(2x-3)$ to be the mapped Chebyshev polynomials with corresponding weights $w_{-1}(x) = (1 - (2x+3)^2)^{-1/2}$ and $w_1(x) = (1 - (2x-3)^2)^{-1/2}$. Using a block-wise version of the quasimatrix notation, we want to solve

$$\begin{pmatrix} \mathcal{L}_{\Gamma_{-1}} & \mathcal{L}_{\Gamma_1}|_{\Gamma_{-1}} \\ \mathcal{L}_{\Gamma_{-1}}|_{\Gamma_1} & \mathcal{L}_{\Gamma_1} \end{pmatrix} \begin{pmatrix} w_{-1}\mathbf{T}_{-1} & \\ & w_1\mathbf{T}_1 \end{pmatrix} \begin{pmatrix} \mathbf{u}_{-1} \\ \mathbf{u}_1 \end{pmatrix} = \begin{pmatrix} \mathbf{T}_{-1}\mathbf{f}_{-1} \\ \mathbf{T}_1\mathbf{f}_1 \end{pmatrix}.$$

As $\log|x-y|$ is a Hilbert–Schmidt operator when $y \in \Gamma_{-1}$ and $x \in \Gamma_1$, it is clear we have

$$\mathcal{L}_{\Gamma_{-1}}|_{\Gamma_1} w_{-1} \mathbf{T}_{-1} = \mathbf{T}_1 L_{T_{-1}}^{T_1},$$

where $L_{T_{-1}}^{T_1}$ is a compact operator in appropriate spaces (in fact, any algebraically decaying space). Thus we can effectively approximate it by a finite-dimensional matrix $L_{T_{-1},m}^{T_1}$. The entries can be calculated via a variety of means such as a tensor product constructors or exploiting the recurrence relationship. Here we use the following approximation:

$$\begin{aligned} & \begin{pmatrix} \mathcal{L}_{\Gamma_{-1}} & \mathcal{L}_{\Gamma_1}|_{\Gamma_{-1}} \\ \mathcal{L}_{\Gamma_{-1}}|_{\Gamma_1} & \mathcal{L}_{\Gamma_1} \end{pmatrix} \begin{pmatrix} w_{-1}\mathbf{T}_{-1} & \\ & w_1\mathbf{T}_1 \end{pmatrix} \\ &= \begin{pmatrix} \mathbf{T}_{-1} & \\ & \mathbf{T}_1 \end{pmatrix} \begin{pmatrix} L_{T_{-1}} & L_{T_1}^{T_{-1}} \\ L_{T_{-1}}^{T_1} & L_{T_1} \end{pmatrix} \\ &\approx \begin{pmatrix} \mathbf{T}_{-1} & \\ & \mathbf{T}_1 \end{pmatrix} \begin{pmatrix} L_{T_{-1}} & L_{T_1,m}^{T_{-1}} \\ L_{T_{-1},m}^{T_1} & L_{T_1} \end{pmatrix}. \end{aligned}$$

The resulting approximation is a finite-dimensional perturbation of a block-diagonal system, and therefore requires only $O(N)$ operations to solve if the right-hand side has N non-zero entries.

The extension to many intervals is also possible. As an example, Figure 7.3(b) shows the acoustic scattering off many metal wires. Note that the complexity does depend on the geometry: the closer the intervals are to each other the larger m must be, and the total perturbation from diagonal

typically grows with the number of intervals. Hierarchical matrix methods may provide an attractive way around this issue.

7.6. Other geometries

It is also possible to calculate Cauchy transforms over smooth arcs by approximating them using a degree d polynomial map $p: [-1, 1] \rightarrow \Gamma$, which is one-to-one. By using an explicit change-of-variable formula found in Olver (2012) and Trogdon and Olver (2016), we have

$$\mathcal{C}_\Gamma f(z) = \sum_{j=1}^d \mathcal{C}_{[-1,1]}[f \circ p](p_j^{-1}(z)),$$

where $p_1^{-1}(z), \dots, p_d^{-1}(z)$ are the d pre-images of p . That is, $p(p_j^{-1}(z)) = z$. This leads to an expression

$$\mathcal{H}_\Gamma f(s) = \mathcal{H}_{[-1,1]}[f \circ p](p^{-1}(s)) + \text{compact},$$

where $p^{-1}: \Gamma \rightarrow [-1, 1]$. This shows that the ultraspherical spectral method for SIEs extends naturally to the mapped basis $\mathbf{T} \circ p^{-1}$:

$$\mathcal{H}_\Gamma(\mathbf{T} \circ p^{-1}) = (\mathbf{U} \circ p^{-1})[H_T^U + \text{compact}].$$

The log kernel can be handled similarly using the expression¹²

$$\mathcal{L}_\Gamma f(z) = \sum_{j=1}^d \mathcal{L}_{[-1,1]} \{(f \circ p)|p'|\}(p_j^{-1}(z)) + \log M \int_\Gamma f ds,$$

where $p(z) = Mz^d + O(z^{d-1})$. What remains unclear is how to handle geometries with corners.

7.7. Fractional differential equations

Orthogonal polynomials have nice recurrence relationships for fractional derivatives, which have been used in the construction of collocation schemes (Zayernouri and Karniadakis 2014, Chen, Shen and Wang 2016), achieving spectral accuracy when the solution is smooth. Hale and Olver (2018) considered equations involving one-sided fractional integral operators and fractional differential equations of Riemann–Liouville or Caputo type of

¹² This expression has not appeared before, but follows from Plemelj’s theorem on a related expression for $\tilde{\mathcal{L}}_\Gamma$. The subject of branch cuts is delicate, however, so we present the formula without the proof.

rational order. This has the interesting feature of considering an expansion in multiple types of singularities at the same time; for example, the solution looks like

$$u(x) = f(x) + \sqrt{x}g(x),$$

where f and g are smooth, and the approximation is achieved using suitably chosen orthogonal expansion bases. Extensions to fractional Laplacians also appear to be feasible. For brevity we refer the readers to the existing publications, but please stay tuned to new publications as this is evolving rapidly.

On the sphere, fractional differential operators have representations in terms of spherical harmonic expansions, which have been successfully used in combination with fast spherical harmonic transforms (Slevinsky, Montanelli and Du 2018) to compute solutions to fractional evolution equations.

8. Conclusions

Despite over a hundred years of research, there are many directions for future work on OPs.

- (1) *Semiclassical orthogonal polynomials.* These are becoming increasingly of interest in applications to frames (Huybrechs 2010), computations on geometries defined by quadratic curves (Olver and Xu 2019b, Olver and Xu 2020b, Olver and Xu 2020a) including solving PDEs on light cones (Olver and Xu 2019a) and disk slices and trapeziums (Snowball and Olver 2020). We anticipate that semiclassical OPs will become increasingly important in applications involving non-standard geometries.
- (2) *Multivariate orthogonal polynomials that are not constructed by tensorization or by Koornwinder's method* (which implies they are separable in *some* coordinate system). This makes transforms much more expensive as even direct non-separable transforms in d dimensions cost $O(N^{2d})$ for N^d degrees of freedom. With separability in all d dimensions, the direct cost is only $O(N^{d+1})$.
- (3) *Full hp-adaptivity in spectral element methods.* For small p , effective h adaptivity is prevalent and effective (Karniadakis and Sherwin 2013). In this review we have focused on large p , where sparsity in the discretization lends itself to effective p -adaptive schemes. What remains open is how to synthesize the two regimes effectively.
- (4) *Eigenproblems with continuous spectra.* Orthogonal polynomials simultaneously solve differential eigenvalue problems with discrete spectra and are eigenfunctions of Jacobi operators with continuous spectra.

This hints at the possibility of their use in solving differential equations with continuous spectra without discretizing the spectrum itself. Promising work along this avenue includes that of Webb (2017).

- (5) *Spectral methods for non-local analogues of calculus operators* (Du, Gunzburger, Lehoucq and Zhou 2012, Du, Gunzburger, Lehoucq and Zhou 2013).

A. Classical orthogonal polynomial recurrence relationships

A.1. Jacobi

The results in this appendix hold when $\alpha, \beta > -1$. Throughout this appendix, ‘DLMF’ refers to the NIST Handbook of Mathematical Functions (Olver, Lozier, Boisvert and Clark 2010).

Proposition A.1 (mass matrix (DLMF, Table 18.3.1)). We have

$$(\mathbf{P}^{(\alpha, \beta)})^\top (1-x)^\alpha (1+x)^\beta \mathbf{P}^{(\alpha, \beta)} = M_{(\alpha, \beta)},$$

where $M_{(\alpha, \beta)}$ is the diagonal matrix satisfying

$$(M_{(\alpha, \beta)})_{n,n} = 2^{\alpha+\beta+1} \frac{\Gamma(n+\alpha+1)\Gamma(n+\beta+1)}{(2n+\alpha+\beta+1)\Gamma(n+\alpha+\beta+1)n!}.$$

Proposition A.2 (reflection). For the reflection operator

$$\mathcal{R}f(x) := f(-x),$$

we have

$$\mathcal{R}\mathbf{P}^{(\alpha, \beta)} = \mathbf{P}^{(\beta, \alpha)} \underbrace{\begin{pmatrix} 1 & & & \\ & -1 & & \\ & & 1 & \\ & & & \ddots \end{pmatrix}}_{\Sigma}.$$

Proposition A.3 (conversion (DLMF, 18.9.5)). For Jacobi polynomials the conversion is given by

$$\mathbf{P}^{(\alpha, \beta)} = \mathbf{P}^{(\alpha, \beta+1)} \underbrace{\begin{pmatrix} 1 & \frac{1+\alpha}{\alpha+\beta+3} & & \\ & \frac{\alpha+\beta+2}{\alpha+\beta+3} & \frac{2+\alpha}{\alpha+\beta+5} & \\ & & \frac{\alpha+\beta+3}{\alpha+\beta+5} & \frac{3+\alpha}{\alpha+\beta+7} \\ & & & \ddots & \ddots \end{pmatrix}}_{R_{(\alpha, \beta)}^{(\alpha, \beta+1)}}.$$

We also have

$$\mathbf{P}^{(\alpha, \beta)} = \underbrace{\mathbf{P}^{(\alpha+1, \beta)} \Sigma R_{(\alpha, \beta)}^{(\alpha, \beta+1)} \Sigma}_{R_{(\alpha, \beta)}^{(\alpha+1, \beta)}},$$

where Σ was defined in Proposition A.2. The conversion operator $R_{(\alpha, \beta)}^{(\alpha+k, \beta+j)}$ for integers $k, j \geq 0$ can be defined by composing conversions together; for example,

$$R_{(\alpha, \beta)}^{(\alpha+k, \beta+j)} := R_{(\alpha+k, \beta+j-1)}^{(\alpha+k, \beta+j)} \cdots R_{(\alpha+k, \beta)}^{(\alpha+k, \beta+1)} R_{(\alpha+k-1, \beta)}^{(\alpha+k, \beta)} \cdots R_{(\alpha, \beta)}^{(\alpha+1, \beta)}.$$

Proposition A.4 (weighted conversion (DLMF, 18.9.5)). For Jacobi polynomials the weighted conversion is given by

$$(1+x)\mathbf{P}^{(\alpha, \beta+1)} = \mathbf{P}^{(\alpha, \beta)} \underbrace{\begin{pmatrix} \frac{2\beta+2}{\alpha+\beta+2} & & & & \\ & \frac{4+2\beta}{\alpha+\beta+4} & & & \\ \frac{2}{\alpha+\beta+2} & & \frac{4}{\alpha+\beta+4} & & \frac{6+2\beta}{\alpha+\beta+6} \\ & & & \frac{6}{\alpha+\beta+6} & \frac{8+2\beta}{\alpha+\beta+8} \\ & & & & \ddots & \ddots \end{pmatrix}}_{L_{(\alpha, \beta+1)}^{(\alpha, \beta)}}.$$

We also have

$$(1-x)\mathbf{P}^{(\alpha+1, \beta)} = \mathbf{P}^{(\alpha, \beta)} \underbrace{\Sigma L_{(\alpha, \beta+1)}^{(\alpha, \beta)} \Sigma}_{L_{(\alpha+1, \beta)}^{(\alpha, \beta)}}.$$

The weighted conversion operator $L_{(a+k, b+j)}^{(a, b)}$ for integers $k, j \geq 0$ can be defined by composing weighted conversions together; for example,

$$L_{(\alpha+k, \beta+j)}^{(\alpha, \beta)} := L_{(\alpha+1, \beta)}^{(\alpha, \beta)} \cdots L_{(\alpha+k, \beta)}^{(\alpha+k-1, \beta)} L_{(\alpha+k, \beta+1)}^{(\alpha+k, \beta)} \cdots L_{(\alpha+k, \beta+j)}^{(\alpha+k, \beta+j-1)}.$$

Proposition A.5 (differentiation (DLMF, 18.9.15)). For the derivative operator

$$\mathcal{D}f(x) := \frac{\partial}{\partial x},$$

we have

$$\mathcal{D} \mathbf{P}^{(\alpha, \beta)} = \mathbf{P}^{(\alpha+1, \beta+1)} \underbrace{\begin{pmatrix} 0 & \frac{\alpha+\beta+2}{2} & & \\ & & \frac{\alpha+\beta+3}{2} & \\ & & & \frac{\alpha+\beta+4}{2} \\ & & & \ddots \end{pmatrix}}_{D_{(\alpha, \beta)}^{(\alpha+1, \beta+1)}}.$$

Proposition A.6 (weighted differentiation (DLMF, 18.9.16)). The weighted differentiation operator is given by

$$\mathcal{D}(1-x)^{\alpha+1}(1+x)^{\beta+1} \mathbf{P}^{(\alpha+1, \beta+1)} = \mathbf{P}^{(\alpha, \beta)} \underbrace{\begin{pmatrix} 0 & & & \\ -2 & -4 & & \\ & & \ddots & \\ & & & \end{pmatrix}}_{W_{(\alpha+1, \beta+1)}^{(\alpha, \beta)}}.$$

Proposition A.7 (one-sided weighted differentiation.). The one-sided weighted differentiation operators are given by

$$\mathcal{D}(1-x)^{\alpha+1} \mathbf{P}^{(\alpha+1, \beta)} = (1-x)^\alpha \mathbf{P}^{(\alpha, \beta+1)} \underbrace{\begin{pmatrix} \alpha+1 & & & \\ & \alpha+2 & & \\ & & \alpha+3 & \\ & & & \ddots \end{pmatrix}}_{W_{R, (\alpha+1, \beta)}^{(\alpha, \beta+1)}},$$

$$\mathcal{D}(1+x)^{\beta+1} \mathbf{P}^{(\alpha, \beta+1)} = (1+x)^\beta \mathbf{P}^{(\alpha, \beta+1)} \underbrace{\begin{pmatrix} \beta+1 & & & \\ & \beta+2 & & \\ & & \beta+3 & \\ & & & \ddots \end{pmatrix}}_{W_{L, (\alpha, \beta+1)}^{(\alpha+1, \beta)}}.$$

Proof. Equivalent to (DLMF, 15.5.4, 15.5.6), as outlined in Olver, Townsend and Vasil (2020). \square

A.2. Ultraspherical

The results in this section hold when $\lambda > 0$.

Proposition A.8 (mass matrix (DLMF, Table 18.3.1)). We have

$$(\mathbf{C}^{(\lambda)})^\top (1-x^2)^{\lambda-1/2} \mathbf{C}^{(\lambda)} = M_{(\lambda)},$$

where $M_{(\lambda)}$ is the diagonal matrix satisfying

$$(M_{(\lambda)})_{n,n} = \frac{2^{1-2\lambda}\pi\Gamma(n+2\lambda)}{(n+\lambda)\Gamma(\lambda)^2 n!}.$$

Proposition A.9 (conversion (DLMF, 18.9.7)). We have

$$\mathbf{C}^{(\lambda)} = \mathbf{C}^{(\lambda+1)} \underbrace{\begin{pmatrix} 1 & 0 & -\frac{\lambda}{2+\lambda} & & \\ & \frac{\lambda}{1+\lambda} & 0 & -\frac{\lambda}{3+\lambda} & \\ & & \frac{\lambda}{2+\lambda} & 0 & -\frac{\lambda}{5+\lambda} \\ & & & \ddots & \ddots \\ & & & & \ddots \end{pmatrix}}_{R_{(\lambda)}^{(\lambda+1)}}$$

for the conversion between ultraspherical polynomials. By composing conversions, we have $R_{(\lambda)}^{(\lambda+m)} = R_{(\lambda+m-1)}^{(\lambda+m)} \cdots R_{(\lambda)}^{(\lambda+1)}$.

Proposition A.10 (weighted conversion (DLMF, 18.9.8)). We have

$$(1-x^2)\mathbf{C}^{(\lambda+1)} = \mathbf{C}^{(\lambda)} L_{(\lambda+1)}^{(\lambda)},$$

for the weighted conversion between ultraspherical polynomials, where the entries of $L_{(\lambda+1)}^{(\lambda)}$ are determined by the recurrence

$$\begin{aligned} 4\lambda(n+\lambda+1)(1-x^2)C_n^{(\lambda+1)}(x) \\ = -(n+1)(n+2)C_{n+2}^{(\lambda)}(x) + (n+2\lambda)(n+2\lambda+1)C_n^{(\lambda)}(x). \end{aligned}$$

By composing weighted conversions, we have $L_{(\lambda+m)}^{(\lambda)} = L_{(\lambda+1)}^{(\lambda)} \cdots L_{(\lambda+m)}^{(\lambda+m-1)}$.

Proposition A.11 (differentiation (DLMF, 18.9.19)). For the derivative operator \mathcal{D} , we have

$$\mathcal{D}\mathbf{C}^{(\lambda)} = \mathbf{C}^{(\lambda+1)} \underbrace{\begin{pmatrix} 0 & 2\lambda & & \\ & 2\lambda & & \\ & & 2\lambda & \\ & & & \ddots \end{pmatrix}}_{D_{(\lambda)}^{(\lambda+1)}}.$$

We also use $D_{(\lambda)}^{(\lambda+m)} = D_{(\lambda+m-1)}^{(\lambda+m)} \cdots D_{(\lambda)}^{(\lambda+1)}$.

A.3. Chebyshev

Proposition A.12 (mass matrix (DLMF, Table 18.3.1)). We have

$$\mathbf{T}^\top (1-x^2)^{-1/2} \mathbf{T} = M_T, \quad \mathbf{U}^\top (1-x^2)^{1/2} \mathbf{U} = M_U,$$

where M_T and M_U are diagonal matrices satisfying

$$\begin{aligned}(M_T)_{0,0} &= \pi, & (M_T)_{n,n} &= \pi/2, & n > 0, \\ (M_U)_{n,n} &= \pi/2, & n \geq 0.\end{aligned}$$

Combining the conversion between Jacobi and Chebyshev/ultraspherical polynomials, we have the following.

Proposition A.13 (conversion (DLMF, 18.9.9, 18.9.11)). We have

$$\begin{aligned}\mathbf{T} &= \mathbf{U} \underbrace{\begin{pmatrix} 1 & 0 & -1/2 \\ & 1/2 & 0 & -1/2 \\ & & 1/2 & 0 & -1/2 \\ & & & \ddots & \ddots & \ddots \end{pmatrix}}_{R_T^U} \\ &= \mathbf{V} \underbrace{\begin{pmatrix} 1 & 1/2 \\ & 1/2 & 1/2 \\ & & 1/2 & 1/2 \\ & & & \ddots & \ddots & \ddots \end{pmatrix}}_{R_T^V}.\end{aligned}$$

We also use $R_U^{(\lambda)} := R_{(1)}^{(\lambda)}$ and $R_T^{(\lambda)} = R_U^{(\lambda)} R_T^U$.

Proposition A.14 (weighted conversion (DLMF, 18.9.10)). For the weighted conversion between Chebyshev polynomials we have

$$(1-x^2)\mathbf{U} = \mathbf{T} \underbrace{\begin{pmatrix} 1/2 \\ 0 & 1/2 \\ -1/2 & 0 & 1/2 \\ & \ddots & \ddots & \ddots \end{pmatrix}}_{L_U^T}.$$

We also have $L_{(\lambda)}^U = L_{(\lambda)}^{(1)}$ and $L_{(\lambda)}^T = L_U^T L_{(\lambda)}^U$.

Proposition A.15 (differentiation (DLMF, 18.9.21)). The differentiation operator between Chebyshev polynomials is given by

$$\mathcal{D}\mathbf{T} = \mathbf{U} \underbrace{\begin{pmatrix} 0 & 1 \\ & 2 \\ & & 3 \\ & & & \ddots \end{pmatrix}}_{D_T^U}.$$

We also have $D_U^{(\lambda)} = D_{(1)}^{(\lambda)}$ and $D_T^{(\lambda)} = D_U^{(\lambda)} D_T^U$.

Proposition A.16 (weighted differentiation (DLMF, 18.9.22)). The weighted differentiation operator for Chebyshev polynomials is given by

$$\mathcal{D}\sqrt{1-x^2}\mathbf{U} = \frac{1}{\sqrt{1-x^2}}\mathbf{T} \underbrace{\begin{pmatrix} 0 & & & \\ -1 & -2 & & \\ & & -3 & \\ & & & \ddots \end{pmatrix}}_{W_U^T}.$$

We also have $W_{(\lambda)}^U = W_{(\lambda)}^{(1)}$ and $W_{(\lambda)}^T = W_U^T W_{(\lambda)}^U$.

Proposition A.17 (integration). Let \mathcal{Q} denote the indefinite integration operator normalized so that the zeroth Chebyshev coefficient vanishes. That is,

$$\mathcal{D}\mathcal{Q} = \mathcal{I}, \quad \int_{-1}^1 \frac{(\mathcal{Q}f)(x)}{\sqrt{1-x^2}} dx = 0.$$

Then we have

$$\int \mathbf{T} = \mathbf{T} \underbrace{\begin{pmatrix} 0 & 0 & & & \\ 1 & 0 & -1/2 & & \\ & 1/4 & 0 & -1/4 & \\ & & 1/6 & 0 & -1/6 \\ & & & \ddots & \ddots & \ddots \end{pmatrix}}_{Q_T}.$$

Proof. This follows from the relation

$$\mathcal{Q}\mathbf{T}Q_T = \mathbf{U}D_T^UQ_T,$$

and observing that $D_T^UQ_T = R_T^U$. □

A.4. Laguerre

We use a tilde above operators for Laguerre polynomials to avoid confusion with ultraspherical polynomials. The results in this section hold when $\alpha > -1$.

Proposition A.18 (mass matrix (DLMF, Table 18.3.1)). We have

$$(\tilde{\mathbf{L}}^{(\alpha)})^\top x^\alpha e^{-x} \tilde{\mathbf{L}}^{(\alpha)} = \tilde{M}_{(\alpha)},$$

where $\tilde{M}_{(\alpha)}$ is the diagonal matrix satisfying

$$(\tilde{M}_{(\alpha)})_{n,n} = \frac{\Gamma(n+\alpha+1)}{n!}.$$

Proposition A.19 (conversion). We have

$$\mathbf{L}^{(\alpha)} = \mathbf{L}^{(\alpha+1)} \underbrace{\begin{pmatrix} 1 & -1 & & & \\ & 1 & -1 & & \\ & & 1 & -1 & \\ & & & \ddots & \ddots \end{pmatrix}}_{\tilde{R}_{(\alpha)}^{(\alpha+1)}}$$

for the conversion operator between Laguerre polynomials. We also have $\tilde{R}_{(\alpha)}^{(\alpha+m)} = \tilde{R}_{(\alpha+m-1)}^{(\alpha+m)} \cdots \tilde{R}_{(\alpha)}^{(\alpha+1)}$.

Proposition A.20 (weighted conversion). For the weighted conversion operator we have

$$x\mathbf{L}^{(\alpha+1)} = \mathbf{L}^{(\alpha)} \underbrace{\begin{pmatrix} \alpha+1 & & & & \\ -1 & \alpha+2 & & & \\ -2 & & \alpha+3 & & \\ & & & \ddots & \ddots \end{pmatrix}}_{\tilde{L}_{(\alpha+1)}^{(\alpha)}}.$$

We also have $\tilde{L}_{(\alpha+m)}^{(\alpha)} = \tilde{L}_{(\alpha+1)}^{(\alpha)} \cdots \tilde{L}_{(\alpha+m)}^{(\alpha+m-1)}$.

Proposition A.21 (differentiation (DLMF, 18.9.23)). The differentiation operator for Laguerre is given by

$$\mathcal{D}\mathbf{L}^{(\alpha)} = \mathbf{L}^{(\alpha)} \underbrace{\begin{pmatrix} 0 & -1 & & & \\ & -1 & & & \\ & & -1 & & \\ & & & \ddots & \ddots \end{pmatrix}}_{\tilde{D}_{(\alpha)}^{(\alpha+1)}}.$$

Proposition A.22 (weighted differentiation (DLMF, 18.9.24)). The weighted differentiation operator for Laguerre is given by

$$\mathcal{D}x^{\alpha+1}e^{-x}\mathbf{L}^{(\alpha+1)} = x^\alpha e^{-x} \mathbf{L}^{(\alpha)} \underbrace{\begin{pmatrix} 0 & & & \\ 1 & 2 & & \\ & 3 & \ddots & \\ & & \ddots & \end{pmatrix}}_{\tilde{W}_{(\alpha+1)}^{(\alpha)}}.$$

A.5. Hermite

Proposition A.23 (mass matrix (DLMF, Table 18.3.1)). We have

$$\mathbf{H}^\top e^{-x^2} \mathbf{H} = M_H,$$

where M_H is the diagonal matrix satisfying $(M_H)_{n,n} = \sqrt{\pi} 2^n n!.$

Proposition A.24 (differentiation (DLMF, 18.9.25)). The differentiation operator for Hermite is given by

$$\mathcal{D}\mathbf{H} = \mathbf{H} \underbrace{\begin{pmatrix} 0 & 2 & & \\ & 4 & & \\ & & 6 & \\ & & & \ddots \end{pmatrix}}_{D_H}.$$

Proposition A.25 (weighted differentiation (DLMF, 18.9.26)). The weighted differentiation operator for Hermite is given by

$$\mathcal{D}e^{-x^2} \mathbf{H} = e^{-x^2} \mathbf{H} \underbrace{\begin{pmatrix} 0 & & & \\ -1 & -1 & & \\ & -1 & \ddots & \\ & & \ddots & \end{pmatrix}}_{W_U^T}.$$

A.6. Associated Legendre polynomials

The associated Legendre polynomials satisfy the relationship

$$P_n^m(x) = (1/2)_m (-2)^m (1-x^2)^{m/2} C_{n-m}^{(m+1/2)}(x)$$

or, in quasimatrix notation, we have

$$\mathbf{P}^m = (1/2)_m (-2)^m (1-x^2)^{m/2} \mathbf{C}^{(m+1/2)}.$$

We thus inherit the recurrence relationships from the ultraspherical polynomials; for example,

$$\begin{aligned} (\mathbf{P}^m)^\top \mathbf{P}^m &= \underbrace{(1/2)_m^2 (-2)^{2m} M_{(m+1/2)}}_{M_m}, \\ x \mathbf{P}^m &= \mathbf{P}^m X_{(m+1/2)}. \end{aligned} \quad (\text{A.1})$$

B. Multivariate orthogonal polynomial and their recurrence relationships

B.1. Jacobi polynomials on the triangle

Here, we outline the recurrence relationships for

$$P_{n,k}^{(a,b,c)}(x,y) := P_{n-k}^{(2k+b+c+1,a)}(2x-1)(1-x)^k P_k^{(c,b)}\left(\frac{2y}{1-x}-1\right),$$

which are derived by Xu (2017) and Olver *et al.* (2020). Here, $a,b,c > -1$ and $0 \leq k \leq n$. We define $z := 1 - x - y$ and

$$\frac{\partial}{\partial z} := \frac{\partial}{\partial y} - \frac{\partial}{\partial x}.$$

Proposition B.1 (mass matrix (Dunkl and Xu 2014)). We have

$$(\mathbf{P}^{(a,b,c)})^\top x^a y^b z^c \mathbf{P}^{(a,b,c)} = M_{(a,b,c)},$$

where $M_{(a,b,c)}$ is a diagonal matrix whose diagonal is a block-vector. The k th entry of the n th block is given by

$$\begin{aligned} &\frac{1}{2^{2k+a+2b+2c+3}} (M_{(2k+b+c+1,a)})_{n-k,n-k} (M_{(c,b)})_{k,k} \\ &= \frac{\Gamma(n+k+b+c+2)\Gamma(n-k+a+1)\Gamma(k+b+1)\Gamma(k+c+1)}{(2n+a+b+c+2)(2k+b+c+1)\Gamma(n+k+a+b+c+2)\Gamma(k+b+c+1)(n-k)!k!}. \end{aligned}$$

Corollary B.2 (Olver *et al.* 2020, Corollary 1). The following recurrence relations for the partial derivatives hold:

$$\begin{aligned} &(2k+b+c+1) \frac{\partial}{\partial x} P_{n,k}^{(a,b,c)} \\ &= (n+k+a+b+c+2)(k+b+c+1) P_{n-1,k}^{(a+1,b,c+1)} \\ &\quad + (k+b)(n+k+b+c+1) P_{n-1,k-1}^{(a+1,b,c+1)}, \end{aligned}$$

$$\frac{\partial}{\partial y} P_{n,k}^{(a,b,c)} = (k+b+c+1) P_{n-1,k-1}^{(a,b+1,c+1)},$$

$$\begin{aligned}
& (2k+b+c+1) \frac{\partial}{\partial z} P_{n,k}^{(a,b,c)} \\
&= -(n+k+a+b+c+2)(k+b+c+1) P_{n-1,k}^{(a+1,b+1,c)} \\
&\quad + (k+c)(n+k+b+c+1) P_{n-1,k-1}^{(a+1,b+1,c)}.
\end{aligned}$$

They give the entries of sparse matrices satisfying

$$\begin{aligned}
\partial_x \mathbf{P}^{(a,b,c)} &= \mathbf{P}^{(a+1,b,c+1)} D_{x,(a,b,c)}^{(a+1,b,c+1)}, \\
\partial_y \mathbf{P}^{(a,b,c)} &= \mathbf{P}^{(a+1,b,c+1)} D_{y,(a,b,c)}^{(a+1,b+1,c+1)}, \\
\partial_z \mathbf{P}^{(a,b,c)} &= \mathbf{P}^{(a+1,b+1,c)} D_{z,(a,b,c)}^{(a+1,b+1,c)}.
\end{aligned}$$

Each matrix is block-super-diagonal with diagonal ($D_{y,(a,b,c)}^{(a+1,b+1,c+1)}$) or upper-bidiagonal ($D_{x,(a,b,c)}^{(a+1,b,c+1)}$ and $D_{z,(a,b,c)}^{(a+1,b+1,c)}$) blocks.

Corollary B.3 (Olver et al. 2020, Corollary 2). The following recurrence relations for the weighted partial derivatives hold:

$$\begin{aligned}
& - (2k+b+c+1) \frac{\partial}{\partial x} (x^a y^b z^c P_{n,k}^{(a,b,c)}) \\
&= x^{a-1} y^b z^{c-1} ((k+c)(n-k+1) P_{n+1,k}^{(a-1,b,c-1)} \\
&\quad + (k+1)(n-k+a) P_{n+1,k+1}^{(a-1,b,c-1)}), \\
& \frac{\partial}{\partial y} (x^a y^b z^c P_{n,k}^{(a,b,c)}) = -(k+1) x^a y^{b-1} z^{c-1} P_{n+1,k+1}^{(a,b-1,c-1)}, \\
& (2k+b+c+1) \frac{\partial}{\partial z} (x^a y^b z^c P_{n,k}^{(a,b,c)}) \\
&= x^{a-1} y^{b-1} z^c ((k+b)(n-k+1) P_{n+1,k}^{(a-1,b-1,c)} \\
&\quad - (k+1)(n-k+a) P_{n+1,k+1}^{(a-1,b-1,c)}).
\end{aligned}$$

They give the entries of sparse matrices satisfying

$$\begin{aligned}
\partial_x x^{a+1} y^b z^{c+1} \mathbf{P}^{(a+1,b,c+1)} &= x^a y^b z^c \mathbf{P}^{(a,b,c)} W_{x,(a+1,b,c+1)}^{(a,b,c)}, \\
\partial_y x^a y^{b+1} z^{c+1} \mathbf{P}^{(a,b+1,c+1)} &= x^a y^b z^c \mathbf{P}^{(a,b,c)} W_{y,(a,b+1,c+1)}^{(a,b,c)}, \\
\partial_z x^{a+1} y^{b+1} z^c \mathbf{P}^{(a+1,b+1,c)} &= x^a y^b z^c \mathbf{P}^{(a,b,c)} W_{z,(a+1,b+1,c)}^{(a,b,c)}.
\end{aligned}$$

Each matrix is block-sub-diagonal with diagonal ($W_{y,(a,b+1,c+1)}^{(a,b,c)}$) or lower-bidiagonal ($W_{x,(a+1,b,c+1)}^{(a,b,c)}$ and $W_{z,(a+1,b+1,c)}^{(a,b,c)}$) blocks.

Corollary B.4 (Olver et al. 2020, Corollary 3). The following recurrence relations for conversions hold:

$$\begin{aligned}
 & (2n + a + b + c + 2)P_{n,k}^{(a,b,c)} \\
 &= (n + k + a + b + c + 2)P_{n,k}^{(a+1,b,c)} \\
 &\quad + (n + k + b + c + 1)P_{n-1,k}^{(a+1,b,c)}, \\
 & (2n + a + b + c + 2)(2k + b + c + 1)P_{n,k}^{(a,b,c)} \\
 &= (n + k + a + b + c + 2)(k + b + c + 1)P_{n,k}^{(a,b+1,c)} \\
 &\quad - (n - k + a)(k + b + c + 1)P_{n-1,k}^{(a,b+1,c)} \\
 &\quad + (k + c)(n + k + b + c + 1)P_{n-1,k-1}^{(a,b+1,c)} \\
 &\quad - (k + c)(n - k + 1)P_{n,k-1}^{(a,b+1,c)}, \\
 & (2n + a + b + c + 2)(2k + b + c + 1)P_{n,k}^{(a,b,c)} \\
 &= (n + k + a + b + c + 2)(k + b + c + 1)P_{n,k}^{(a,b,c+1)} \\
 &\quad - (n - k + a)(k + b + c + 1)P_{n-1,k}^{(a,b,c+1)} \\
 &\quad - (k + b)(n + k + b + c + 1)P_{n-1,k-1}^{(a,b,c+1)} \\
 &\quad + (k + b)(n - k + 1)P_{n,k-1}^{(a,b,c+1)}.
 \end{aligned}$$

They give the entries of sparse matrices satisfying

$$\begin{aligned}
 \mathbf{P}^{(a,b,c)} &= \mathbf{P}^{(a+1,b,c)} R_{(a,b,c)}^{(a+1,b,c)}, \\
 \mathbf{P}^{(a,b,c)} &= \mathbf{P}^{(a,b+1,c)} R_{(a,b,c)}^{(a,b+1,c)}, \\
 \mathbf{P}^{(a,b,c)} &= \mathbf{P}^{(a,b,c+1)} R_{(a,b,c)}^{(a,b,c+1)}.
 \end{aligned}$$

Each matrix is block-upper-bidiagonal with diagonal ($R_{(a,b,c)}^{(a+1,b,c)}$) or upper-bidiagonal ($R_{(a,b,c)}^{(a,b+1,c)}$ and $R_{(a,b,c)}^{(a,b,c+1)}$) blocks.

Corollary B.5 (Olver et al. 2020, Corollary 4). The following recurrence relations for lowering operators hold:

$$\begin{aligned}
 & (2n + a + b + c + 2)xP_{n,k}^{(a,b,c)} \\
 &= (n - k + a)P_{n,k}^{(a-1,b,c)} + (n - k + 1)P_{n+1,k}^{(a-1,b,c)},
 \end{aligned}$$

$$\begin{aligned}
& (2k+b+c+1)(2n+a+b+c+2)yP_{n,k}^{(a,b,c)} \\
&= (k+b)(n+k+b+c+1)P_{n,k}^{(a,b-1,c)} \\
&\quad - (k+1)(n-k+a)P_{n,k+1}^{(a,b-1,c)} \\
&\quad - (k+b)(n-k+1)P_{n+1,k}^{(a,b-1,c)} \\
&\quad + (k+1)(n+k+a+b+c+2)P_{n+1,k+1}^{(a,b-1,c)}, \\
& (2k+b+c+1)(2n+a+b+c+2)zP_{n,k}^{(a,b,c)} \\
&= (k+c)(n+k+b+c+1)P_{n,k}^{(a,b,c-1)} \\
&\quad + (k+1)(n-k+a)P_{n,k+1}^{(a,b,c-1)} \\
&\quad - (k+c)(n-k+1)P_{n+1,k}^{(a,b,c-1)} \\
&\quad - (k+1)(n+k+a+b+c+2)P_{n+1,k+1}^{(a,b,c-1)}.
\end{aligned}$$

They give the entries of sparse matrices satisfying

$$\begin{aligned}
x\mathbf{P}^{(a+1,b,c)} &= \mathbf{P}^{(a,b,c)}L_{(a+1,b,c)}^{(a,b,c)}, \\
y\mathbf{P}^{(a,b+1,c)} &= \mathbf{P}^{(a,b,c)}L_{(a,b+1,c)}^{(a,b,c)}, \\
z\mathbf{P}^{(a,b,c+1)} &= \mathbf{P}^{(a,b,c)}L_{(a,b,c+1)}^{(a,b,c)}.
\end{aligned}$$

Each matrix is block-lower-bidiagonal with diagonal ($L_{(a+1,b,c)}^{(a,b,c)}$) or lower-bidiagonal ($L_{(a,b+1,c)}^{(a,b,c)}$ and $L_{(a,b,c+1)}^{(a,b,c)}$) blocks.

B.2. Dirichlet polynomials on the triangle

Imposing Dirichlet conditions requires a special basis, which we denote by $Q_{n,k}^{(a,b,c)}$, and forms the quasimatrix $\mathbf{Q}^{(a,b,c)}$. It is related to negative parameter OPs on the triangle, which are also Sobolev OPs; see Xu (2017). We define $\tilde{P}(x) = P(2x - 1)$ to be the shifted Legendre polynomials (Legendre polynomials transplanted onto $[0, 1]$.)

Definition B.6. The following polynomials vanish at $x = 0$ apart from when $k = n$:

$$\begin{aligned}
Q_{n,k}^{(1,0,0)}(x,y) &:= xP_{n-1,k}^{(1,0,0)}(x,y), \quad 0 \leq k \leq n-1, \\
Q_{n,n}^{(1,0,0)}(x,y) &:= P_{n,n}(x,y).
\end{aligned}$$

The following polynomials vanish at $y = 0$ apart from when $k = 0$:

$$\begin{aligned}
Q_{n,0}^{(0,1,0)}(x,y) &:= \tilde{P}_n^{(0,0)}(x), \\
Q_{n,k}^{(0,1,0)}(x,y) &:= yP_{n-1,k-1}^{(0,1,0)}(x,y), \quad 1 \leq k \leq n.
\end{aligned}$$

The following polynomials vanish at $z = 0$ (*i.e.* $y = 1 - x$) apart from when $k = 0$:

$$\begin{aligned} Q_{n,0}^{(0,0,1)}(x,y) &:= \tilde{P}_n^{(0,0)}(x), \\ Q_{n,k}^{(0,0,1)}(x,y) &:= z P_{n-1,k-1}^{(0,0,1)}(x,y), \quad 1 \leq k \leq n. \end{aligned}$$

The following polynomials vanish at $x = 0$ and $y = 0$ apart from when $k = 0, n$:

$$\begin{aligned} Q_{0,0}^{(1,1,0)}(x,y) &:= 1, \\ Q_{n,0}^{(1,1,0)}(x,y) &:= x \tilde{P}_{n-1}^{(0,1)}(x), \\ Q_{n,k}^{(1,1,0)}(x,y) &:= xy P_{n-2,k-1}^{(1,1,0)}(x,y), \quad 1 \leq k \leq n-1, \\ Q_{n,n}^{(1,1,0)}(x,y) &:= y P_{n-1,n-1}^{(0,1,0)}(x,y). \end{aligned}$$

The following polynomials vanish at $x = 0$ and $z = 0$ apart from when $k = 0, n$:

$$\begin{aligned} Q_{0,0}^{(1,0,1)}(x,y) &:= 1, \\ Q_{n,0}^{(1,0,1)}(x,y) &:= x \tilde{P}_{n-1}^{(0,1)}(x), \\ Q_{n,k}^{(1,0,1)}(x,y) &:= xz P_{n-2,k-1}^{(1,0,1)}(x,y), \quad 1 \leq k \leq n-1, \\ Q_{n,n}^{(1,0,1)}(x,y) &:= z P_{n-1,n-1}^{(0,0,1)}(x,y). \end{aligned}$$

The following polynomials vanish at $y = 0$ and $z = 0$ apart from when $k = 0, 1$:

$$\begin{aligned} Q_{0,0}^{(0,1,1)}(x,y) &:= 1, \\ Q_{n,0}^{(0,1,1)}(x,y) &:= (1-x) P_{n-1,0}(x,y) = (1-x) \tilde{P}_{n-1}^{(1,0)}(x), \\ Q_{n,1}^{(0,1,1)}(x,y) &:= (1-x-2y) P_{n-1,0}(x,y) = (1-x-2y) \tilde{P}_{n-1}^{(1,0)}(x), \\ Q_{n,k}^{(0,1,1)}(x,y) &:= yz P_{n-2,k-2}^{(0,1,1)}(x,y), \quad 2 \leq k \leq n. \end{aligned}$$

The following polynomials vanish at $x = 0$, $y = 0$ and $z = 0$ apart from when $k = 0, 1$, and n :

$$\begin{aligned} Q_{0,0}^{(1,1,1)}(x,y) &:= 1, \\ Q_{1,0}^{(1,1,1)}(x,y) &:= 1-2x, \\ Q_{1,1}^{(1,1,1)}(x,y) &:= 1-x-2y, \\ Q_{n,0}^{(1,1,1)}(x,y) &:= x(1-x) P_{n-2,0}^{(1,0,0)}(x,y) = x(1-x) P_{n-2}^{(1,1)}(x), \\ Q_{n,1}^{(1,1,1)}(x,y) &:= x(1-x-2y) P_{n-2,0}^{(1,0,0)}(x,y) = x(1-x-2y) P_{n-2}^{(1,1)}(x), \end{aligned}$$

$$\begin{aligned} Q_{n,k}^{(1,1,1)}(x,y) &:= xyzP_{n-3,k-2}^{(1,1,1)}(x,y), \quad 2 \leq k \leq n-1, \\ Q_{n,n}^{(1,1,1)}(x,y) &:= yzP_{n-2,n-2}^{(0,1,1)}(x,y). \end{aligned}$$

Proposition B.7 (conversion (Olver *et al.* 2019, Corollary C.1–3)).
The following recurrence relationships hold:

$$\begin{aligned} Q_{0,0}^{(1,0,0)}(x,y) &= P_{0,0}(x,y), \\ (2n+1)Q_{n,k}^{(1,0,0)}(x,y) &= (n-k)[P_{n,k}(x,y) + P_{n-1,k}(x,y)], \\ Q_{n,n}^{(1,0,0)}(x,y) &= P_{n,n}(x,y), \\ (2n+1)Q_{n,0}^{(0,1,0)}(x,y) &= (n+1)P_{n,0}(x,y) - nP_{n-1,0}(x,y), \\ (2n+1)Q_{n,k}^{(0,1,0)}(x,y) &= (n+k+1)P_{n,k}(x,y) - (n-k+1)P_{n,k-1}(x,y) \\ &\quad - (n-k)P_{n-1,k}(x,y) + (n+k)P_{n-1,k-1}(x,y), \\ (2n+1)Q_{n,0}^{(0,0,1)}(x,y) &= (n+1)P_{n,0}(x,y) - nP_{n-1,0}(x,y), \\ (2n+1)Q_{n,k}^{(0,0,1)}(x,y) &= -(n+k+1)P_{n,k}(x,y) - (n-k+1)P_{n,k-1}(x,y) \\ &\quad + (n-k)P_{n-1,k}(x,y) + (n+k)P_{n-1,k-1}(x,y), \\ Q_{0,0}^{(1,1,0)}(x,y) &= Q_{0,0}^{(1,0,0)}(x,y), \\ 2nQ_{n,0}^{(1,1,0)}(x,y) &= (n+1)Q_{n,0}^{(1,0,0)}(x,y) - nQ_{n-1,0}^{(1,0,0)}(x,y), \\ 4nQ_{n,k}^{(1,1,0)}(x,y) &= (n+k+1)Q_{n,k}^{(1,0,0)}(x,y) - (n-k)Q_{n,k-1}^{(1,0,0)}(x,y) \\ &\quad + (k-n)Q_{n-1,k}^{(1,0,0)}(x,y) + (n+k-1)Q_{n-1,k-1}^{(1,0,0)}(x,y), \\ 2Q_{n,n}^{(1,1,0)}(x,y) &= Q_{n,n}^{(1,0,0)}(x,y) - Q_{n,n-1}^{(1,0,0)}(x,y) + Q_{n-1,n-1}^{(1,0,0)}(x,y), \\ 2Q_{n,0}^{(1,1,0)}(x,y) &= Q_{n,0}^{(0,1,0)}(x,y) + Q_{n-1,0}^{(0,1,0)}(x,y), \\ 2nQ_{n,k}^{(1,1,0)}(x,y) &= (n-k)[Q_{n,k}^{(0,1,0)}(x,y) + Q_{n-1,k}^{(0,1,0)}(x,y)], \\ Q_{n,n}^{(1,1,0)}(x,y) &= Q_{n,n}^{(0,1,0)}(x,y), \\ Q_{0,0}^{(1,0,1)}(x,y) &= Q_{0,0}^{(1,0,0)}(x,y), \\ 2nQ_{n,0}^{(1,0,1)}(x,y) &= (n+1)Q_{n,0}^{(1,0,0)}(x,y) - nQ_{n-1,0}^{(1,0,0)}(x,y), \\ 4nQ_{n,k}^{(1,0,1)}(x,y) &= -(n+k+1)Q_{n,k}^{(1,0,0)}(x,y) - (n-k)Q_{n,k-1}^{(1,0,0)}(x,y) \\ &\quad + (n-k)Q_{n-1,k}^{(1,0,0)}(x,y) + (n+k-1)Q_{n-1,k-1}^{(1,0,0)}(x,y), \\ 2Q_{n,n}^{(1,0,1)}(x,y) &= -Q_{n,n}^{(1,0,0)}(x,y) - Q_{n,n-1}^{(1,0,0)}(x,y) + Q_{n-1,n-1}^{(1,0,0)}(x,y), \\ 2Q_{n,0}^{(1,0,1)}(x,y) &= Q_{n,0}^{(0,0,1)}(x,y) + Q_{n-1,0}^{(0,0,1)}(x,y), \\ 2nQ_{n,k}^{(1,0,1)}(x,y) &= (n-k)[Q_{n,k}^{(0,0,1)}(x,y) + Q_{n-1,k}^{(0,0,1)}(x,y)], \end{aligned}$$

$$\begin{aligned}
2Q_{n,0}^{(0,1,1)}(x,y) &= -Q_{n,0}^{(0,1,0)}(x,y) + Q_{n-1,0}^{(0,1,0)}(x,y), \\
2nQ_{n,1}^{(0,1,1)}(x,y) &= -2(n+1)Q_{n,1}^{(0,1,0)}(x,y) - nQ_{n,0}^{(0,1,0)}(x,y) \\
&\quad + 2(n-1)Q_{n-1,1}^{(0,1,0)}(x,y) + nQ_{n-1,0}^{(0,1,0)}(x,y), \\
2n(2k-1)Q_{n,k}^{(0,1,1)}(x,y) &= -(k-1)(n+k)Q_{n,k}^{(0,1,0)}(x,y) - (k-1)(n-k+1)Q_{n,k-1}^{(0,1,0)}(x,y) \\
&\quad + (k-1)(n-k)Q_{n-1,k}^{(0,1,0)}(x,y) + (k-1)(n+k-1)Q_{n-1,k-1}^{(0,1,0)}(x,y), \\
2Q_{n,0}^{(0,1,1)}(x,y) &= -Q_{n,0}^{(0,0,1)}(x,y) + Q_{n-1,0}^{(0,0,1)}(x,y), \\
2nQ_{n,1}^{(0,1,1)}(x,y) &= 2(n+1)Q_{n,1}^{(0,1,0)}(x,y) + nQ_{n,0}^{(0,0,1)}(x,y) \\
&\quad - 2(n-1)Q_{n-1,1}^{(0,1,0)}(x,y) - nQ_{n-1,0}^{(0,0,1)}(x,y), \\
2n(2k-1)Q_{n,k}^{(0,1,1)}(x,y) &= (k-1)(n+k)Q_{n,k}^{(0,0,1)}(x,y) - (k-1)(n-k+1)Q_{n,k-1}^{(0,0,1)}(x,y) \\
&\quad - (k-1)(n-k)Q_{n-1,k}^{(0,0,1)}(x,y) + (k-1)(n+k-1)Q_{n-1,k-1}^{(0,0,1)}(x,y), \\
Q_{0,0}^{(1,1,1)}(x,y) &= Q_{0,0}^{(0,1,1)}(x,y), \\
Q_{1,0}^{(1,1,1)}(x,y) &= 2Q_{1,0}^{(0,1,1)}(x,y) - Q_{0,0}^{(0,1,1)}(x,y), \\
Q_{1,1}^{(1,1,1)}(x,y) &= Q_{1,1}^{(0,1,1)}(x,y), \\
(2n-1)Q_{n,0}^{(1,1,1)}(x,y) &= (n-1)[Q_{n,0}^{(0,1,1)}(x,y) + Q_{n-1,0}^{(0,1,1)}(x,y)], \\
(2n-1)Q_{n,k}^{(1,1,1)}(x,y) &= (n-k)[Q_{n,k}^{(0,1,1)}(x,y) + Q_{n-1,k}^{(0,1,1)}(x,y)], \\
Q_{n,n}^{(1,1,1)}(x,y) &= Q_{n,n}^{(0,1,1)}(x,y), \\
Q_{0,0}^{(1,1,1)}(x,y) &= Q_{0,0}^{(1,0,1)}(x,y), \\
Q_{1,0}^{(1,1,1)}(x,y) &= -2Q_{1,0}^{(1,0,1)}(x,y) + Q_{0,0}^{(1,0,1)}(x,y), \\
Q_{1,1}^{(1,1,1)}(x,y) &= 2Q_{1,1}^{(1,0,1)}(x,y) + Q_{1,0}^{(1,0,1)}(x,y) - Q_{0,0}^{(1,0,1)}(x,y), \\
(2n-1)Q_{n,0}^{(1,1,1)}(x,y) &= (n-1)[-Q_{n,0}^{(1,0,1)}(x,y) + Q_{n-1,0}^{(1,0,1)}(x,y)], \\
(2n-1)Q_{n,1}^{(1,1,1)}(x,y) &= 2(n+1)Q_{n,1}^{(1,0,1)}(x,y) + (n-1)Q_{n,0}^{(1,0,1)}(x,y) \\
&\quad - 2(n-1)Q_{n-1,1}^{(1,0,1)}(x,y) - (n-1)Q_{n-1,0}^{(1,0,1)}(x,y)], \\
(2n-1)(2k-1)Q_{n,k}^{(1,1,1)}(x,y) &= (n+k)(k-1)Q_{n,k}^{(1,0,1)}(x,y) - (n-k)(k-1)Q_{n,k-1}^{(1,0,1)}(x,y) \\
&\quad - (n-k)(k-1)Q_{n-1,k}^{(1,0,1)}(x,y) + (n+k-2)(k-1)Q_{n-1,k-1}^{(1,0,1)}(x,y),
\end{aligned}$$

$$\begin{aligned}
& (2n-1)Q_{n,n}^{(1,1,1)}(x,y) \\
&= (n-1)[Q_{n,n}^{(1,0,1)}(x,y) - Q_{n,n-1}^{(1,0,1)}(x,y) + Q_{n-1,n-1}^{(1,0,1)}(x,y)], \\
& Q_{0,0}^{(1,1,1)}(x,y) = Q_{0,0}^{(1,1,0)}(x,y), \\
& Q_{1,0}^{(1,1,1)}(x,y) = -2Q_{1,0}^{(1,1,0)}(x,y) + Q_{0,0}^{(1,1,0)}(x,y), \\
& Q_{1,1}^{(1,1,1)}(x,y) = -2Q_{1,1}^{(1,1,0)}(x,y) - Q_{1,0}^{(1,1,0)}(x,y) + Q_{0,0}^{(1,1,0)}(x,y), \\
& (2n-1)Q_{n,0}^{(1,1,1)}(x,y) = (n-1)[-Q_{n,0}^{(1,1,0)}(x,y) + Q_{n-1,0}^{(1,1,0)}(x,y)], \\
& (2n-1)Q_{n,1}^{(1,1,1)}(x,y) = -2(n+1)Q_{n,1}^{(1,1,0)}(x,y) - (n-1)Q_{n,0}^{(1,1,0)}(x,y) \\
&\quad + 2(n-1)Q_{n-1,1}^{(1,1,0)}(x,y) + (n-1)Q_{n-1,0}^{(1,1,0)}(x,y)], \\
& (2n-1)(2k-1)Q_{n,k}^{(1,1,1)}(x,y) \\
&= -(n+k)(k-1)Q_{n,k}^{(1,1,0)}(x,y) - (n-k)(k-1)Q_{n,k-1}^{(1,1,0)}(x,y) \\
&\quad + (n-k)(k-1)Q_{n-1,k}^{(1,1,0)}(x,y) + (n+k-2)(k-1)Q_{n-1,k-1}^{(1,1,0)}(x,y), \\
& (2n-1)Q_{n,n}^{(1,1,1)}(x,y) \\
&= (n-1)[-Q_{n,n}^{(1,1,0)}(x,y) - Q_{n,n-1}^{(1,1,0)}(x,y) + Q_{n-1,n-1}^{(1,1,0)}(x,y)].
\end{aligned}$$

These define the conversion operators

$$\mathbf{Q}^{(a,b,c)} = \mathbf{P}^{(0,0,0)} \tilde{\mathbf{R}}_{(a,b,c)}^{(0,0,0)}.$$

Proposition B.8 (derivatives (Olver *et al.* 2019, Corollary C.4)).
The following recurrence relationships hold:

$$\begin{aligned}
& \frac{\partial}{\partial y} Q_{n,0}^{(0,1,1)}(x,y) = 0, \\
& \frac{\partial}{\partial y} Q_{n,1}^{(0,1,1)}(x,y) = -2P_{n-1,0}(x,y), \\
& \frac{\partial}{\partial y} Q_{n,k}^{(0,1,1)}(x,y) = (1-k)P_{n-1,k-1}(x,y), \\
& \frac{\partial}{\partial x} Q_{n,0}^{(1,0,1)}(x,y) = nP_{n-1,0}(x,y), \\
& \frac{\partial}{\partial x} Q_{n,k}^{(1,0,1)}(x,y) = \frac{k-n}{2}[P_{n-1,k-1}(x,y) + P_{n-1,k}(x,y)], \\
& \frac{\partial}{\partial x} Q_{n,n}^{(1,0,1)}(x,y) = -nP_{n-1,n-1}(x,y), \\
& \frac{\partial}{\partial z} Q_{n,0}^{(1,1,0)}(x,y) = -nP_{n-1,0}(x,y),
\end{aligned}$$

$$\begin{aligned}\frac{\partial}{\partial z} Q_{n,k}^{(1,1,0)}(x,y) &= \frac{n-k}{2}[P_{n-1,k-1}(x,y) - P_{n-1,k}(x,y)], \\ \frac{\partial}{\partial z} Q_{n,n}^{(1,1,0)}(x,y) &= nP_{n-1,n-1}(x,y).\end{aligned}$$

These define the derivative operators

$$\begin{aligned}\mathcal{D}_x \mathbf{Q}^{(1,0,1)} &= \mathbf{P}^{(0,0,0)} \tilde{D}_{x,(1,0,1)}^{(0,0,0)}, \\ \mathcal{D}_y \mathbf{Q}^{(0,1,1)} &= \mathbf{P}^{(0,0,0)} \tilde{D}_{x,(0,1,1)}^{(0,0,0)}, \\ \mathcal{D}_z \mathbf{Q}^{(1,1,0)} &= \mathbf{P}^{(0,0,0)} \tilde{D}_{z,(1,1,0)}^{(0,0,0)}, \\ \mathcal{D}_x \mathbf{Q}^{(1,1,1)} &= \mathbf{P}^{(0,0,0)} \underbrace{\tilde{D}_{x,(1,0,1)}^{(0,0,0)}}_{\tilde{D}_{x,(1,1,1)}^{(0,0,0)}} \tilde{R}_{(1,1,1)}^{(1,0,1)}, \\ \mathcal{D}_y \mathbf{Q}^{(1,1,1)} &= \mathbf{P}^{(0,0,0)} \underbrace{\tilde{D}_{y,(0,1,1)}^{(0,0,0)}}_{\tilde{D}_{y,(1,1,1)}^{(0,0,0)}} \tilde{R}_{(1,1,1)}^{(0,1,1)}, \\ \mathcal{D}_z \mathbf{Q}^{(1,1,1)} &= \mathbf{P}^{(0,0,0)} \underbrace{\tilde{D}_{x,(1,1,0)}^{(0,0,0)}}_{\tilde{D}_{z,(1,1,1)}^{(0,0,0)}} \tilde{R}_{(1,1,1)}^{(1,1,0)}.\end{aligned}$$

Proposition B.9 (restriction (Olver *et al.* 2019, Proposition B.2)).
The restriction operator to $x = 0$ is given by

$$\begin{aligned}Q_{n,n}^{(1,0,0)}(0,y) &:= \tilde{P}_n^{(0,0)}(y), \\ Q_{n,k}^{(1,0,0)}(0,y) &:= 0, \quad k \neq n.\end{aligned}$$

The restriction operator to $y = 0$ is given by

$$\begin{aligned}Q_{n,0}^{(0,1,0)}(x,0) &:= \tilde{P}_n^{(0,0)}(x), \\ Q_{n,k}^{(0,1,0)}(x,0) &:= 0, \quad k \neq 0.\end{aligned}$$

The restriction operator to $z = 0$ is given by

$$\begin{aligned}Q_{n,0}^{(0,0,1)}(x,1-x) &:= \tilde{P}_n^{(0,0)}(x), \\ Q_{n,k}^{(0,0,1)}(x,1-x) &:= 0, \quad k \neq 0.\end{aligned}$$

We denote these as

$$\begin{aligned}\mathbf{Q}^{(0,1,0)}(x,0) &= \mathbf{P}(2x-1)\tilde{R}_{(0,1,0)}^{P,x=0}, \\ \mathbf{Q}^{(1,0,0)}(0,y) &= \mathbf{P}(2y-1)\tilde{R}_{(1,0,0)}^{P,y=0}, \\ \mathbf{Q}^{(0,0,1)}(x,1-x) &= \mathbf{P}(2x-1)\tilde{R}_{(0,0,1)}^{P,z=0}.\end{aligned}$$

These then give

$$\begin{aligned} \mathbf{Q}^{(1,1,1)}(x,0) &= \mathbf{P}(2x-1) \underbrace{\tilde{R}_{(0,1,0)}^{P,x=0}}_{\tilde{R}_{(1,1,1)}^{P,x=0}} \tilde{R}_{(1,1,1)}^{(0,1,0)}, \\ \mathbf{Q}^{(1,1,1)}(0,y) &= \mathbf{P}(2y-1) \underbrace{\tilde{R}_{(1,0,0)}^{P,y=0}}_{\tilde{R}_{(1,1,1)}^{P,y=0}} \tilde{R}_{(1,1,1)}^{(1,0,0)}, \\ \mathbf{Q}^{(1,1,1)}(x,1-x) &= \mathbf{P}(2x-1) \underbrace{\tilde{R}_{(0,0,1)}^{P,z=0}}_{\tilde{R}_{(1,1,1)}^{P,z=0}} \tilde{R}_{(1,1,1)}^{(0,0,1)}. \end{aligned}$$

C. Semiclassical orthogonal polynomials

Given a family of classical orthogonal polynomials

$$p_{n+1}(x) = (A_n x + B_n) p_n(x) - C_n p_{n-1}(x), \quad p_0(x) = 1,$$

the *associated* orthogonal polynomials are those polynomials with the same initial conditions that use the classical recurrence coefficients with indices offset by $c \in \mathbb{N}$ units:

$$p_{n+1}(x; c) = (A_{n+c} x + B_{n+c}) p_n(x; c) - C_{n+c} p_{n-1}(x; c).$$

They too satisfy a linear homogeneous differential equation (Zarzo, Ronveaux and Godoy 1993):

$$\begin{aligned} &- \sigma^2 p_n^{(iv)} - 5\sigma\sigma' p_n''' \\ &+ [\tau^2 - 2\tau\sigma' - 3\sigma'^2 + (2n+4c)\sigma\tau' - (4+n-n^2+4c-2nc-2c^2)\sigma\sigma''] p_n'' \\ &+ \frac{3}{2}[2\tau\tau' + (2n-2+4c)\sigma'\tau' - 2\tau\sigma'' + (n^2-n-4c+2nc+2c^2)\sigma'\sigma''] p_n' \\ &= \frac{1}{4}n(n+2)[(n+2c-3)\sigma'' + 2\tau'] [(n+2c-1)\sigma'' + 2\tau'] p_n. \end{aligned}$$

Associated orthogonal polynomials have applications through their diagonalization of the following Stieltjes transforms (Erdélyi *et al.* 1953, vol. 2, p. 162 (6)):

$$p_n(x; c+1) = \frac{1}{A_c \int_D d\mu(t; c)} \int_D \frac{p_{n+1}(x; c) - p_{n+1}(t; c)}{x-t} d\mu(t; c).$$

(Generalized) co-recursive orthogonal polynomials are related to the associated orthogonal polynomials. They too satisfy fourth-order linear homogeneous differential equations, though we refer the interested reader to Ronveaux and Marcellán (1989) and in particular to Ronveaux, Zarzo and Godoy (1995) for their coefficients.

Konoplev orthogonal polynomials (Konoplev 1961) are another semiclassical case that are orthogonal with respect to $L^2([-1,1],|x|^\gamma(1-x^2)^\alpha dx)$. They are the polynomial eigenfunctions of

$$\{(1-x^2)^{-\alpha}|x|^{2-\gamma}(-\mathcal{D})[|x|^\gamma(1-x^2)^{\alpha+1}\mathcal{D}]+\delta_n\gamma\}u=n(n+2\alpha+\gamma+1)x^2u,$$

where $\delta_n=(1-(-1)^n)/2$. Note that there is a sign error in Konoplev (1961). Konoplev orthogonal polynomials are an example where the weight is a product of algebraic factors. Other examples arise when the domain is a subset of $[-1,1]$. It is also possible to develop a linear differential equation with polynomial coefficients for these orthogonal polynomials (Magnus 1995).

Bochner (1929) and Krall (1938) characterized orthogonal polynomial systems that are solutions of linear differential eigenvalue problems with polynomial coefficients independent of the degree. Their generalized polynomials are all orthogonal with respect to measures of classical orthogonal polynomials with added Dirac distributions at the endpoints. They may satisfy fourth, sixth or infinite-order linear differential eigenvalue problems.

D. Notation

D.1. Continuous operators

$$\mathcal{D}u=\frac{du}{dx}, \quad \mathcal{D}_x u=\frac{\partial u}{\partial x}, \quad \mathcal{D}_y u=\frac{\partial u}{\partial y},$$

$$\mathcal{Q}u=\int u(x)dx \text{ with the constant defined in Proposition A.17,}$$

$$\mathcal{H}u=\frac{1}{\pi}\int_a^b \frac{u(t)}{x-t}dt,$$

$$\mathcal{L}u=\frac{1}{\pi}\int_a^b u(t)\log|t-x|dt.$$

D.2. Univariate

Quasimatrix notation:

$$\mathbf{P}^{(\alpha,\beta)}=(P_0^{(\alpha,\beta)} \mid P_1^{(\alpha,\beta)} \mid \dots)$$

where $P_n^{(\alpha,\beta)}(x)$ are Jacobi polynomials,

$$\mathbf{P}=(P_0 \mid P_1 \mid \dots)$$

where $P_n(x)$ are Legendre polynomials,

$$\mathbf{C}^{(\lambda)}=(C_0^{(\lambda)} \mid C_1^{(\lambda)} \mid \dots)$$

where $C_n^{(\lambda)}(x)$ are ultraspherical polynomials,

$$\mathbf{T} = (T_0 \mid T_1 \mid \cdots)$$

where $T_n(x)$ are Chebyshev polynomials (1st kind),

$$\mathbf{U} = (U_0 \mid U_1 \mid \cdots)$$

where $U_n(x)$ are Chebyshev polynomials (2nd kind),

$$\tilde{\mathbf{P}}^{(a,b)} = \mathbf{P}^{(a,b)} M_{(a,b)}^{-1/2}$$

normalized Jacobi polynomials,

$$\tilde{\mathbf{P}} = \mathbf{P} M_P^{-1/2}$$

normalized Legendre polynomials,

$$\tilde{\mathbf{C}}^{(\lambda)} = \mathbf{C}^{(\lambda)} M_{(\lambda)}^{-1/2}$$

normalized ultraspherical polynomials,

$$\tilde{\mathbf{T}} = \mathbf{T} M_T^{-1/2}$$

normalized Chebyshev polynomials (1st kind),

$$\tilde{\mathbf{U}} = \mathbf{U} M_U^{-1/2}$$

normalized Chebyshev polynomials (1st kind),

$$\hat{\mathbf{P}}^{(a,b)} = (1-x)^{a/2}(1+x)^{b/2} \mathbf{P}$$

normalized, weighted Jacobi polynomials,

$$\hat{\mathbf{C}}^{(\lambda)} = (1-x^2)^{\lambda/2-1/4} \mathbf{C}^{(\lambda)}$$

normalized, weighted ultraspherical polynomials.

Multiplication operators:

$$x\mathbf{P}^{(\alpha,\beta)} = \mathbf{P}^{(\alpha,\beta)} X_{(\alpha,\beta)},$$

$$x\mathbf{C}^{(\lambda)} = \mathbf{C}^{(\lambda)} X_{(\lambda)},$$

$$x\mathbf{P} = \mathbf{P} X_P,$$

$$x\mathbf{T} = \mathbf{T} X_T,$$

$$x\mathbf{L}^{(\alpha)} = \mathbf{L}^{(\alpha)} \tilde{X}_{(\alpha)},$$

$$x\mathbf{H} = \mathbf{H} X_H.$$

Conversion operators:

$$\mathbf{P}^{(\alpha,\beta)} = \mathbf{P}^{(\alpha+\kappa,\beta+\nu)} R_{(\alpha,\beta)}^{(\alpha+\kappa,\beta+\nu)},$$

$$\mathbf{C}^{(\lambda)} = \mathbf{C}^{(\lambda+\kappa)} R_{(\lambda)}^{(\lambda+\kappa)},$$

$$\mathbf{T} = \mathbf{U} R_T^U,$$

$$\mathbf{L}^{(\alpha)} = \mathbf{L}^{(\alpha+\kappa)} \tilde{R}_{(\alpha)}^{(\alpha+\kappa)}.$$

Weighted conversion operators:

$$\begin{aligned}(1-x)^\kappa(1+x)^\nu \mathbf{P}^{(\alpha+\kappa, \beta+\nu)} &= \mathbf{P}^{(\alpha, \beta)} L_{(\alpha+\kappa, \beta+\nu)}^{(\alpha, \beta)}, \\ (1-x^2)^\kappa \mathbf{C}^{(\lambda+\kappa)} &= \mathbf{C}^{(\lambda)} L_{(\lambda+\kappa)}^{(\lambda)}, \\ (1-x^2)\mathbf{U} &= \mathbf{T} L_U^T, \\ x^\kappa \mathbf{L}^{(\alpha+\kappa)} &= \mathbf{L}^{(\alpha)} \tilde{L}_{(\alpha+\kappa)}^{(\alpha)}.\end{aligned}$$

Differentiation operators:

$$\begin{aligned}\mathcal{D}^\kappa \mathbf{P}^{(\alpha, \beta)} &= \mathbf{P}^{(\alpha+\kappa, \beta+\kappa)} D_{(\alpha, \beta)}^{(\alpha+\kappa, \beta+\kappa)}, \\ \mathcal{D}^\kappa \mathbf{C}^{(\lambda)} &= \mathbf{C}^{(\lambda+\kappa)} D_{(\lambda)}^{(\lambda+\kappa)}, \\ \mathcal{D}^\kappa \mathbf{P} &= \mathbf{C}^{(\kappa+1/2)} D_P^{(\kappa+1/2)} = \mathbf{P}^{(\kappa, \kappa)} D_P^{(\kappa, \kappa)}, \\ \mathcal{D}\mathbf{T} &= \mathbf{U} D_T^U, \quad \mathcal{D}^\kappa \mathbf{T} = \mathbf{C}^{(\lambda)} D_T^{(\kappa)}, \\ \mathcal{D}^\kappa \mathbf{L}^{(\alpha)} &= \mathbf{L}^{(\alpha+\kappa)} \tilde{D}_{(\alpha)}^{(\alpha+\kappa)}, \\ \mathcal{D}\mathbf{H} &= \mathbf{H} D_H.\end{aligned}$$

Weighted differentiation operators:

$$\begin{aligned}\mathcal{D}^\kappa (1-x)^{\alpha+\kappa} (1+x)^{\beta+\kappa} \mathbf{P}^{(\alpha+\kappa, \beta+\kappa)} &= (1-x)^\alpha (1+x)^\beta \mathbf{P}^{(\alpha, \beta)} W_{(\alpha+\kappa, \beta+\kappa)}^{(\alpha, \beta)}, \\ \mathcal{D}^\kappa (1-x^2)^{\lambda+\kappa} \mathbf{C}^{(\lambda+\kappa)} &= (1-x^2)^\lambda \mathbf{C}^{(\lambda)} W_{(\lambda+\kappa)}^{(\lambda)}, \\ \mathcal{D} \sqrt{1-x^2} \mathbf{U} &= \frac{\mathbf{T}}{\sqrt{1-x^2}} W_U^T, \\ \mathcal{D}^\kappa x^{\alpha+\kappa} e^{-x} \mathbf{L}^{(\alpha+\kappa)} &= x^\alpha e^{-x} \mathbf{L}^{(\alpha)} \tilde{W}_{(\alpha+\kappa)}^{(\alpha)}, \\ \mathcal{D} e^{-x^2} \mathbf{H} &= e^{-x^2} \mathbf{H} W_H.\end{aligned}$$

One-sided weighted differentiation operators:

$$\begin{aligned}\mathcal{D}^\kappa (1-x)^{\alpha+\kappa} \mathbf{P}^{(\alpha+\kappa, \beta)} &= (1-x)^\alpha \mathbf{P}^{(\alpha, \beta+\kappa)} W_{\text{L}, (\alpha, \beta)}^{(\alpha, \beta+\kappa)}, \\ \mathcal{D} (1+x)^{\beta+\kappa} \mathbf{P}^{(\alpha, \beta+\kappa)} &= (1+x)^\beta \mathbf{P}^{(\alpha+\kappa, \beta)} W_{\text{R}, (\alpha, \beta)}^{(\alpha+\kappa, \beta)}.\end{aligned}$$

Singular integral operators:

$$\begin{aligned}\mathcal{H} \frac{\mathbf{T}}{\sqrt{1-x^2}} &= \mathbf{U} H_T^U, \\ \mathcal{H} \sqrt{1-x^2} \mathbf{U} &= \mathbf{T} H_U^T, \\ \mathcal{L} \frac{\mathbf{T}}{\sqrt{1-x^2}} &= \mathbf{T} L_T.\end{aligned}$$

D.3. Bivariate orthogonal polynomials

Quasimatrix notation:

$$\mathbf{P}^{(a,b,c,d)} = \mathbf{P}^{(a,b)} \otimes \mathbf{P}^{(c,d)},$$

$$\mathbf{C}^{(\lambda,\mu)} = \mathbf{C}^{(\lambda)} \otimes \mathbf{C}^{(\mu)},$$

$$\mathbf{T}^2 = \mathbf{T} \otimes \mathbf{T},$$

$$\mathbf{P}^{(a,b,c)} = (P_0^{(a,b,c)} | \dots)$$

where $P_n^{(a,b,c)}(x,y)$ are Jacobi polynomials on the unit simplex.

Multiplication operators:

$$x\mathbf{P}^{(a,b,c,d)} = \mathbf{P}^{(a,b,c,d)} X_{(a,b,c,d)}, \quad y\mathbf{P}^{(a,b,c,d)} = \mathbf{P}^{(a,b,c,d)} Y_{(a,b,c,d)},$$

$$x\mathbf{C}^{(\lambda,\mu)} = \mathbf{C}^{(\lambda,\mu)} \hat{X}_{(\lambda,\mu)}, \quad y\mathbf{C}^{(\lambda,\mu)} = \mathbf{C}^{(\lambda,\mu)} \hat{Y}_{(\lambda,\mu)},$$

$$x\mathbf{T}^2 = \mathbf{T}^2 X_{T,T}, \quad y\mathbf{T}^2 = \mathbf{T}^2 Y_{T,T},$$

$$x\mathbf{P}^{(a,b,c)} = \mathbf{P}^{(a,b,c)} X_{(a,b,c)}, \quad y\mathbf{P}^{(a,b,c)} = \mathbf{P}^{(a,b,c)} Y_{(a,b,c)}.$$

Conversion operators:

$$\mathbf{P}^{(a,b,c,d)} = \mathbf{P}^{(a+\kappa,b+\nu,c+\eta,d+\delta)} R_{(a,b,c,d)}^{(a+\kappa,b+\nu,c+\eta,d+\delta)},$$

$$\mathbf{C}^{(\lambda,\mu)} = \mathbf{C}^{(\lambda+\kappa,\mu+\nu)} \hat{R}_{(\lambda,\mu)}^{(\lambda+\kappa,\mu+\nu)},$$

$$\mathbf{T}^2 = \mathbf{C}^{(\lambda,\mu)} R_{T,T}^{(\lambda,\mu)},$$

$$\mathbf{P}^{(a,b,c)} = \mathbf{P}^{(a+\kappa,b+\nu,c+\eta)} R_{(a,b,c)}^{(a+\kappa,b+\nu,c+\eta)}.$$

Weighted conversion operators ($z = 1 - x - y$):

$$x^\kappa y^\nu z^\eta \mathbf{P}^{(a+\kappa,b+\nu,c+\eta)} = \mathbf{P}^{(a,b,c)} L_{(a+\kappa,b+\nu,c+\eta)}^{(a,b,c)}.$$

Differentiation operators ($\mathcal{D}_z = \mathcal{D}_y - \mathcal{D}_x$, $\Delta = \mathcal{D}_x^2 + \mathcal{D}_y^2$):

$$\mathcal{D}_x^\kappa \mathbf{P}^{(a,b,c)} = \mathbf{P}^{(a+\kappa,b+\nu,c+\kappa)} D_{x,(a,b,c)}^{(a+\kappa,b+\nu,c+\kappa)},$$

$$\mathcal{D}_y^\kappa \mathbf{P}^{(a,b,c)} = \mathbf{P}^{(a+\nu,b+\kappa,c+\kappa)} D_{y,(a,b,c)}^{(a+\nu,b+\kappa,c+\kappa)},$$

$$\mathcal{D}_z^\kappa \mathbf{P}^{(a,b,c)} = \mathbf{P}^{(a+\kappa,b+\kappa,c+\nu)} D_{z,(a,b,c)}^{(a+\kappa,b+\kappa,c+\nu)},$$

$$\Delta \mathbf{P}^{(a,b,c)} = \mathbf{P}^{(a+2,b+2,c+2)} \Delta_{(a,b,c)}^{(a+2,b+2,c+2)}.$$

Weighted differentiation operators ($z = 1 - x - y$, $\mathcal{D}_z = \mathcal{D}_y - \mathcal{D}_x$, $\Delta = \mathcal{D}_x^2 + \mathcal{D}_y^2$):

$$\mathcal{D}_x x^\kappa y^\nu z^\kappa \mathbf{P}^{(a+\kappa,b+\nu,c+\kappa)} = \mathbf{P}^{(a,b,c)} W_{x,(a+\kappa,b+\nu,c+\kappa)}^{(a,b,c)},$$

$$\mathcal{D}_y x^\nu y^\kappa z^\kappa \mathbf{P}^{(a+\nu,b+\kappa,c+\kappa)} = \mathbf{P}^{(a,b,c)} W_{y,(a+\nu,b+\kappa,c+\kappa)}^{(a,b,c)},$$

$$\begin{aligned} \mathcal{D}_z x^\kappa y^\kappa z^\nu P^{(a+\kappa, b+\kappa, c+\nu)} &= P^{(a, b, c)} W_{z, (a+\kappa, b+\kappa, c+\nu)}^{(a, b, c)}, \\ \Delta xyz P^{(1, 1, 1)} &= P^{(1, 1, 1)} \Delta_W. \end{aligned}$$

REFERENCES¹³

- B. K. Alpert and V. Rokhlin (1991), ‘A fast algorithm for the evaluation of Legendre expansions’, *SIAM J. Sci. Statist. Comput.* **12**, 158–179.
- G. E. Andrews, R. Askey and R. Roy (1999), *Special Functions*, Cambridge University Press.
- K. Atkinson and W. Han (2012), *Spherical Harmonics and Approximations on the Unit Sphere: An Introduction*, Vol. 2044 of *Lecture Notes in Mathematics*, Springer.
- J. L. Aurentz and R. M. Slevinsky (2020), ‘On symmetrizing the ultraspherical spectral method for self-adjoint problems’, *J. Comput. Phys.* **410**, 109383.
- J. L. Aurentz and L. N. Trefethen (2017), ‘Chopping a Chebyshev series’, *ACM Trans. Math. Software* **43**, 33: 1–33:21.
- D. H. Bailey, K. Jeyabalan and X. S. Li (2004), ‘A comparison of three high-precision quadrature schemes’, *Exp. Math.* **14**, 317–329.
- M. Berry (2007), ‘Focused tsunami waves’, *Proc. Roy. Soc. A Math. Phys. Engrg Sci.* **463**(2007), 3055–3071.
- S. Beuchler and J. Schöberl (2006), ‘New shape functions for triangular p -FEM using integrated Jacobi polynomials’, *Numer. Math.* **103**, 339–366.
- S. Bochner (1929), ‘Über Sturm–Liouville Polynomsysteme’, *Math. Z.* **29**, 730–736.
- I. Bogaert (2014), ‘Iteration-free computation of Gauss–Legendre quadrature nodes and weights’, *SIAM J. Sci. Comput.* **36**, A1008–A1026.
- I. Bogaert, B. Michiels and J. Fostier (2012), ‘ $\mathcal{O}(1)$ computation of Legendre polynomials and Gauss–Legendre nodes and weights for parallel computing’, *SIAM J. Sci. Comput.* **34**, C83–C101.
- C. F. Borges and W. B. Gragg (1993), A parallel divide and conquer algorithm for the generalized real symmetric definite tridiagonal eigenproblem. In *Numerical Linear Algebra and Scientific Computing* (L. Reichel, A. Ruttan and R. S. Varga, eds), de Gruyter, pp. 11–29.
- J. P. Boyd (2013), ‘Finding the zeros of a univariate equation: Proxy rootfinders, Chebyshev interpolation, and the companion matrix’, *SIAM Rev.* **55**, 375–396.
- J. Bremer and H. Yang (2019), ‘Fast algorithms for Jacobi expansions via nonoscillatory phase functions’, *IMA J. Numer. Anal.* doi:10.1093/imanum/drz016
- J. R. Bunch, C. P. Nielsen and D. C. Sorensen (1978), ‘Rank-one modification of the symmetric eigenproblem’, *Numer. Math.* **31**, 31–48.

¹³ The URLs cited in this work were correct at the time of going to press, but the publisher and the authors make no undertaking that the citations remain live or are accurate or appropriate.

- K. J. Burns, G. M. Vasil, J. S. Oishi, D. Lecoanet and B. P. Brown (2020), ‘Dedalus: A flexible framework for numerical simulations with spectral methods’, *Phys. Rev. Research* **2**, 023068.
- S. Chandrasekaran and M. Gu (2004), ‘A divide-and-conquer algorithm for the eigendecomposition of symmetric block-diagonal plus semiseparable matrices’, *Numer. Math.* **96**, 723–731.
- S. Chen, J. Shen and L.-L. Wang (2016), ‘Generalized Jacobi functions and their applications to fractional differential equations’, *Math. Comp.* **85**, 1603–1638.
- C. W. Clenshaw (1955), ‘A note on the summation of Chebyshev series’, *Math. Comp.* **9**, 118–120.
- C. W. Clenshaw (1957), ‘The numerical solution of linear differential equations in Chebyshev series’, *Math. Proc. Cambridge Philos. Soc.* **53**, 134–149.
- C. W. Clenshaw and A. R. Curtis (1960), ‘A method for numerical integration on an automatic computer’, *Numer. Math.* **2**, 197–205.
- J. W. Cooley and J. W. Tukey (1965), ‘An algorithm for the machine calculation of complex Fourier series’, *Math. Comp.* **19**, 297–301.
- E. Coutsias, T. Hagstrom and D. Torres (1996), ‘An efficient spectral method for ordinary differential equations with rational function coefficients’, *Math. Comp.* **65**, 611–635.
- C. R. Crawford (1973), ‘Reduction of a band-symmetric generalized eigenvalue problem’, *Comm. Assoc. Comput. Mach.* **16**, 41–44.
- J. J. M. Cuppen (1981), ‘A divide and conquer method for the symmetric tridiagonal eigenproblem’, *Numer. Math.* **36**, 177–195.
- F. Dai and Y. Xu (2013), *Approximation Theory and Harmonic Analysis on Spheres and Balls*, Monographs in Mathematics, Springer.
- P. Davis and P. Rabinowitz (1956), ‘Abscissas and weights for Gaussian quadratures of high order’, *J. Res. Nat. Bur. Standards* **56**, 35–37.
- P. Davis and P. Rabinowitz (1958), ‘Additional abscissas and weights for Gaussian quadratures of high order: Values for $n = 64, 80$, and 96 ’, *J. Res. Nat. Bur. Standards* **60**, 613–614.
- B. de F. Bayly (1938), ‘(iii) Gauss’ quadratic formula with twelve ordinates’, *Biometrika* **30**, 193–194.
- P. Deift (1999), *Orthogonal Polynomials and Random Matrices: A Riemann–Hilbert Approach*, American Mathematical Society.
- E. H. Doha and W. M. Abd-Elhameed (2002), ‘Efficient spectral-Galerkin algorithms for direct solution of second-order equations using ultraspherical polynomials’, *SIAM J. Sci. Comput.* **24**, 548–571.
- J. R. Driscoll, D. M. Healy Jr and D. N. Rockmore (1997), ‘Fast discrete polynomial transforms with applications to data analysis for distance transitive graphs’, *SIAM J. Sci. Comput.* **26**, 1066–1099.
- T. A. Driscoll, N. Hale and L. N. Trefethen, eds (2014), *Chebfun Guide*, Pafnuty Publications.
- Q. Du, M. Gunzburger, R. B. Lehoucq and K. Zhou (2012), ‘Analysis and approximation of nonlocal diffusion problems with volume constraints’, *SIAM Rev.* **54**, 667–696.

- Q. Du, M. Gunzburger, R. B. Lehoucq and K. Zhou (2013), ‘A nonlocal vector calculus, nonlocal volume-constrained problems, and nonlocal balance laws’, *Math. Mod. Meth. Appl. Sci.* **23**, 493–540.
- M. Dubiner (1991), ‘Spectral methods on triangles and other domains’, *J. Sci. Comput.* **6**, 345–390.
- C. F. Dunkl and Y. Xu (2014), *Orthogonal Polynomials of Several Variables*, second edition, Encyclopedia of Mathematics and its Applications, Cambridge University Press.
- A. Dutt and V. Rokhlin (1993), ‘Fast Fourier transforms for nonequispaced data’, *SIAM J. Sci. Comput.* **14**, 1368–1393.
- A. Edelman and B. D. Sutton (2007), ‘From random matrices to stochastic operators’, *J. Statist. Phys.* **127**, 1121–1165.
- D. Elliott (1982), ‘The classical collocation method for singular integral equations’, *SIAM J. Numer. Anal.* **19**, 816–832.
- L. Elsner, A. Fasse and E. Langmann (1997), ‘A divide-and-conquer method for the tridiagonal generalized eigenvalue problem’, *J. Comput. Appl. Math.* **86**, 141–148.
- A. Erdélyi *et al.*, eds (1953), *Higher Transcendental Functions*, Vol. 2, McGraw-Hill.
- L. Fejér (1933), ‘On the infinite sequences arising in the theories of harmonic analysis, of interpolation, and of mechanical quadratures’, *Bull. Amer. Math. Soc.* **39**, 521–534.
- K.-J. Förster and K. Petras (1990), ‘On estimates for the weights in Gaussian quadrature in the ultraspherical case’, *Math. Comp.* **55**, 243–264.
- K.-J. Förster and K. Petras (1993), ‘Inequalities for the zeros of ultraspherical polynomials and Bessel functions’, *Z. Angew. Math. Mech.* **73**, 232–236.
- D. Fortunato and A. Townsend (2019), ‘Fast Poisson solvers for spectral methods’, *IMA J. Numer. Anal.* doi:10.1093/imanum/drz034
- L. Gatteschi (1979), On the construction of some Gaussian quadrature rules. In *Numerische Integration*, Springer, pp. 138–146.
- C. F. Gauss (1815), *Methodus nova integralium valores per approximationem inveniendi*. Apud Henricum Dieterich.
- W. Gautschi (1967), ‘Computational aspects of three-term recurrence relations’, *SIAM Rev.* **9**, 24–82.
- W. Gautschi (1981a), *A survey of Gauss–Christoffel quadrature formulae*. In *E. B. Christoffel* (P. L. Butzer and F. Fehér, eds), Birkhäuser.
- W. Gautschi (1981b), ‘Minimal solutions of three-term recurrence relations and orthogonal polynomials’, *Math. Comp.* **36**, 547–554.
- W. Gautschi (1982), ‘On generating orthogonal polynomials’, *SIAM J. Sci. Statist. Comput.* **3**, 289–317.
- W. Gautschi (1992), ‘On mean convergence of extended Lagrange interpolation’, *J. Comput. Appl. Math.* **43**, 19–35.
- W. Gautschi (2004), *Orthogonal Polynomials: Computation and Approximation*, Clarendon Press.
- W. Gautschi and J. Wimp (1987), ‘Computing the Hilbert transform of a Jacobi weight function’, *BIT Numer. Math.* **27**, 203–215.

- H. J. Gawlik (1958), *Zeros of Legendre Polynomials of Orders 2–64 and Weight Coefficients of Gauss Quadrature Formulae*, Armament Research and Development Establishment.
- W. M. Gentleman (1972), ‘Implementing Clenshaw–Curtis quadrature, II: Computing the cosine transformation’, *Comm. Assoc. Comput. Mach.* **15**, 343–346.
- A. Gil, J. Segura and N. M. Temme (2018), ‘Asymptotic approximations to the nodes and weights of Gauss–Hermite and Gauss–Laguerre quadratures’, *Stud. Appl. Math.* **140**, 298–332.
- A. Gil, J. Segura and N. M. Temme (2019a), ‘Fast, reliable and unrestricted iterative computation of Gauss–Hermite and Gauss–Laguerre quadratures’, *Numer. Math.* **143**, 649–682.
- A. Gil, J. Segura and N. M. Temme (2019b), ‘Noniterative computation of Gauss–Jacobi quadrature’, *SIAM J. Sci. Comput.* **41**, A668–A693.
- A. Glaser, X. Liu and V. Rokhlin (2007), ‘A fast algorithm for the calculation of the roots of special functions’, *SIAM J. Sci. Comput.* **29**, 1420–1438.
- G. H. Golub and C. F. V. Loan (2013), *Matrix Computations*, fourth edition, The Johns Hopkins University Press.
- G. H. Golub and J. H. Welsch (1969), ‘Calculation of Gauss quadrature rules’, *Math. Comput.* **23**, 221–230.
- L. Greengard (1991), ‘Spectral integration and two-point boundary value problems’, *SIAM J. Numer. Anal.* **28**, 1071–1080.
- L. Greengard and J.-Y. Lee (2004), ‘Accelerating the nonuniform fast Fourier transform’, *SIAM Rev.* **46**, 443–454.
- L. Greengard and V. Rokhlin (1987), ‘A fast algorithm for particle simulations’, *J. Comput. Phys.* **73**, 325–348.
- M. Gu and S. C. Eisenstat (1994), ‘A stable and efficient algorithm for the rank-one modification of the symmetric eigenproblem’, *SIAM J. Matrix Anal. Appl.* **15**, 1266–1276.
- M. Gu and S. C. Eisenstat (1995), ‘A divide-and-conquer algorithm for the symmetric tridiagonal eigenproblem’, *SIAM J. Matrix Anal. Appl.* **16**, 172–191.
- R. H. Gutierrez and P. A. A. Laura (1995), ‘Vibrations of non-uniform rings studied by means of the differential quadrature method’, *J. Sound Vibration* **185**, 507–513.
- N. Hale and S. Olver (2018), ‘A fast and spectrally convergent algorithm for rational-order fractional integral and differential equations’, *SIAM J. Sci. Comput.* **40**, A2456–A2491.
- N. Hale and A. Townsend (2013), ‘Fast and accurate computation of Gauss–Legendre and Gauss–Jacobi quadrature nodes and weights’, *SIAM J. Sci. Comput.* **35**, A652–A674.
- N. Hale and A. Townsend (2014), ‘A fast, simple, and stable Chebyshev–Legendre transform using an asymptotic formula’, *SIAM J. Sci. Comput.* **36**, A148–A167.
- D. M. Healy Jr, D. N. Rockmore, P. J. Kostelec and S. Moore (2003), ‘FFTs for the 2-sphere: Improvements and variations’, *J. Fourier Anal. Appl.* **9**, 341–385.
- W. Heinrichs (1989), ‘Spectral methods with sparse matrices’, *Numer. Math.* **56**, 25–41.

- D. Huybrechs (2010), ‘On the Fourier extension of nonperiodic functions’, *SIAM J. Numer. Anal.* **47**, 4326–4355.
- E. A. Hylleraas (1929), ‘Neue Berechnung der Energie des Heliums im Grundzustande, sowie des tiefsten Terms von Ortho-Helium’, *Z. Physik* **54**, 347–366.
- K. Ishioka (2018), ‘A new recurrence formula for efficient computation of spherical harmonic transform’, *J. Met. Soc. Japan* **96**, 241–249.
- C. G. J. Jacobi (1826), ‘Ueber Gauss neue Methode, die Werthe der Integrale näherungsweise zu finden’, *J. Reine Angew. Math.* **1**, 301–308.
- K. Julien and M. Watson (2009), ‘Efficient multi-dimensional solution of PDEs using Chebyshev spectral methods’, *J. Comput. Phys.* **228**, 1480–1503.
- G. Karniadakis and S. Sherwin (2013), *Spectral/hp Element Methods for Computational Fluid Dynamics*, Oxford University Press.
- L. Kaufman (1993), ‘An algorithm for the banded symmetric generalized matrix eigenvalue problem’, *SIAM J. Matrix Anal. Appl.* **14**, 372–389.
- L. Kaufman (2000), ‘Band reduction algorithms revisited’, *ACM Trans. Math. Software* **26**, 551–567.
- J. Keiner (2008), ‘Gegenbauer polynomials and semiseparable matrices’, *Electron. Trans. Numer. Anal.* **30**, 26–53.
- J. Keiner (2009), ‘Computing with expansions in Gegenbauer polynomials’, *SIAM J. Sci. Comput.* **31**, 2151–2171.
- J. Keiner (2011), *Fast Polynomial Transforms*, Logos.
- A. Klöckner, A. Barnett, L. Greengard and M. O’Neil (2013), ‘Quadrature by expansion: A new method for the evaluation of layer potentials’, *J. Comput. Phys.* **252**, 332–349.
- V. P. Konoplev (1961), ‘Polynomials orthogonal with respect to weight functions which are zero or infinite at isolated points of the interval of orthogonality’, *Dokl. Akad. Nauk SSSR* **141**, 781–784.
- T. Koornwinder (1975), Two-variable analogues of the classical orthogonal polynomials. In *Theory and Application of Special Functions* (R. Askey, ed.), pp. 435–495.
- H. L. Krall (1938), ‘Certain differential equations for Tchebycheff polynomials’, *Duke Math. J.* **4**, 705–718.
- M. Krishnapur, B. Rider and B. Virág (2016), ‘Universality of the stochastic Airy operator’, *Commun. Pure Appl. Math.* **69**, 145–199.
- S. Kunis and D. Potts (2003), ‘Fast spherical Fourier algorithms’, *J. Comp. Appl. Math.* **161**, 75–98.
- C. Lanczos (1938), ‘Trigonometric interpolation of empirical and analytical functions’, *J. Math. Phys.* **17**, 123–199.
- J.-Y. Lee and L. Greengard (1997), ‘A fast adaptive numerical method for stiff two-point boundary value problems’, *SIAM J. Sci. Comput.* **18**, 403–429.
- F. G. Lether (1978), ‘On the construction of Gauss–Legendre quadrature rules’, *J. Comput. Appl. Math.* **4**, 47–52.
- H. Li and J. Shen (2010), ‘Optimal error estimates in Jacobi-weighted Sobolev spaces for polynomial approximations on the triangle’, *Math. Comp.* **79**, 1621–1646.

- P. W. Livermore (2010), ‘Galerkin orthogonal polynomials’, *J. Comput. Phys.* **229**, 2046–2060.
- C. H. Love (1966), Abscissas and weights for Gaussian quadrature for $N = 2$ to 100, and $N = 125, 150, 175, 200$. United States Department Of Commerce, National Bureau of Standards Monograph.
- A. N. Lowan, N. Davids and A. Levenson (1942), ‘Table of the zeros of the Legendre polynomials of order 1–16 and the weight coefficients for Gauss’ mechanical quadrature formula’, *Bull. Amer. Math. Soc.* **48**, 739–742.
- P.-O. Löwdin and H. Shull (1956), ‘Natural orbitals in the quantum theory of two-electron systems’, *Phys. Rev.* **101**, 1730–1739.
- D. W. Lozier (1980), *Numerical solution of linear difference equations. NBSIR Technical Report 80-1976*, National Bureau of Standards.
- A. P. Magnus (1995), ‘Painlevé-type differential equations for the recurrence coefficients of semi-classical orthogonal polynomials’, *J. Comput. Appl. Math.* **57**, 215–237.
- P. G. Martinsson and V. Rokhlin (2007), ‘An accelerated kernel-independent fast multipole method in one dimension’, *SIAM J. Sci. Comput.* **29**, 1160–1178.
- J. Mason (1993), ‘Chebyshev polynomials of the second, third and fourth kinds in approximation, indefinite integration, and integral transforms’, *J. Comput. Appl. Math.* **49**, 169–178.
- M. J. Mohlenkamp (1999), ‘A fast transform for spherical harmonics’, *J. Fourier Anal. Appl.* **5**, 159–184.
- B. P. Moors (1905), *Valeur approximative d’une intégrale définie*, Gauthier-Villars.
- A. Mori, R. Suda and M. Sugihara (1999), ‘An improvement on Orszag’s fast algorithm for Legendre polynomial transform’, *Trans. Info. Process. Soc. Japan* **40**, 3612–3615.
- N. I. Muskhelishvili (1953), *Singular Integral Equations*, second edition, Dover.
- A. Narayan and J. S. Hesthaven (2012), ‘Computation of connection coefficients and measure modifications for orthogonal polynomials’, *BIT Numer. Math.* **52**, 457–483.
- E. J. Nyström (1930), ‘Über die praktische auflösung von Integralgleichungen mit Anwendungen auf Randwertaufgaben’, *Acta Math.* **54**, 185–204.
- F. W. J. Olver (1967), ‘Numerical solution of second-order linear difference equations’, *J. Res. Nat. Bur. Standards* **71B**, 111–129.
- F. W. J. Olver, D. W. Lozier, R. F. Boisvert and C. W. Clark, eds (2010), *NIST Handbook of Mathematical Functions*, Cambridge University Press.
- S. Olver (2011), ‘Computing the Hilbert transform and its inverse’, *Math. Comp.* **80**, 1745–1767.
- S. Olver (2012), ‘A general framework for solving Riemann–Hilbert problems numerically’, *Numer. Math.* **122**, 305–340.
- S. Olver and A. Townsend (2013), ‘A fast and well-conditioned spectral method’, *SIAM Rev.* **55**, 462–489.
- S. Olver and Y. Xu (2019a), Non-homogeneous wave equation on a cone. arXiv:1907.08286
- S. Olver and Y. Xu (2019b), ‘Orthogonal structure on a wedge and on the boundary of a square’, *Found. Comput. Math.* **19**, 561–589.

- S. Olver and Y. Xu (2020a), ‘Orthogonal polynomials in and on a quadratic surface of revolution’, *Math. Comp.*, to appear.
- S. Olver and Y. Xu (2020b), ‘Orthogonal structure on a quadratic curve’, *IMA J. Numer. Anal.*, to appear.
- S. Olver, G. Goretkin, R. M. Slevinsky and A. Townsend (2016), Julia package for function approximation. GitHub. <https://github.com/ApproxFun/ApproxFun.jl>
- S. Olver, A. Townsend and G. Vasil (2019), ‘A sparse spectral method on triangles’, *SIAM J. Sci. Comput.* **41**, A3728–A3756.
- S. Olver, A. Townsend and G. M. Vasil (2020), *Recurrence relations for a family of orthogonal polynomials on a triangle*. In *Spectral and High Order Methods for Partial Differential Equations (ICOSAHOM) 2018*, Springer.
- S. A. Orszag (1986), Fast eigenfunction transforms. *Science and Computers* **10**, 13–30.
- B. N. Parlett (1974), ‘The Rayleigh quotient iteration and some generalizations for nonnormal matrices’, *Math. Comp.* **28**, 679–693.
- T. Pearcey (1946), ‘The structure of an electromagnetic field in the neighbourhood of a cusp of a caustic’, *Philos. Mag.* **37**, 311–317.
- K. Petras (1999), ‘On the computation of the Gauss–Legendre quadrature formula with a given precision’, *J. Comput. Appl. Math.* **112**, 253–267.
- D. Potts and G. Steidl (2003), ‘Fast summation at nonequispaced knots by NFFTs’, *SIAM J. Sci. Comput.* **24**, 2013–2037.
- D. Potts, G. Steidl and M. Tasche (1998), ‘Fast and stable algorithms for discrete spherical Fourier transforms’, *Linear Algebra Appl.* **275–276**, 433–450.
- M. Püschel and J. M. Moura (2003), ‘The algebraic approach to the discrete cosine and sine transforms and their fast algorithms’, *SIAM J. Comput.* **32**, 1280–1316.
- L. Rayleigh (1937), *The Theory of Sound*, Macmillan.
- M. Reinecke and D. S. Seljebotn (2013), ‘Libsharp: Spherical harmonic transforms revisited’, *Astronomy Astrophys.* **554**, A112.
- V. Rokhlin and M. Tygert (2006), ‘Fast algorithms for spherical harmonic expansions’, *SIAM J. Sci. Comput.* **27**, 1903–1928.
- A. Ronveaux and F. Marcellán (1989), ‘Co-recursive orthogonal polynomials and fourth-order differential equation’, *J. Comput. Appl. Math.* **25**, 105–109.
- A. Ronveaux, A. Zarzo and E. Godoy (1995), ‘Fourth-order differential equations satisfied by the generalized co-recursive of all classical orthogonal polynomials: A study of their distribution of zeros’, *J. Comput. Appl. Math.* **59**, 295–328.
- D. Ruiz-Antolín and A. Townsend (2018), ‘A nonuniform fast Fourier transform based on low rank approximation’, *SIAM J. Sci. Comput.* **40**, A529–A547.
- H. Rutishauser (1962), ‘On a modification of the QD-algorithm with Graeffe-type convergence’, *Z. Angew. Math. Phys.* **13**, 493–496.
- N. Schaeffer (2013), ‘Efficient spherical harmonic transforms aimed at pseudospectral numerical simulations’, *Geochem. Geophys. Geosyst.* **14**, 751–758.
- H. R. Schwarz (1968), ‘Tridiagonalization of a symmetric band matrix’, *Numer. Math.* **12**, 231–241.
- J. Shen (1994), ‘Efficient spectral-Galerkin method I: Direct solvers of second- and fourth-order equations using Legendre polynomials’, *SIAM J. Sci. Comput.* **15**, 1489–1505.

- H. Shull and P.-O. Löwdin (1955), ‘Role of the continuum in superposition of configurations’, *J. Chem. Phys.* **23**, 1362–1363.
- R. M. Slevinsky (2016), Julia package for fast orthogonal polynomial transforms. GitHub. <https://github.com/MikaelSlevinsky/FastTransforms.jl>
- R. M. Slevinsky (2017), Conquering the pre-computation in two-dimensional harmonic polynomial transforms. [arXiv:1711.07866](https://arxiv.org/abs/1711.07866)
- R. M. Slevinsky (2018a), Fast orthogonal polynomial transforms. GitHub. <https://github.com/MikaelSlevinsky/FastTransforms>
- R. M. Slevinsky (2018b), ‘On the use of Hahn’s asymptotic formula and stabilized recurrence for a fast, simple, and stable Chebyshev–Jacobi transform’, *IMA J. Numer. Anal.* **38**, 102–124.
- R. M. Slevinsky (2019), ‘Fast and backward stable transforms between spherical harmonic expansions and bivariate Fourier series’, *Appl. Comput. Harmon. Anal.* **47**, 585–606.
- R. M. Slevinsky and S. Olver (2017), ‘A fast and well-conditioned spectral method for singular integral equations’, *J. Comput. Phys.* **332**, 290–315.
- R. M. Slevinsky, H. Montanelli and Q. Du (2018), ‘A spectral method for nonlocal diffusion operators on the sphere’, *J. Comput. Phys.* **372**, 893–911.
- B. Snowball and S. Olver (2020), ‘Sparse spectral and p -finite element methods for partial differential equations on disk slices and trapeziums’, *Stud. Appl. Math.*, to appear.
- A. Sommariva (2013), ‘Fast construction of Fejér and Clenshaw–Curtis rules for general weight functions’, *Comput. Math. Appl.* **65**, 682–693.
- N. J. Sonine (1887), ‘Über die angenäherte Berechnung der bestimmten Integrale und über die dabei vorkommenden ganzen Functionen’, *Warsaw Univ. Izv.* **18**, 1–76.
- A. H. Stroud and D. Secrest (1966), *Gaussian Quadrature Formulas*, Prentice Hall.
- R. Suda and M. Takami (2002), ‘A fast spherical harmonics transform algorithm’, *Math. Comp.* **71**, 703–715.
- P. N. Swarztrauber (2003), ‘On computing the points and weights for Gauss–Legendre quadrature’, *SIAM J. Sci. Comput.* **24**, 945–954.
- G. Szegő (1975), *Orthogonal Polynomials*, fourth edition, American Mathematical Society.
- H. Tallqvist (1905), ‘Grunderna af teorin fioer sferiska funktioner jöoemte anvioendninggar inom fysiken’, *Finska litteratur-sööellskapets tryckeri*.
- D. J. Torres and E. A. Coutsias (1999), ‘Pseudospectral solution of the two-dimensional Navier–Stokes equations in a disk’, *SIAM J. Sci. Comput.* **21**, 378–403.
- A. Townsend (2015), ‘The race for high order Gauss–Legendre quadrature’, *SIAM News* **48**, 1–3.
- A. Townsend and S. Olver (2015), ‘The automatic solution of partial differential equations using a global spectral method’, *J. Comput. Phys.* **299**, 106–123.
- A. Townsend and L. N. Trefethen (2013), ‘An extension of Chebfun to two dimensions’, *SIAM J. Sci. Comput.* **35**, C495–C518.
- A. Townsend, T. Trogdon and S. Olver (2016), ‘Fast computation of Gauss quadrature nodes and weights on the whole real line’, *IMA J. Numer. Anal.* **36**, 337–358.

- A. Townsend, M. Webb and S. Olver (2018), ‘Fast polynomial transforms based on Toeplitz and Hankel matrices’, *Math. Comp.* **87**, 1913–1934.
- L. N. Trefethen (2008), ‘Is Gauss quadrature better than Clenshaw–Curtis?’, *SIAM Rev.* **50**, 67–87.
- T. Trogdon and S. Olver (2016), *Riemann–Hilbert Problems, Their Numerical Solution and the Computation of Nonlinear Special Functions*, SIAM.
- M. Tygert (2008), ‘Fast algorithms for spherical harmonic expansions, II’, *J. Comput. Phys.* **227**, 4260–4279.
- M. Tygert (2010a), ‘Fast algorithms for spherical harmonic expansions, III’, *J. Comput. Phys.* **229**, 6181–6192.
- M. Tygert (2010b), ‘Recurrence relations and fast algorithms’, *Appl. Comput. Harmon. Anal.* **28**, 121–128.
- G. M. Vasil, K. J. Burns, D. Lecoanet, S. Olver, B. P. Brown and J. S. Oishi (2016), ‘Tensor calculus in polar coordinates using Jacobi polynomials’, *J. Comput. Phys.* **325**, 53–73.
- D. Viswanath (2015), ‘Spectral integration of linear boundary value problems’, *J. Comput. Appl. Math.* **290**, 159–173.
- J. Vogel, J. Xia, S. Cauley and V. Balakrishnan (2016), ‘Superfast divide-and-conquer method and perturbation analysis for structured eigenvalue solutions’, *SIAM J. Sci. Comput.* **38**, A1358–A1382.
- J. Waldvogel (2003), ‘Fast construction of the Fejér and Clenshaw–Curtis quadrature rules’, *BIT Numer. Math.* **43**, 001–018.
- M. Webb (2017), Isospectral algorithms, Toeplitz matrices and orthogonal polynomials. PhD thesis, University of Cambridge.
- E. Wegert (2012), *Visual Complex Functions: An Introduction with Phase Portraits*, Springer Science & Business Media.
- E. J. Weniger (2019), ‘Comment on “Fourier transform of hydrogen-type atomic orbitals”’, *Canad. J. Phys.* **97**, 1349–1360.
- H. Wilber, A. Townsend and Gs. B. Wright (2017), ‘Computing with functions in spherical and polar geometries II: The disk’, *SIAM J. Sci. Comput.* **39**, C238–C262.
- Y. Xu (2017), ‘Approximation and orthogonality in Sobolev spaces on a triangle’, *Const. Approx.* **46**, 349–434.
- E. Yakimiw (1996), ‘Accurate computation of weights in classical Gauss–Christoffel quadrature rules’, *J. Comput. Phys.* **129**, 403–430.
- A. Zarzo, A. Ronveaux and E. Godoy (1993), ‘Fourth-order differential equation satisfied by the associated of any order of all classical orthogonal polynomials: A study of their distribution of zeros’, *J. Comput. Appl. Math.* **49**, 349–359.
- M. Zayernouri and G. E. Karniadakis (2014), ‘Fractional spectral collocation method’, *SIAM J. Sci. Comput.* **36**, A40–A62.
- A. Zebib (1984), ‘A Chebyshev method for the solution of boundary value problems’, *J. Comput. Phys.* **53**, 443–455.
- F. Zernike (1934), ‘Beugungstheorie des Schneidenverfahrens und seiner verbesserten Form, der Phasenkontrastmethode’, *Physica* **1**, 689–704.

